

ABG-VBG Analysis

Brian Locke, Anila Mehta

Table of contents

| | |
|---|-----------|
| 1 Data Pre-processing | 1 |
| 1.0.1 Package Set Up | 1 |
| 1.1 Helper functions for model diagnostics | 1 |
| 1.1.1 Configuration for the IPW models | 2 |
| 1.1.2 Outcome Variable Creation | 5 |
| 1.2 2) Baseline tables | 6 |
| 1.3 3) Three-level PCO ₂ categories (unweighted) | 16 |
| 1.4 4) Restricted cubic spline regressions (unweighted) | 20 |
| 1.4.1 4.1 Unweighted, Restricted Cubic Spline Regression - ABG by PaCO ₂ | 21 |
| 1.4.2 4.2 Unweighted, Restricted Cubic Spline - VBG | 23 |
| 2 Inverse Propensity Weighting | 29 |
| 2.0.1 5.1 ABG IPW weighting and diagnostics | 29 |
| 2.0.2 5.2 ABG IPW spline models | 31 |
| 2.0.3 5.3 ABG IPW spline models (2–98th percentile) | 33 |
| 2.0.4 5.4 VBG IPW weighting and spline models | 37 |
| 2.0.5 5.5 Three-level PCO ₂ categories (weighted; ABG, VBG) | 43 |
| 2.1 6) Propensity score diagnostics | 47 |
| 3 Multiple Imputation Analysis | 50 |
| 3.0.1 7.2 Missingness structure and drivers | 50 |
| 3.0.2 7.3 Monte Carlo error check after MI | 52 |
| 3.1 8) Pre-imputation data prep (consistent types & predictors) | 52 |

| | | |
|-------|---|-----|
| 3.2 | 9) Imputation model specification (MICE) | 52 |
| 3.2.1 | 9.1 Predictor matrix & methods. Run MICE (moderate settings for scale) | 52 |
| 3.3 | 10) Refit propensity models within each imputation | 84 |
| 3.3.1 | FAIL-FAST CHECKS | 84 |
| 3.3.2 | 10.1 ABG propensity (has_abg) | 84 |
| 3.3.3 | 10.2 Balance diagnostics across imputations | 84 |
| 3.3.4 | 10.3 VBG propensity (has_vbg) | 86 |
| 3.3.5 | 10.4 VBG balance | 87 |
| 3.4 | 11) Weighted outcome models within each imputation + pooling | 88 |
| 3.4.1 | 11.1 Helper: fit + extract log-OR and SE from svyglm | 89 |
| 3.4.2 | 11.3 VBG: MI pooled spline models (treated cohort only) | 89 |
| 3.5 | 12) Explainability on one representative imputation | 89 |
| 3.6 | 13) MI three-level PCO2 helpers and checks | 89 |
| 3.7 | 14) MI + IPW three-level PCO2 (ABG & VBG) | 90 |
| 3.7.1 | 14.1 ABG: MI + IPW, three-level PCO2 outcomes | 90 |
| 3.7.2 | 14.2 VBG: MI + IPW, three-level PCO2 outcomes | 90 |
| 3.7.3 | 14.3 Table 3: MI-pooled IPW associations (3-level CO) | 91 |
| 3.7.4 | 14.4 Summary: adjusted CO2-category associations across analysis tracks | 91 |
| 3.8 | Manuscript outputs summary | 92 |
| 3.8.1 | 14.3 Visualization: pooled three-level ORs | 97 |
| 3.8.2 | 15.3 Visualization | 99 |
| 3.9 | Diagnostics | 102 |
| 3.9.1 | MI convergence and mixing | 102 |
| 3.9.2 | MI stability across m | 102 |
| 3.9.3 | MI maxit sensitivity (sampled) | 102 |
| 3.9.4 | Balance diagnostics | 102 |
| 3.9.5 | Outcome diagnostics | 104 |
| 3.9.6 | Diagnostics summary and audit | 104 |
| 3.9.7 | Performance / runtime log | 104 |
| 3.9.8 | Performance / runtime log | 106 |
| 3.10 | 16) Save, export, and session info | 107 |

1 Data Pre-processing

This code pulls in the master database (a STATA file) and does some initial cleaning - this will only need to be run once, and then the data can be accessed in the usual way.

1.0.1 Package Set Up

Cohort flow diagram will be generated externally and inserted here in the final manuscript packet.

1.1 Helper functions for model diagnostics

Converts the data from a STATA format to rdata if the rdata file does not exist. If it does already exist, it just loads that.

```
# data_dir_name resolved in setup-packages from params/env
stata_file <- file.path(data_dir_name, "full_db.dta")

stata_data <- read_dta(stata_file)
var_labels <- var_label(stata_data)
value_labels <- lapply(stata_data, function(x) if (is.labelled(x)) val_labels(x))
saveRDS(
  list(var_labels = var_labels, value_labels = value_labels),
  results_path("stata_labels.rds")
)
```

1.1.1 Configuration for the IPW models

```
drop_vars_ultra_missing <- c("bpn", "spo2")
cat_vars <- c("sex", "race_ethnicity", "location", "encounter_type")
numeric_vars <- c(
  "age_at_encounter", "curr_bmi", "temp_new", "sbp", "dbp", "hr",
  "sodium", "serum_cr", "serum_hco3", "serum_cl", "serum_lac", "serum_k",
  "wbc", "plt", "serum_phos", "serum_ca"
)
```

```

co2_vars <- c("paco2", "vbg_co2", "vbg_o2sat")

covars_gbm <- c(
  "age_at_encounter", "sex", "race_ethnicity", "curr_bmi",
  "copd", "asthma", "osa", "chf", "acute_nmd", "phtn", "ckd", "dm",
  "location", "encounter_type", "temp_new", "sbp", "dbp", "hr",
  "sodium", "serum_cr", "serum_hco3", "serum_cl", "serum_lac", "serum_k",
  "wbc", "plt", "serum_phos", "serum_ca"
)
covars_gbm <- setdiff(covars_gbm, drop_vars_ultra_missing)
covars_ps <- covars_gbm

# Core adjustment set for conditional prognostic models
adj_core <- c("age_at_encounter", "sex", "race_ethnicity", "location", "encounter_type")

gbm_params <- list(
  n.trees          = 800,
  interaction.depth = 3,
  shrinkage        = 0.01,
  bag.fraction     = 0.8,
  n.minobsinnode   = 10,
  cv.folds         = 0,
  stop.method      = "smd.max",
  n.cores          = 1L
)
stopifnot(gbm_params$stop.method == "smd.max")
SPLINE_BASIS <- "ns"
SPLINE_DF <- 4L
stopifnot(SPLINE_BASIS %in% c("ns", "rcs"))
get_gbm_cores <- function() {
  n_rows <- nrow(subset_data)
  if (is.finite(n_rows) && n_rows > 200000L) return(1L)
  gbm_params$n.cores
}
ps_trunc_quantile <- 0.01
stopifnot(ps_trunc_quantile > 0, ps_trunc_quantile < 0.5)

```

```

formula_abg      <- reformulate(covars_gbm, response = "has_abg")
formula_vbg      <- reformulate(covars_gbm, response = "has_vbg")

# Model diagrams: propensity models (GBM PS)
register_model_diagram("PS model: ABG test (GBM)", formula_abg, width = 10, height = 7)
register_model_diagram("PS model: VBG test (GBM)", formula_vbg, width = 10, height = 7)

# TODO: consider stop.method = "es.mean" (ATS version) if smd.max remains unstable on full N.

```

```

run_meta <- tibble::tibble(
  run_id      = diag_run_id,
  run_mode    = RUN_MODE,
  pilot_frac = PILOT_FRAC,
  mi_batch_threshold = MI_BATCH_THRESHOLD,
  full_n     = nrow(stata_data),
  subset_n   = nrow(subset_data)
)
render_table_pdf(run_meta,
                 "Run metadata (pilot vs full)",
                 "run_metadata",
                 preview_rows = 5,
                 digits = 0)

```

Table 1: Run metadata (pilot vs full)

| run_id | run_mode | pilot_frac | mi_batch_threshold | full_n | subset_n |
|-----------------|----------|------------|--------------------|--------|----------|
| 20260204_072236 | pilot | 0 | 5000 | 833476 | 25852 |

```

# Write run config JSON for portability
json_escape <- function(x) gsub("\\\"", "\\\\"\"", x)
run_cfg <- list(
  run_id = diag_run_id,
  run_mode = RUN_MODE,
  pilot_frac = PILOT_FRAC,

```

```

mi_batch_threshold = MI_BATCH_THRESHOLD,
data_dir = data_dir_name,
results_dir = results_dir,
full_n = nrow(stata_data),
subset_n = nrow(subset_data)
)
json_lines <- c(
  "{",
  paste0("  \"run_id\": \"", json_escape(run_cfg$run_id), "\","),
  paste0("  \"run_mode\": \"", json_escape(run_cfg$run_mode), "\","),
  paste0("  \"pilot_frac\": ", run_cfg$pilot_frac, ","),
  paste0("  \"mi_batch_threshold\": ", run_cfg$mi_batch_threshold, ","),
  paste0("  \"data_dir\": \"", json_escape(run_cfg$data_dir), "\","),
  paste0("  \"results_dir\": \"", json_escape(run_cfg$results_dir), "\","),
  paste0("  \"full_n\": ", run_cfg$full_n, ","),
  paste0("  \"subset_n\": ", run_cfg$subset_n),
  "}"
)
writeLines(json_lines, results_path("run_config.json"))

```

Codebook exported to Results/codebookr.docx.

1.1.2 Outcome Variable Creation

```

subset_data <- subset_data %>%
  mutate(
    ## 1. Did the patient die?
    died = if_else(!is.na(death_date), 1L, 0L),

    ## 2. Absolute death date (if death_date is an offset)
    death_abs = if_else(!is.na(death_date),
      encounter_date + death_date,
      as.Date(NA)),

```

```

## 3. Year month (YM) for encounter and death
enc_ym    = floor_date(encounter_date, unit = "month"),
death_ym  = floor_date(death_abs      , unit = "month"),

## 4. Reference censoring date: 1 Jun 2024
ref_ym = ymd("2024-06-01"),

## 5. Months from encounter to death or censoring
months_death_or_cens = case_when(
  !is.na(death_ym) ~ interval(enc_ym, death_ym) %/% months(1),
  TRUE           ~ interval(enc_ym, ref_ym)   %/% months(1)
),

## 6. Remove impossible values
months_death_or_cens = if_else(
  months_death_or_cens < 0 | months_death_or_cens > 16,
  NA_integer_, months_death_or_cens
),

## 7. Death within one or two months
died_1mo = if_else(died == 1 & months_death_or_cens < 1, 1L, 0L),
died_2mo = if_else(died == 1 & months_death_or_cens <= 1, 1L, 0L),

## 8. Month of death (missing if censored)
death_time = if_else(died == 1, months_death_or_cens, NA_integer_),

## 9. Death within 60 days (new variable)
death_60d = if_else(died == 1 & death_abs <= (encounter_date + days(60)), 1L, 0L)
) %>%
select(-enc_ym, -death_ym)

subset_data <- subset_data %>%
  mutate(
    death_60d = if_else(died == 1 & death_abs <= (encounter_date + days(60)), 1L, 0L)
)

```

1.2 2) Baseline tables

```
# Robust derivation of analysis variables + helper for Table 1 production
# -----
#
# helper: label binary 0/1 → "No"/"Yes"
bin_lab <- function(x) {
  y <- to01(x)
  if (all(is.na(y))) {
    return(factor(y, levels = c(0, 1), labels = c("No", "Yes")))
  }
  factor(y, levels = c(0, 1), labels = c("No", "Yes"))
}

# helper: preserve labeled factors if already present; otherwise map numeric codes
label_from_codes <- function(x, codes, labels) {
  if (is.factor(x)) {
    lev <- levels(x)
    if (all(lev %in% labels)) {
      return(factor(x, levels = labels))
    }
    lev_num <- suppressWarnings(as.numeric(lev))
    if (all(!is.na(lev_num)) && all(lev_num %in% codes)) {
      return(factor(as.numeric(as.character(x)), levels = codes, labels = labels))
    }
    return(x)
  }
  x_num <- suppressWarnings(as.numeric(as.character(x)))
  if (all(is.na(x_num))) return(factor(x, levels = labels))
  if (all(x_num %in% codes, na.rm = TRUE)) {
    return(factor(x_num, levels = codes, labels = labels))
  }
  factor(x)
}
```

```

subset_data <- subset_data %>%
  mutate(
    ## ensure 0/1 numerics (avoids factor-level coercion)
    across(c(has_abg, has_vbg),
           ~ to01(.)),

    ## derive ABG / VBG status groups (binary test status only)
    abg_group = case_when(
      has_abg == 0 ~ "No ABG",
      has_abg == 1 ~ "ABG",
      TRUE          ~ "Missing"
    ),
    vbg_group = case_when(
      has_vbg == 0 ~ "No VBG",
      has_vbg == 1 ~ "VBG",
      TRUE          ~ "Missing"
    ),

    ## factorise groups with explicit NA/Missing level
    abg_group = factor(
      abg_group,
      levels = c("No ABG", "ABG", "Missing")
    ),
    vbg_group = factor(
      vbg_group,
      levels = c("No VBG", "VBG", "Missing")
    ),

    ## labelled covariates (robust to factor or numeric codes)
    sex_label = label_from_codes(sex, c(0, 1), c("Female", "Male")),
    race_ethnicity_label = label_from_codes(
      race_ethnicity,
      0:6,
      c("White", "Black or African American", "Hispanic",
        "Asian", "American Indian", "Pacific Islander", "Unknown")
    ),
  )

```

```

location_label = label_from_codes(
  location,
  0:3,
  c("South", "Northeast", "Midwest", "West")
),
encounter_type_label = label_from_codes(
  encounter_type,
  c(2, 3),
  c("Emergency", "Inpatient")
),
osa_label      = bin_lab(osa),
asthma_label   = bin_lab(asthma),
copd_label     = bin_lab(copd),
chf_label      = bin_lab(chf),
nmd_label      = bin_lab(nmd),
phtn_label     = bin_lab(phtn),
ckd_label      = bin_lab(ckd),
diabetes_label = bin_lab(dm)
)

# variables to summarise
vars <- c(
  "age_at_encounter", "curr_bmi", "sex_label", "race_ethnicity_label", "location_label",
  "osa_label", "asthma_label", "copd_label", "chf_label", "nmd_label",
  "phtn_label", "ckd_label", "diabetes_label", "encounter_type_label", "vbg_co2", "paco2"
)
vars_baseline <- setdiff(vars, c("vbg_co2", "paco2"))
vars_abg <- c(vars_baseline, "paco2")
vars_vbg <- c(vars_baseline, "vbg_co2")

# Table 1 constructor
make_table1 <- function(data, group_var, caption = "") {
  group_sym <- rlang::sym(group_var)

  df <- data %>%
    filter(!is.na (!!group_sym), # drop explicit NA

```

```

  !!group_sym != "Missing") %>%          # drop "Missing" cohort
  mutate (!!group_sym := droplevels (!!group_sym)) %>%  # only drop group levels
  select(all_of(c(group_var, vars_baseline)))

empty_fac <- names(which(vapply(df, function(z) is.factor(z) && length(levels(z)) == 0L, logical(1))))
if (length(empty_fac) > 0) {
  warning("0-level factor columns detected: ", paste(empty_fac, collapse = ", "),
  ". Converting to character to prevent gtsummary failure.", call. = FALSE)
  df[empty_fac] <- lapply(df[empty_fac], as.character)
}

df %>%
  gtsummary::tbl_summary(
    by     = !!group_sym,
    type   = list(sex_label ~ "categorical"),
    statistic = list(
      gtsummary::all_continuous() ~ "{mean} ± {sd}; {N_miss}/{N_obs} missing ({p_miss}%)",
      gtsummary::all_categorical() ~ "{n} ({p}%)"
    ),
    digits  = list(gtsummary::all_continuous() ~ 1),
    missing  = "no"                                # no gtsummary missing column/row
  ) %>%
  gtsummary::modify_header(label = "***Variable***") %>%
  gtsummary::modify_caption(strip_manual_table_number(caption))
}

if (sum(!is.na(subset_data$sex_label)) == 0L || length(levels(subset_data$sex_label)) == 0L) {
  warning("sex_label is all NA or has zero levels; check sex normalization/mapping.", call. = FALSE)
  stopifnot("sex" %in% names(subset_data))
}

# build tables
table1A <- make_table1(subset_data, "abg_group", caption = "Table 1A: ABG cohorts")
table1B <- make_table1(subset_data, "vbg_group", caption = "Table 1B: VBG cohorts")

tbl1a_pdf <- to_pdf_table(table1A, font_size = 7, landscape = FALSE, label_col_width = "2.0in",

```

```

    longtable = TRUE)
tbl1b_pdf <- to_pdf_table(table1B, font_size = 7, landscape = FALSE, label_col_width = "2.0in",
                           longtable = TRUE)
tbl1a_pdf

```

Table 2: ABG cohorts

| **Variable** | **No ABG** N = 16,490 | **ABG** N = 9,362 |
|---------------------------|--|---|
| age_at_encounter | 58.2 ± 18.1; 0.0/16,490.0 missing (0.0%) | 61.1 ± 16.9; 0.0/9,362.0 missing (0.0%) |
| curr_bmi | 32.5 ± 8.7; 9,223.0/16,490.0 missing (55.9%) | 29.0 ± 7.1; 5,500.0/9,362.0 missing (58.7%) |
| sex_label | | |
| Female | 8,598 (52%) | 4,218 (45%) |
| Male | 7,892 (48%) | 5,144 (55%) |
| race_ethnicity_label | | |
| White | 10,061 (61%) | 6,090 (65%) |
| Black or African American | 3,149 (19%) | 1,387 (15%) |
| Hispanic | 1,168 (7.1%) | 500 (5.3%) |
| Asian | 252 (1.5%) | 179 (1.9%) |
| American Indian | 88 (0.5%) | 101 (1.1%) |
| Pacific Islander | 25 (0.2%) | 14 (0.1%) |
| Unknown | 1,747 (11%) | 1,091 (12%) |
| location_label | | |
| South | 6,962 (42%) | 5,171 (55%) |
| Northeast | 4,707 (29%) | 1,831 (20%) |
| Midwest | 1,143 (6.9%) | 783 (8.4%) |
| West | 3,678 (22%) | 1,577 (17%) |
| osa_label | 3,047 (18%) | 1,455 (16%) |
| asthma_label | 2,439 (15%) | 1,086 (12%) |
| copd_label | 3,037 (18%) | 2,052 (22%) |
| chf_label | 2,964 (18%) | 2,104 (22%) |
| nmd_label | 599 (3.6%) | 428 (4.6%) |
| phtn_label | 1,231 (7.5%) | 897 (9.6%) |
| ckd_label | 2,825 (17%) | 1,837 (20%) |
| diabetes_label | 4,838 (29%) | 2,776 (30%) |
| encounter_type_label | | |
| Emergency | 7,055 (43%) | 1,451 (15%) |
| Inpatient | 9,435 (57%) | 7,911 (85%) |

```
tbl1b_pdf
```

Table 3: VBG cohorts

| **Variable** | **No VBG** N = 18,392 | **VBG** N = 7,460 |
|---------------------------|--|---|
| age_at_encounter | 59.4 ± 17.7; 0.0/18,392.0 missing (0.0%) | 58.9 ± 17.7; 0.0/7,460.0 missing (0.0%) |
| curr_bmi | 31.9 ± 8.5; 9,691.0/18,392.0 missing (52.7%) | 29.0 ± 7.5; 5,032.0/7,460.0 missing (67.5%) |
| sex_label | | |
| Female | 9,331 (51%) | 3,485 (47%) |
| Male | 9,061 (49%) | 3,975 (53%) |
| race_ethnicity_label | | |
| White | 12,190 (66%) | 3,961 (53%) |
| Black or African American | 3,109 (17%) | 1,427 (19%) |
| Hispanic | 1,149 (6.2%) | 519 (7.0%) |
| Asian | 279 (1.5%) | 152 (2.0%) |
| American Indian | 91 (0.5%) | 98 (1.3%) |
| Pacific Islander | 33 (0.2%) | 6 (<0.1%) |
| Unknown | 1,541 (8.4%) | 1,297 (17%) |
| location_label | | |
| South | 9,882 (54%) | 2,251 (30%) |
| Northeast | 3,358 (18%) | 3,180 (43%) |
| Midwest | 1,241 (6.7%) | 685 (9.2%) |
| West | 3,911 (21%) | 1,344 (18%) |
| osa_label | 3,273 (18%) | 1,229 (16%) |
| asthma_label | 2,520 (14%) | 1,005 (13%) |
| copd_label | 3,596 (20%) | 1,493 (20%) |
| chf_label | 3,495 (19%) | 1,573 (21%) |
| nmd_label | 757 (4.1%) | 270 (3.6%) |
| phtn_label | 1,405 (7.6%) | 723 (9.7%) |
| ckd_label | 3,119 (17%) | 1,543 (21%) |
| diabetes_label | 5,137 (28%) | 2,477 (33%) |
| encounter_type_label | | |
| Emergency | 6,183 (34%) | 2,323 (31%) |
| Inpatient | 12,209 (66%) | 5,137 (69%) |

```
float_barrier()
```

```
# Status factors (column labels are taken from factor levels)
subset_data <- subset_data %>%
  mutate(
    abg_status = factor(has_abg, levels = c(0, 1),
                         labels = c("Did not get ABG", "Did get ABG")),
    vbg_status = factor(has_vbg, levels = c(0, 1),
                         labels = c("Did not get VBG", "Did get VBG"))
  )
```

```

# ABG table with "Everyone" column first
tbl1_abg <- subset_data %>%
  select(all_of(vars_baseline), abg_status) %>%
  gtsummary::tbl_summary(
    by = abg_status,
    type = list(sex_label ~ "categorical"),
    statistic = list(
      gtsummary::all_continuous() ~ "{mean} ± {sd}; {N_miss}/{N_obs} missing ({p_miss}%)",
      gtsummary::all_categorical() ~ "{n} ({p}%)"
    ),
    digits = list(gtsummary::all_continuous() ~ 1),
    missing = "no"
  ) %>%
  gtsummary::add_overall(last = FALSE, col_label = "Everyone") %>%
  gtsummary::modify_header(label = "***Variable***")

# VBG table (no "Everyone" here)
tbl1_vbg <- subset_data %>%
  select(all_of(vars_baseline), vbg_status) %>%
  gtsummary::tbl_summary(
    by = vbg_status,
    type = list(sex_label ~ "categorical"),
    statistic = list(
      gtsummary::all_continuous() ~ "{mean} ± {sd}; {N_miss}/{N_obs} missing ({p_miss}%)",
      gtsummary::all_categorical() ~ "{n} ({p}%)"
    ),
    digits = list(gtsummary::all_continuous() ~ 1),
    missing = "no"
  ) %>%
  gtsummary::modify_header(label = "***Variable***")

library(gtsummary)

tbl1 <- tbl_merge(
  tbls = list(tbl1_abg, tbl1_vbg)
) %>%

```

```

modify_caption(strip_manual_table_number("**Table 1. Baseline summary: Everyone, ABG status, and VBG status**"))

tbl1_pdf <- to_pdf_table(tbl1, font_size = 7, landscape = FALSE, label_col_width = "2.0in",
                           longtable = TRUE)
tbl1_pdf

```

Table 4: Baseline summary: Everyone, ABG status, and VBG status**

| **Variable** | Everyone | **Did not get ABG** N = 16,490 | **Did get ABG** N = 9,362 | **Did not get VBG** N = 12,203 |
|---------------------------|---|--|---|--|
| age_at_encounter | 59.2 ± 17.7; 0.0/25,852.0 missing (0.0%) | 58.2 ± 18.1; 0.0/16,490.0 missing (0.0%) | 61.1 ± 16.9; 0.0/9,362.0 missing (0.0%) | 59.4 ± 17.7; 0.0/12,203.0 missing (0.0%) |
| curr_bmi | 31.3 ± 8.4; 14,723.0/25,852.0 missing (57.0%) | 32.5 ± 8.7; 9,223.0/16,490.0 missing (55.9%) | 29.0 ± 7.1; 5,500.0/9,362.0 missing (58.7%) | 31.9 ± 8.5; 9,691.0/12,203.0 missing (57.0%) |
| sex_label | | | | |
| Female | 12,816 (50%) | 8,598 (52%) | 4,218 (45%) | 9,331 (76%) |
| Male | 13,036 (50%) | 7,892 (48%) | 5,144 (55%) | 3,062 (24%) |
| race_ethnicity_label | | | | |
| White | 16,151 (62%) | 10,061 (61%) | 6,090 (65%) | 12,188 (98%) |
| Black or African American | 4,536 (18%) | 3,149 (19%) | 1,387 (15%) | 3,101 (25%) |
| Hispanic | 1,668 (6.5%) | 1,168 (7.1%) | 500 (5.3%) | 1,143 (9%) |
| Asian | 431 (1.7%) | 252 (1.5%) | 179 (1.9%) | 279 (2%) |
| American Indian | 189 (0.7%) | 88 (0.5%) | 101 (1.1%) | 91 (0.7%) |
| Pacific Islander | 39 (0.2%) | 25 (0.2%) | 14 (0.1%) | 33 (0.2%) |
| Unknown | 2,838 (11%) | 1,747 (11%) | 1,091 (12%) | 1,546 (12%) |
| location_label | | | | |
| South | 12,133 (47%) | 6,962 (42%) | 5,171 (55%) | 9,886 (81%) |
| Northeast | 6,538 (25%) | 4,707 (29%) | 1,831 (20%) | 3,355 (27%) |
| Midwest | 1,926 (7.5%) | 1,143 (6.9%) | 783 (8.4%) | 1,244 (10%) |
| West | 5,255 (20%) | 3,678 (22%) | 1,577 (17%) | 3,912 (32%) |
| osa_label | 4,502 (17%) | 3,047 (18%) | 1,455 (16%) | 3,270 (27%) |
| asthma_label | 3,525 (14%) | 2,439 (15%) | 1,086 (12%) | 2,520 (21%) |
| copd_label | 5,089 (20%) | 3,037 (18%) | 2,052 (22%) | 3,590 (30%) |
| chf_label | 5,068 (20%) | 2,964 (18%) | 2,104 (22%) | 3,490 (30%) |
| nmd_label | 1,027 (4.0%) | 599 (3.6%) | 428 (4.6%) | 757 (6%) |
| phtn_label | 2,128 (8.2%) | 1,231 (7.5%) | 897 (9.6%) | 1,401 (11%) |
| ckd_label | 4,662 (18%) | 2,825 (17%) | 1,837 (20%) | 3,111 (25%) |
| diabetes_label | 7,614 (29%) | 4,838 (29%) | 2,776 (30%) | 5,131 (42%) |
| encounter_type_label | | | | |
| Emergency | 8,506 (33%) | 7,055 (43%) | 1,451 (15%) | 6,183 (51%) |
| Inpatient | 17,346 (67%) | 9,435 (57%) | 7,911 (85%) | 12,203 (49%) |

```
float_barrier()
```

```

# ABG cohort (has_abg == 1)
tbl2_abg <- subset_data %>%
  filter(has_abg == 1) %>%
  select(all_of(vars_abg), pco2_cat_abg) %>%
  gtsummary::tbl_summary(
    by = pco2_cat_abg,
    type = list(sex_label ~ "categorical"),
    statistic = list(
      gtsummary::all_continuous() ~ "{mean} ± {sd}; {N_miss}/{N_obs} missing ({p_miss}%)",
      gtsummary::all_categorical() ~ "{n} ({p}%)"
    ),
    digits = list(gtsummary::all_continuous() ~ 1),
    missing = "no"
  ) %>%
  gtsummary::modify_header(
    label = "***Variable***",
    stat_1 = "***Normal***",
    stat_2 = "***Low***",
    stat_3 = "***High***"
  )

# VBG cohort (has_vbg == 1)
tbl2_vbg <- subset_data %>%
  filter(has_vbg == 1) %>%
  select(all_of(vars_vbg), pco2_cat_vbg) %>%
  gtsummary::tbl_summary(
    by = pco2_cat_vbg,
    type = list(sex_label ~ "categorical"),
    statistic = list(
      gtsummary::all_continuous() ~ "{mean} ± {sd}; {N_miss}/{N_obs} missing ({p_miss}%)",
      gtsummary::all_categorical() ~ "{n} ({p}%)"
    ),
    digits = list(gtsummary::all_continuous() ~ 1),
    missing = "no"
  ) %>%
  gtsummary::modify_header()

```

```

label  = "##Variable**",
stat_1 = "##Normal**",
stat_2 = "##Low**",
stat_3 = "##High**"
)

tbl2 <- gtsummary::tbl_merge(
  tbls = list(tbl2_abg, tbl2_vbg),
  tab_spanner = c("##ABG (PaCO2)**", "##VBG (PvCO2)**")
) %>%
  gtsummary::modify_caption(strip_manual_table_number("##Table 2. Baseline summary by CO2 category within ABG and VBG cohorts##"))

tbl2_pdf <- to_pdf_table(tbl2, font_size = 7, landscape = FALSE, label_col_width = "2.0in",
                           longtable = TRUE)
tbl2_pdf

```

Table 5: Baseline summary by CO2 category within ABG and VBG cohorts**

| ##Variable## | ##Normal## | ##Low## | ##High## | ##Normal## |
|---------------------------|---|---|---|---|
| age_at_encounter | 60.9 ± 16.9; 0.0/4,229.0 missing (0.0%) | 60.4 ± 17.4; 0.0/2,488.0 missing (0.0%) | 62.0 ± 16.3; 0.0/2,645.0 missing (0.0%) | 58.0 ± 18.0; 0.0/3,427.0 missing (0.0%) |
| curr_bmi | 29.0 ± 6.9; 2,440.0/4,229.0 missing (57.7%) | 28.3 ± 6.8; 1,493.0/2,488.0 missing (60.0%) | 29.8 ± 7.6; 1,567.0/2,645.0 missing (59.2%) | 29.1 ± 7.3; 2,398.0/3,427.0 missing (59.2%) |
| sex_label | | | | |
| Female | 1,878 (44%) | 1,123 (45%) | 1,217 (46%) | 1,627 (47%) |
| Male | 2,351 (56%) | 1,365 (55%) | 1,428 (54%) | 1,800 (53%) |
| race_ethnicity_label | | | | |
| White | 2,742 (65%) | 1,510 (61%) | 1,838 (69%) | 1,825 (53%) |
| Black or African American | 609 (14%) | 400 (16%) | 378 (14%) | 648 (19%) |
| Hispanic | 229 (5.4%) | 144 (5.8%) | 127 (4.8%) | 243 (7.1%) |
| Asian | 82 (1.9%) | 59 (2.4%) | 38 (1.4%) | 71 (2.1%) |
| American Indian | 58 (1.4%) | 32 (1.3%) | 11 (0.4%) | 34 (1.0%) |
| Pacific Islander | 8 (0.2%) | 5 (0.2%) | 1 (<0.1%) | 3 (<0.1%) |
| Unknown | 501 (12%) | 338 (14%) | 252 (9.5%) | 603 (18%) |
| location_label | | | | |
| South | 2,334 (55%) | 1,399 (56%) | 1,438 (54%) | 1,106 (32%) |
| Northeast | 795 (19%) | 405 (16%) | 631 (24%) | 1,510 (44%) |
| Midwest | 373 (8.8%) | 184 (7.4%) | 226 (8.5%) | 312 (9.1%) |
| West | 727 (17%) | 500 (20%) | 350 (13%) | 499 (15%) |
| osa_label | 602 (14%) | 291 (12%) | 562 (21%) | 528 (15%) |
| asthma_label | 474 (11%) | 245 (9.8%) | 367 (14%) | 433 (13%) |
| copd_label | 779 (18%) | 374 (15%) | 899 (34%) | 548 (16%) |
| chf_label | 834 (20%) | 503 (20%) | 767 (29%) | 639 (19%) |
| nmd_label | 210 (5.0%) | 105 (4.2%) | 113 (4.3%) | 114 (3.3%) |
| phtn_label | 346 (8.2%) | 207 (8.3%) | 344 (13%) | 286 (8.3%) |
| ckd_label | 799 (19%) | 497 (20%) | 541 (20%) | 650 (19%) |

Table 5: Baseline summary by CO2 category within ABG and VBG cohorts** (*continu*

| **Variable** | **Normal** | **Low** | **High** | **Normal** |
|----------------------|--|--|---|--|
| diabetes_label | 1,211 (29%) | 750 (30%) | 815 (31%) | 1,070 (31%) |
| encounter_type_label | | | | |
| Emergency | 627 (15%) | 345 (14%) | 479 (18%) | 1,173 (34%) |
| Inpatient | 3,602 (85%) | 2,143 (86%) | 2,166 (82%) | 2,254 (66%) |
| paco2 | 39.7 ± 3.0; 0.0/4,229.0 missing (0.0%) | 29.4 ± 4.4; 0.0/2,488.0 missing (0.0%) | 59.6 ± 20.6; 0.0/2,645.0 missing (0.0%) | |
| vbg_co2 | | | | 44.6 ± 3.0; 0.0/3,427.0 missing (0.0%) |

```
float_barrier()
```

Table 6: Table 2a. Crude outcomes by CO2 category

| Cohort | Outcome | Normal | Low | High |
|--------|----------------|------------------|------------------|------------------|
| ABG | IMV | 979/4229 (23.1%) | 691/2488 (27.8%) | 721/2645 (27.3%) |
| ABG | NIV | 288/4229 (6.8%) | 187/2488 (7.5%) | 421/2645 (15.9%) |
| ABG | Death (60d) | 604/4229 (14.3%) | 537/2488 (21.6%) | 504/2645 (19.1%) |
| ABG | Hypercapnic RF | 207/4229 (4.9%) | 119/2488 (4.8%) | 666/2645 (25.2%) |
| VBG | IMV | 432/3427 (12.6%) | 357/2058 (17.3%) | 373/1975 (18.9%) |
| VBG | NIV | 175/3427 (5.1%) | 122/2058 (5.9%) | 257/1975 (13.0%) |
| VBG | Death (60d) | 378/3427 (11.0%) | 356/2058 (17.3%) | 310/1975 (15.7%) |
| VBG | Hypercapnic RF | 146/3427 (4.3%) | 70/2058 (3.4%) | 476/1975 (24.1%) |

1.3 3) Three-level PCO2 categories (unweighted)

Three groups using low/normal/high CO2 categories

```
stopifnot(all(c("pco2_cat_abg", "pco2_cat_vbg") %in% names(subset_data)))

library(broom)
library(tidyr)
library(dplyr)

run_logit <- function(data, outcome, exposure, group_name, adj_vars = NULL, model_type = "Crude") {
  f <- if (length(adj_vars)) {
    reformulate(c(exposure, adj_vars), response = outcome)
```

```

} else {
  as.formula(paste(outcome, "~", exposure))
}
fit_res <- fit_with_diagnostics(
  function() glm(f, data = data, family = binomial, control = glm.control(maxit = 50)),
  context = make_context(
    stage = "outcome",
    component = "cat3",
    analysis_variant = "unweighted",
    model_type = "cat3",
    group = group_name,
    outcome = outcome,
    imputation = NA_integer_,
    batch = NA_integer_
  )
)
append_outcome_diag(fit_res$diag)
if (is.null(fit_res$fit)) {
  stop("run_logit: model fit failed for outcome=", outcome,
       " exposure=", exposure, " group=", group_name)
}
tidy(fit_res$fit, exponentiate = TRUE, conf.int = TRUE) %>%
  filter(term != "(Intercept)", startsWith(term, exposure)) %>%
  mutate(
    outcome = outcome,
    group = group_name,
    model = model_type
  )
}

outcomes_unw <- c("imv_proc", "niv_proc", "death_60d", "hypercap_resp_failure")

unw_three_level_forms <- list(
  "ABG 3-level: IMV ~ CO2 category + X"      = reformulate(c("pco2_cat_abg", adj_core), response = "imv_proc"),
  "ABG 3-level: NIV ~ CO2 category + X"      = reformulate(c("pco2_cat_abg", adj_core), response = "niv_proc"),
  "ABG 3-level: Death60d ~ CO2 category + X" = reformulate(c("pco2_cat_abg", adj_core), response = "death_60d"),
)

```

```

"ABG 3-level: HCRF ~ CO2 category + X"      = reformulate(c("pc02_cat_abg", adj_core), response = "hypercap_resp_failure"),
"VBG 3-level: IMV ~ CO2 category + X"        = reformulate(c("pc02_cat_vbg", adj_core), response = "imv_proc"),
"VBG 3-level: NIV ~ CO2 category + X"        = reformulate(c("pc02_cat_vbg", adj_core), response = "niv_proc"),
"VBG 3-level: Death60d ~ CO2 category + X"   = reformulate(c("pc02_cat_vbg", adj_core), response = "death_60d"),
"VBG 3-level: HCRF ~ CO2 category + X"        = reformulate(c("pc02_cat_vbg", adj_core), response = "hypercap_resp_failure")
)
register_model_diagrams(unw_three_level_forms)

unw_results_crude <- bind_rows(
  lapply(outcomes_unw, function(o) run_logit(subset_data, o, "pc02_cat_abg", "ABG")),
  lapply(outcomes_unw, function(o) run_logit(subset_data, o, "pc02_cat_vbg", "VBG"))
)
unw_results_adj <- bind_rows(
  lapply(outcomes_unw, function(o) run_logit(subset_data, o, "pc02_cat_abg", "ABG", adj_core, "Adjusted")),
  lapply(outcomes_unw, function(o) run_logit(subset_data, o, "pc02_cat_vbg", "VBG", adj_core, "Adjusted"))
)

unw_threellevel_results <- unw_results_adj %>%
  mutate(method = "Unweighted adjusted")

unw_combined_or_df <- unw_results_adj %>%
  mutate(
    outcome = recode(outcome,
                     imv_proc = "Intubation",
                     niv_proc = "NIV",
                     death_60d = "Death (60d)",
                     hypercap_resp_failure = "Hypercapnic RF")
  )
unw_combined_or_df <- map_or_exposure(unw_combined_or_df, "or-plot-three-level-unweighted") |>
  select(outcome, group, exposure, estimate, conf.low, conf.high)

library(scales)

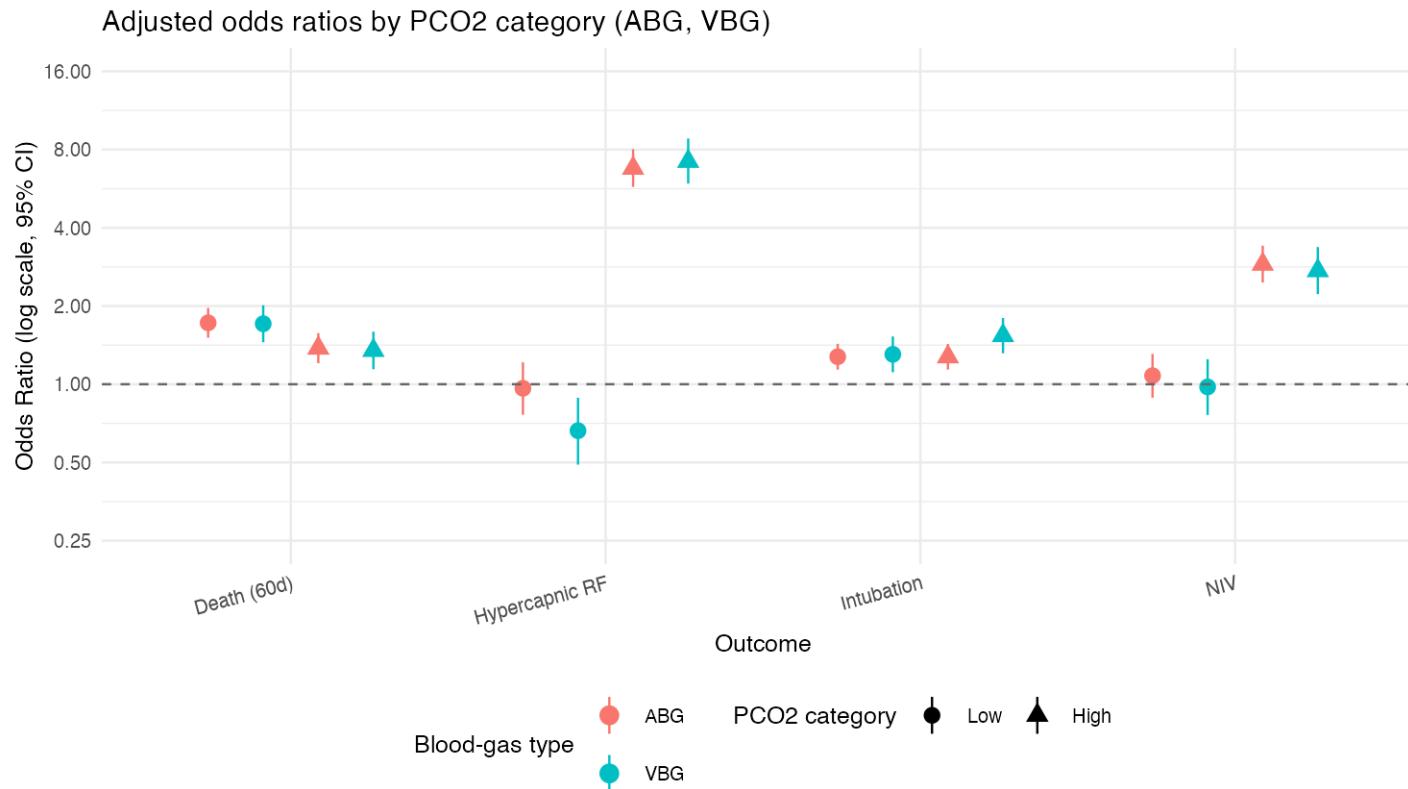
unw_combined_or_df$group <- factor(
  unw_combined_or_df$group,
  levels = c("ABG", "VBG")

```

```

)
unw_plot_df <- build_or_plot_df(unw_combined_or_df, "or-plot-three-level-unweighted",
                                 expected_exposure_levels = CO2_CAT_CONTRAST_LEVELS)
stopifnot(is.data.frame(unw_plot_df))
unw_axis_spec <- compute_or_axis_spec(unw_plot_df, lo_col = "conf.low", hi_col = "conf.high",
                                       default_limits = OR_XLIM)
unw_p_or <- plot_or_safe(
  unw_plot_df,
  plot_name = "or-plot-three-level-unweighted",
  axis_spec = unw_axis_spec,
  title = "Adjusted odds ratios by PCO2 category (ABG, VBG)",
  caption = paste(
    "Adjusted for age, sex, race/ethnicity, location, and encounter type.",
    "Reference = patients in the normal PCO2 range.",
    "Low: <35 mmHg (ABG) or <40 mmHg (VBG); High: >45 mmHg (ABG) or >50 mmHg (VBG).",
    "Because the underlying cohorts differ (ABG, VBG), denominators are not identical across groups.",
    sep = "\n"
  )
)
print_plot_once(unw_p_or, "or-plot-three-level-unweighted", width = 7.5, height = 4.8)

```



1.4 4) Restricted cubic spline regressions (unweighted)

Spline curves are shown as odds ratios relative to CO₂_ref (midpoint of the normal range), holding covariates at the reference profile.

```
# ABG spline dataset
subset_data_abg <- subset_data %>%
  filter(has_abg == 1, !is.na(paco2)) %>%
  select(paco2, imv_proc, niv_proc, death_60d, hypercap_resp_failure, all_of(adj_core)) %>%
  filter(complete.cases(.))
```

1.4.1 4.1 Unweighted, Restricted Cubic Spline Regression - ABG by PaCO2

```
make_spline_fml <- function(outcome, co2_var, adj_vars) {
  spline_term <- if (SPLINE_BASIS == "rcs") {
    sprintf("rms::rcs(%s, %d)", co2_var, SPLINE_DF)
  } else {
    sprintf("splines::ns(%s, %d)", co2_var, SPLINE_DF)
  }
  stats::as.formula(
    paste0(outcome, " ~ ", spline_term,
           if (length(adj_vars)) paste0(" + ", paste(adj_vars, collapse = " + ")) else ""))
}

#| code-block-title: "Unweighted ABG spline models (adjusted)"
abg_spline_forms <- list(
  "ABG spline (adjusted): IMV ~ CO2 spline + X"      = make_spline_fml("imv_proc", "paco2", adj_core),
  "ABG spline (adjusted): NIV ~ CO2 spline + X"       = make_spline_fml("niv_proc", "paco2", adj_core),
  "ABG spline (adjusted): Death60d ~ CO2 spline + X"  = make_spline_fml("death_60d", "paco2", adj_core),
  "ABG spline (adjusted): HCRF ~ CO2 spline + X"      = make_spline_fml("hypercap_resp_failure", "paco2", adj_core)
)
register_model_diagrams(abg_spline_forms)

co2_seq_abg <- stats::quantile(subset_data_abg$paco2, probs = c(0.02, 0.98), na.rm = TRUE)
grid_abg_info_unw <- make_co2_grid_ref(
  "paco2",
  seq(co2_seq_abg[1], co2_seq_abg[2], length.out = SPLINE_GRID_N),
  x_ref_abg,
  ABG_CO2_REF
)
grid_abg_unw <- grid_abg_info_unw$grid
ref_idx_abg_unw <- grid_abg_info_unw$ref_idx
co2_ref_abg_unw <- grid_abg_info_unw$co2_ref

fit_spline_glm <- function(outcome, co2_var, data, group_label) {
```

```

fit_res <- fit_with_diagnostics(
  function() glm(make_spline_fml(outcome, co2_var, adj_core),
    data = data, family = binomial,
    control = glm.control(maxit = 50)),
  context = make_context(
    stage = "outcome",
    component = "spline",
    analysis_variant = "unweighted",
    model_type = "spline",
    group = group_label,
    outcome = outcome,
    imputation = NA_integer_,
    batch = NA_integer_
  )
)
append_outcome_diag(fit_res$diag)
fit_res$fit
}

fit_imv <- fit_spline_glm("imv_proc", "paco2", subset_data_abg, "ABG")
fit_niv <- fit_spline_glm("niv_proc", "paco2", subset_data_abg, "ABG")
fit_death <- fit_spline_glm("death_60d", "paco2", subset_data_abg, "ABG")
fit_hcrcf <- fit_spline_glm("hypercap_resp_failure", "paco2", subset_data_abg, "ABG")
if (any(vapply(list(fit_imv, fit_niv, fit_death, fit_hcrcf), is.null, logical(1)))) {
  stop("Unweighted ABG spline fits failed; see model_fit_diagnostics.csv.")
}

pred_imv <- predict_or_curve_from_fit(fit_imv, grid_abg_unw, ref_idx_abg_unw, "paco2")
pred_niv <- predict_or_curve_from_fit(fit_niv, grid_abg_unw, ref_idx_abg_unw, "paco2")
pred_death <- predict_or_curve_from_fit(fit_death, grid_abg_unw, ref_idx_abg_unw, "paco2")
pred_hcrcf <- predict_or_curve_from_fit(fit_hcrcf, grid_abg_unw, ref_idx_abg_unw, "paco2")
## Plotting deferred until VBG curves are computed so axes can be shared.

```

1.4.2 4.2 Unweighted, Restricted Cubic Spline - VBG

```

# --- VBG dataset ---
subset_data_vbg <- subset_data %>%
  dplyr::filter(has_vbg == 1, !is.na(vbg_co2)) %>%
  dplyr::select(vbg_co2, imv_proc, niv_proc, death_60d, hypercap_resp_failure, all_of(adj_core)) %>%
  dplyr::filter(complete.cases(.))

vbg_spline_forms <- list(
  "VBG spline (adjusted): IMV ~ CO2 spline + X"      = make_spline_fml("imv_proc", "vbg_co2", adj_core),
  "VBG spline (adjusted): NIV ~ CO2 spline + X"      = make_spline_fml("niv_proc", "vbg_co2", adj_core),
  "VBG spline (adjusted): Death60d ~ CO2 spline + X" = make_spline_fml("death_60d", "vbg_co2", adj_core),
  "VBG spline (adjusted): HCRF ~ CO2 spline + X"      = make_spline_fml("hypercap_resp_failure", "vbg_co2", adj_core)
)
register_model_diagrams(vbg_spline_forms)

co2_seq_vbg <- stats::quantile(subset_data_vbg$vbg_co2, probs = c(0.02, 0.98), na.rm = TRUE)
grid_vbg_info_unw <- make_co2_grid_ref(
  "vbg_co2",
  seq(co2_seq_vbg[1], co2_seq_vbg[2], length.out = SPLINE_GRID_N),
  x_ref_vbg,
  VBG_CO2_REF
)
grid_vbg_unw <- grid_vbg_info_unw$grid
ref_idx_vbg_unw <- grid_vbg_info_unw$ref_idx
co2_ref_vbg_unw <- grid_vbg_info_unw$co2_ref

fit_imv_vbg <- fit_spline_glm("imv_proc", "vbg_co2", subset_data_vbg, "VBG")
fit_niv_vbg <- fit_spline_glm("niv_proc", "vbg_co2", subset_data_vbg, "VBG")
fit_death_vbg <- fit_spline_glm("death_60d", "vbg_co2", subset_data_vbg, "VBG")
fit_hcrf_vbg <- fit_spline_glm("hypercap_resp_failure", "vbg_co2", subset_data_vbg, "VBG")
if (any(vapply(list(fit_imv_vbg, fit_niv_vbg, fit_death_vbg, fit_hcrf_vbg), is.null, logical(1)))) {
  stop("Unweighted VBG spline fits failed; see model_fit_diagnostics.csv.")
}

```

```

pred_imv_vbg <- predict_or_curve_from_fit(fit_imv_vbg, grid_vbg_unw, ref_idx_vbg_unw, "vbg_co2")
pred_niv_vbg <- predict_or_curve_from_fit(fit_niv_vbg, grid_vbg_unw, ref_idx_vbg_unw, "vbg_co2")
pred_death_vbg <- predict_or_curve_from_fit(fit_death_vbg, grid_vbg_unw, ref_idx_vbg_unw, "vbg_co2")
pred_hcrf_vbg <- predict_or_curve_from_fit(fit_hcrf_vbg, grid_vbg_unw, ref_idx_vbg_unw, "vbg_co2")
axis_unw_common <- compute_or_axis_spec(
  list(pred_imv, pred_niv, pred_death, pred_hcrf,
    pred_imv_vbg, pred_niv_vbg, pred_death_vbg, pred_hcrf_vbg),
  lo_col = "LCL", hi_col = "UCL"
)

plot_imv <- ggplot(pred_imv, aes(x = paco2, y = OR)) +
  geom_line(color = "blue", linewidth = 1.2) +
  geom_ribbon(aes(ymin = LCL, ymax = UCL), fill = "blue", alpha = 0.2) +
  geom_hline(yintercept = 1, linetype = "dashed", color = "grey40") +
  or_axis_scale(axis_unw_common) +
  labs(title = "Intubation (adjusted)", x = "PaCO2 (mmHg)",
    y = paste0("Odds ratio (ref PaCO2 = ", co2_ref_abg_unw, "; log scale)")) +
  theme_minimal()

plot_niv <- ggplot(pred_niv, aes(x = paco2, y = OR)) +
  geom_line(color = "green", linewidth = 1.2) +
  geom_ribbon(aes(ymin = LCL, ymax = UCL), fill = "green", alpha = 0.2) +
  geom_hline(yintercept = 1, linetype = "dashed", color = "grey40") +
  or_axis_scale(axis_unw_common) +
  labs(title = "NIV (adjusted)", x = "PaCO2 (mmHg)",
    y = paste0("Odds ratio (ref PaCO2 = ", co2_ref_abg_unw, "; log scale)")) +
  theme_minimal()

plot_death <- ggplot(pred_death, aes(x = paco2, y = OR)) +
  geom_line(color = "red", linewidth = 1.2) +
  geom_ribbon(aes(ymin = LCL, ymax = UCL), fill = "red", alpha = 0.2) +
  geom_hline(yintercept = 1, linetype = "dashed", color = "grey40") +
  or_axis_scale(axis_unw_common) +
  labs(title = "Death (60d, adjusted)", x = "PaCO2 (mmHg)",
    y = paste0("Odds ratio (ref PaCO2 = ", co2_ref_abg_unw, "; log scale)")) +
  theme_minimal()

```

```

plot_hcrcf <- ggplot(pred_hcrcf, aes(x = paco2, y = OR)) +
  geom_line(color = "purple", linewidth = 1.2) +
  geom_ribbon(aes(ymin = LCL, ymax = UCL), fill = "purple", alpha = 0.2) +
  geom_hline(yintercept = 1, linetype = "dashed", color = "grey40") +
  or_axis_scale(axis_unw_common) +
  labs(title = "Hypercapnic RF (adjusted)", x = "PaCO2 (mmHg)",
       y = paste0("Odds ratio (ref PaCO2 = ", co2_ref_abg_unw, "; log scale)")) +
  theme_minimal()

plot_imv_vbg <- ggplot(pred_imv_vbg, aes(x = vbg_co2, y = OR)) +
  geom_line(color = "blue") +
  geom_ribbon(aes(ymin = LCL, ymax = UCL), fill = "blue", alpha = 0.2) +
  geom_hline(yintercept = 1, linetype = "dashed", color = "grey40") +
  or_axis_scale(axis_unw_common) +
  labs(title = "IMV (adjusted)", x = "VBG CO2 (mmHg)",
       y = paste0("Odds ratio (ref VBG CO2 = ", co2_ref_vbg_unw, "; log scale)")) +
  theme_minimal()

plot_niv_vbg <- ggplot(pred_niv_vbg, aes(x = vbg_co2, y = OR)) +
  geom_line(color = "green") +
  geom_ribbon(aes(ymin = LCL, ymax = UCL), fill = "green", alpha = 0.2) +
  geom_hline(yintercept = 1, linetype = "dashed", color = "grey40") +
  or_axis_scale(axis_unw_common) +
  labs(title = "NIV (adjusted)", x = "VBG CO2 (mmHg)",
       y = paste0("Odds ratio (ref VBG CO2 = ", co2_ref_vbg_unw, "; log scale)")) +
  theme_minimal()

plot_death_vbg <- ggplot(pred_death_vbg, aes(x = vbg_co2, y = OR)) +
  geom_line(color = "red") +
  geom_ribbon(aes(ymin = LCL, ymax = UCL), fill = "red", alpha = 0.2) +
  geom_hline(yintercept = 1, linetype = "dashed", color = "grey40") +
  or_axis_scale(axis_unw_common) +
  labs(title = "Death (60d, adjusted)", x = "VBG CO2 (mmHg)",
       y = paste0("Odds ratio (ref VBG CO2 = ", co2_ref_vbg_unw, "; log scale)")) +
  theme_minimal()

```

```

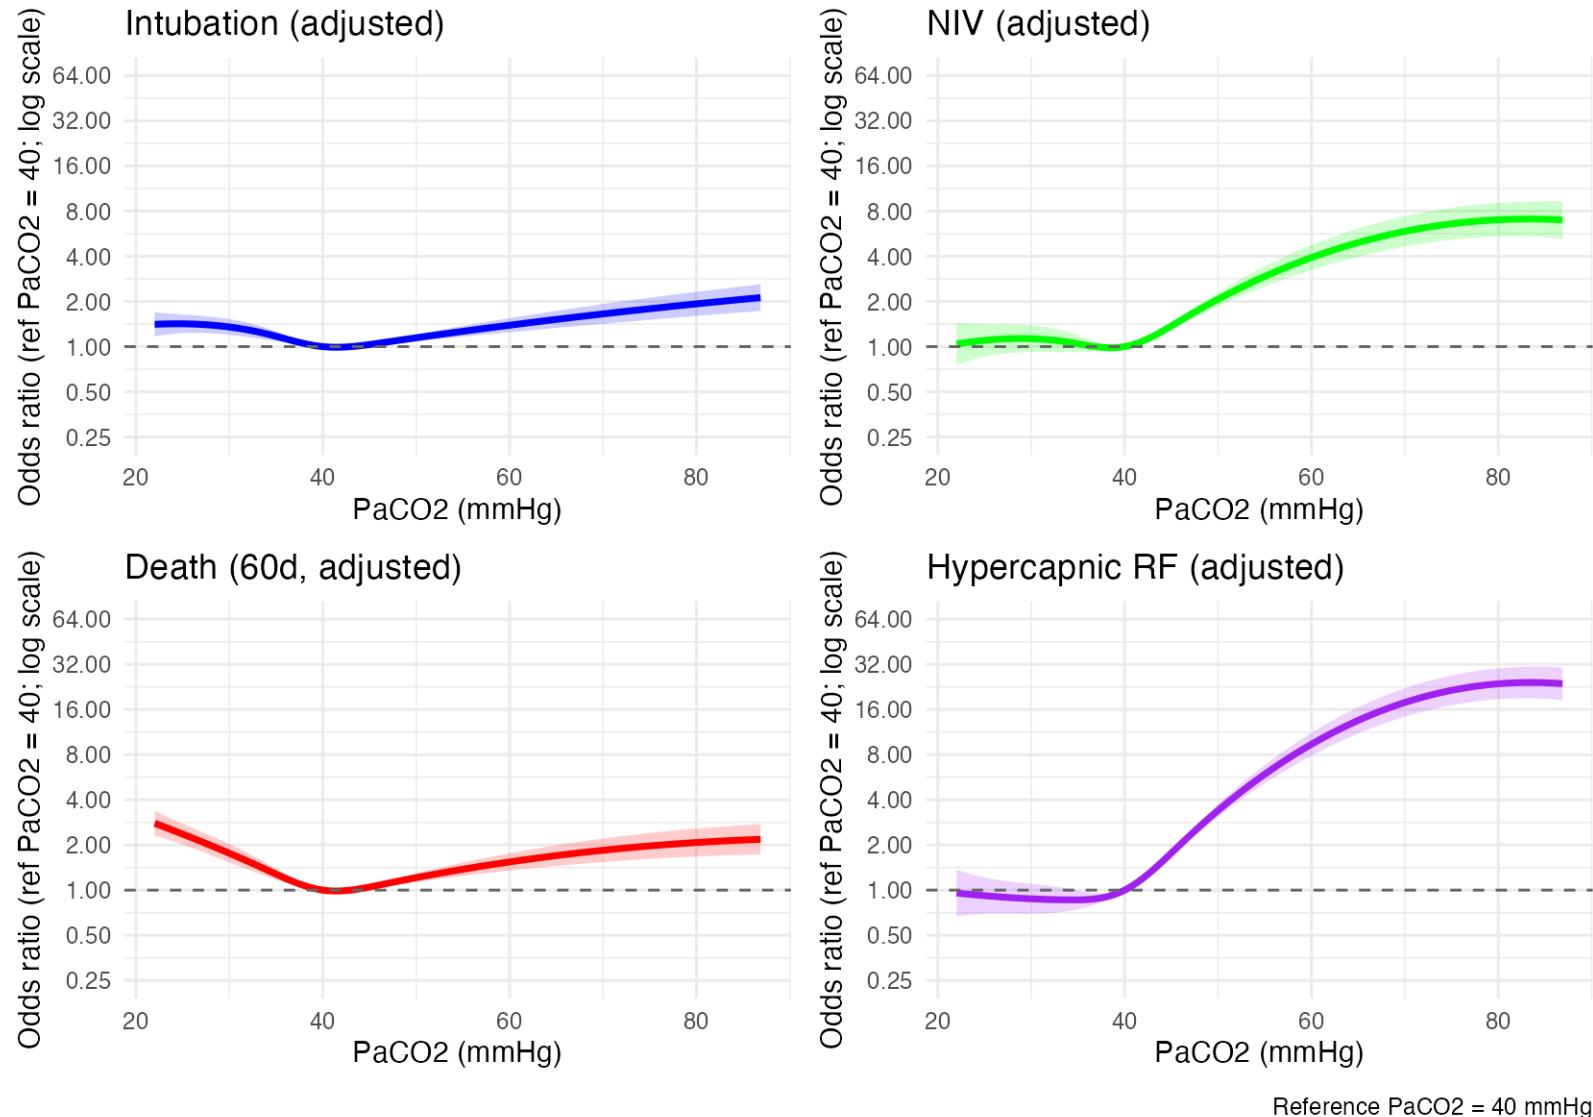
plot_hcrf_vbg <- ggplot(pred_hcrf_vbg, aes(x = vbg_co2, y = OR)) +
  geom_line(color = "purple") +
  geom_ribbon(aes(ymin = LCL, ymax = UCL), fill = "purple", alpha = 0.2) +
  geom_hline(yintercept = 1, linetype = "dashed", color = "grey40") +
  or_axis_scale(axis_unw_common) +
  labs(title = "Hypercapnic RF (adjusted)", x = "VBG CO2 (mmHg)",
       y = paste0("Odds ratio (ref VBG CO2 = ", co2_ref_vbg_unw, "; log scale)")) +
  theme_minimal()

unw_abg_panel <- (plot_imv | plot_niv) / (plot_death | plot_hcrf) +
  plot_annotation(caption = paste0("Reference PaCO2 = ", co2_ref_abg_unw, " mmHg"))

unw_vbg_panel <- ((plot_imv_vbg | plot_niv_vbg) /
  (plot_death_vbg | plot_hcrf_vbg)) +
  plot_annotation(
    title = paste0("Adjusted odds ratios by VBG CO2 (ref = ", co2_ref_vbg_unw, ")"),
    caption = paste0("Reference VBG CO2 = ", co2_ref_vbg_unw, " mmHg")
  )

print_plot_once(unw_abg_panel, "spline-unweighted-abg", width = 8.5, height = 6)

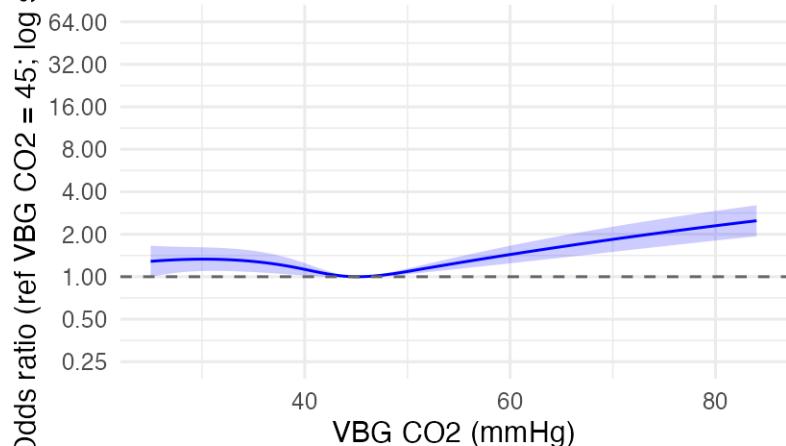
```



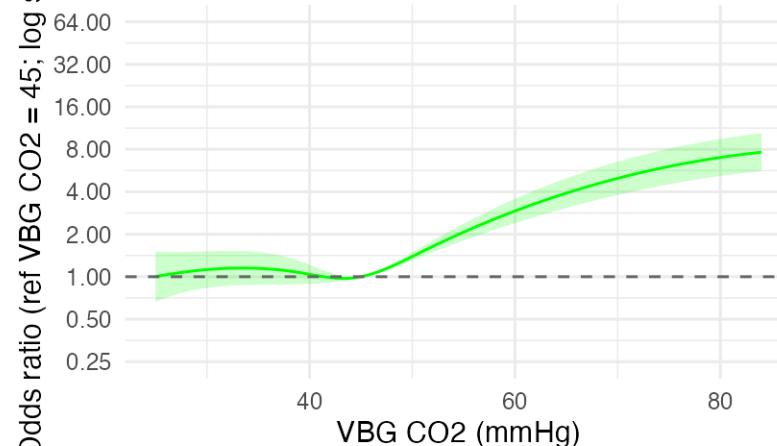
```
print_plot_once(unw_vbg_panel, "spline-unweighted-vbg", width = 8.5, height = 6)
```

Adjusted odds ratios by VBG CO₂ (ref = 45)

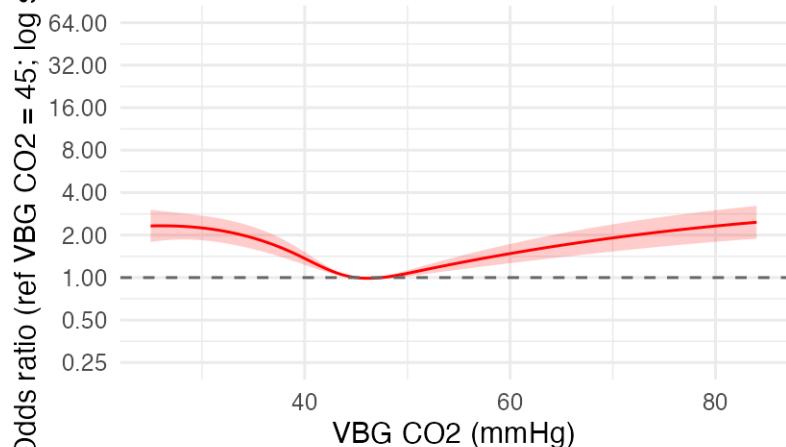
IMV (adjusted)



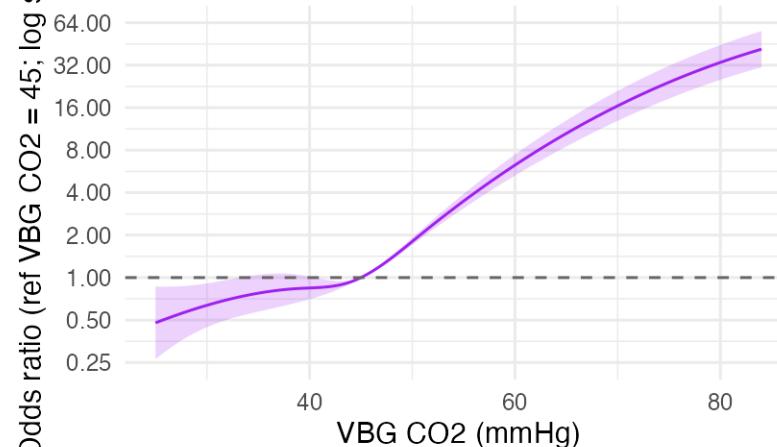
NIV (adjusted)



Death (60d, adjusted)



Hypercapnic RF (adjusted)



Reference VBG CO₂ = 45 mmHg

```
float_barrier()
```

2 Inverse Propensity Weighting

IPW done using Gradient Boosting Methods (GBM) - a type of decision-tree based machine learning. “***Random forests and GBM are designed to automatically include relevant interactions for variables included in the model.*** As such, using a GBM to estimate the PS model, can reduce model misspecification, since ***the analyst is not required to identify relevant interactions or nonlinearities.***” from this citation: PMID: 39947224<https://pmc.ncbi.nlm.nih.gov/articles/PMC11825193/>

Current propensity score uses `covars_gbm` (demographics, comorbidities, encounter type, vitals, labs) as defined above; in this block only `encounter_type` is explicitly factored before weighting.

Note: for all these, I suggested new GBM adjustments that accomplish the following:

1. Smaller GBM & balance-based stopping (`stop.method = "smd.max"`) → faster fit, avoids over-fitting, lighter tails (which lead to extreme weights that are problematic).
2. Target balance compares weighted treated cohort to the full sample; aim for $|SMD| < 0.1$.
3. Weight stabilization (divide by mean) mitigates a few huge weights. We use one-sided truncation at very small propensities (caps large weights only).
4. Uses robust variance estimation (e.g. allows the variances to change by PaCO₂) for IP-weighted GLM; works with splines via `rcs()`. This is a bit nuanced but I think good to change even though it adds complexity
5. Deterministic seed ensures result replication.

2.0.1 5.1 ABG IPW weighting and diagnostics

```
# Already normalized globally; just drop unused levels
subset_data$encounter_type <- droplevels(subset_data$encounter_type)
```

GBM tuning is shared across ABG and VBG via `gbm_params` to keep symmetry; update there if needed.

```
#   1. fit GBM propensity model, ABG
set.seed(42)
gbm_df_abg <- subset_data[, c("has_abg", "has_vbg", covars_gbm), drop = FALSE]
gbm_df_abg <- normalize_types(gbm_df_abg, levels_ref)
gbm_df_abg <- droplevels_all(gbm_df_abg)
```

```

gbm_preflight(gbm_df_abg, covars_gbm, "unimp_abg")
append_mem_snapshot("gbm_unimp", "unimp_abg", "pre")

weight_model <- do.call(
  weightit,
  c(
    list(
      formula_abg,
      data      = gbm_df_abg,
      method    = "gbm",
      estimand  = "ATE",
      missing   = "ind",
      include.obj = FALSE
    ),
    gbm_params
  )
)
append_mem_snapshot("gbm_unimp", "unimp_abg", "post")

# 2. One-sided IPSW (ABG observed only) + truncation of small propensities
ipow_abg <- compute_ipow_weights(
  weight_model,
  treat = gbm_df_abg$has_abg,
  ps_floor_quantile = ps_trunc_quantile,
  stabilize = TRUE
)
w_abg <- ipow_abg$weights
ps_floor_abg <- ipow_abg$ps_floor
subset_data$trunc_abg <- ipow_abg$truncated
subset_data$ps_abg <- ipow_abg$ps
subset_data$w_abg <- w_abg
assert_finite_weights(w_abg[subset_data$has_abg == 1], "w_abg")
rm(weight_model, gbm_df_abg)
invisible(gc())

# Balance diagnostics and treated-only outcome models are handled later.

```

Inverse Propensity-Weighted Logistic Regressions with CO2 predictor represented as a restricted cubic spline.

These are covariate-adjusted outcome models ($\text{outcome} \sim \text{spline}(\text{CO2}) + \text{X}$), fit separately for ABG and VBG cohorts using `survey::svyglm` with robust (design-based) SEs. Spline curves are shown as odds ratios relative to `CO2_ref` (midpoint of the normal range).

2.0.2 5.2 ABG IPW spline models

```
# set.seed(42) # reproducible GBM fit
#
# # 1. inverse-probability weights for receiving an ABG
#
# # done in the last block, so not needed
#
#
# Model diagrams: IPW ABG spline models
ipw_abg_rcs_forms <- list(
  "ABG IPW spline (adjusted): IMV ~ CO2 spline + X"      = make_spline_fml("imv_proc", "paco2", adj_core),
  "ABG IPW spline (adjusted): NIV ~ CO2 spline + X"      = make_spline_fml("niv_proc", "paco2", adj_core),
  "ABG IPW spline (adjusted): Death60d ~ CO2 spline + X" = make_spline_fml("death_60d", "paco2", adj_core),
  "ABG IPW spline (adjusted): HCRF ~ CO2 spline + X"     = make_spline_fml("hypercap_resp_failure", "paco2", adj_core)
)
register_model_diagrams(ipw_abg_rcs_forms)

# 2. analysis sample: rows with a measured PaCO2
subset_data_abg <- subset_data %>%
  filter(!is.na(paco2)) %>% # implies has_abg == 1
  select(paco2, imv_proc, niv_proc, death_60d,
         hypercap_resp_failure, w_abg, all_of(adj_core)) %>%
  filter(complete.cases(.))

#
# 3. weighted logistic spline models with robust SEs
fitfun <- function(formula, outcome) {
  fit_res <- fit_with_diagnostics(
    function() svyglm(
```

```

formula,
design = svydesign(ids = ~1, weights = ~w_abg, data = subset_data_abg),
family = quasibinomial(),
control = glm.control(maxit = 50)
),
context = make_context(
  stage = "outcome",
  component = "spline",
  analysis_variant = "ipw",
  model_type = "spline",
  group = "ABG",
  outcome = outcome,
  imputation = NA_integer_,
  batch = NA_integer_
)
)
append_outcome_diag(fit_res$diag)
fit_res$fit
}

fit_imv_abg  <- fitfun(make_spline_fml("imv_proc", "paco2", adj_core), "imv_proc")
fit_niv_abg  <- fitfun(make_spline_fml("niv_proc", "paco2", adj_core), "niv_proc")
fit_death_abg <- fitfun(make_spline_fml("death_60d", "paco2", adj_core), "death_60d")
fit_hcrf_abg <- fitfun(make_spline_fml("hypercap_resp_failure", "paco2", adj_core),
                         "hypercap_resp_failure")
if (any(vapply(list(fit_imv_abg, fit_niv_abg, fit_death_abg, fit_hcrf_abg), is.null, logical(1)))) {
  stop("IPW ABG spline fits failed; see model_fit_diagnostics.csv.")
}

# 4. prediction helper
mkpred <- function(fit, data_ref, co2_var, ref_df, co2_ref) {
  co2_seq <- stats::quantile(data_ref[[co2_var]], probs = c(0.02, 0.98), na.rm = TRUE)
  grid_info <- make_co2_grid_ref(
    co2_var,
    seq(co2_seq[1], co2_seq[2], length.out = SPLINE_GRID_N),
    ref_df,

```

```

    co2_ref
  )
  predict_or_curve_from_fit(fit, grid_info$grid, grid_info$ref_idx, co2_var)
}

pred_imv_abg  <- mkpred(fit_imv_abg, subset_data_abg, "paco2", x_ref_abg, ABG_CO2_REF)
pred_niv_abg  <- mkpred(fit_niv_abg, subset_data_abg, "paco2", x_ref_abg, ABG_CO2_REF)
pred_death_abg <- mkpred(fit_death_abg, subset_data_abg, "paco2", x_ref_abg, ABG_CO2_REF)
pred_hcrf_abg <- mkpred(fit_hcrf_abg, subset_data_abg, "paco2", x_ref_abg, ABG_CO2_REF)
axis_abg_ipw_trim <- compute_or_axis_spec(
  list(pred_imv_abg, pred_niv_abg, pred_death_abg, pred_hcrf_abg),
  lo_col = "LCL", hi_col = "UCL"
)
# 5. plotting
# Plotting deferred until VBG curves are computed so axes can be shared.

```

Restricting plots bewtween 0.02 and 0.98

2.0.3 5.3 ABG IPW spline models (2–98th percentile)

```

subset_data_abg <- subset_data %>%
  filter(!is.na(paco2)) %>%                                # implies has_abg == 1
  select(paco2, imv_proc, niv_proc, death_60d,
         hypercap_resp_failure, w_abg, all_of(adj_core)) %>%
  filter(complete.cases(.))

fitfun <- function(formula, outcome) {
  fit_res <- fit_with_diagnostics(
    function() svyglm(
      formula,
      design = svydesign(ids = ~1, weights = ~w_abg, data = subset_data_abg),
      family = quasibinomial(),
      control = glm.control(maxit = 50)
    ),

```

```

context = make_context(
  stage = "outcome",
  component = "spline",
  analysis_variant = "ipw",
  model_type = "spline",
  group = "ABG",
  outcome = outcome,
  imputation = NA_integer_,
  batch = NA_integer_
)
)
append_outcome_diag(fit_res$diag)
fit_res$fit
}

fit_imv_abg <- fitfun(make_spline_fml("imv_proc", "paco2", adj_core), "imv_proc")
fit_niv_abg <- fitfun(make_spline_fml("niv_proc", "paco2", adj_core), "niv_proc")
fit_death_abg <- fitfun(make_spline_fml("death_60d", "paco2", adj_core), "death_60d")
fit_hcrcf_abg <- fitfun(make_spline_fml("hypercap_resp_failure", "paco2", adj_core),
                           "hypercap_resp_failure")
if (any(vapply(list(fit_imv_abg, fit_niv_abg, fit_death_abg, fit_hcrcf_abg), is.null, logical(1)))) {
  stop("IPW ABG spline fits (trimmed) failed; see model_fit_diagnostics.csv.")
}

# 4. prediction helper
mkpred <- function(fit, data_ref, co2_var, ref_df, co2_ref) {
  q <- stats::quantile(data_ref[[co2_var]], probs = c(0.02, 0.98), na.rm = TRUE)
  grid_info <- make_co2_grid_ref(
    co2_var,
    seq(q[1], q[2], length.out = SPLINE_GRID_N),
    ref_df,
    co2_ref
  )
  predict_or_curve_from_fit(fit, grid_info$grid, grid_info$ref_idx, co2_var)
}

```

```

pred_imv_abg   <- mkpred(fit_imv_abg,   subset_data_abg, "paco2", x_ref_abg, ABG_CO2_REF)
pred_niv_abg   <- mkpred(fit_niv_abg,   subset_data_abg, "paco2", x_ref_abg, ABG_CO2_REF)
pred_death_abg <- mkpred(fit_death_abg, subset_data_abg, "paco2", x_ref_abg, ABG_CO2_REF)
pred_hcrcf_abg <- mkpred(fit_hcrcf_abg, subset_data_abg, "paco2", x_ref_abg, ABG_CO2_REF)

# 5. plotting
xlab <- expression(paste("ABG CO" [2], " (mmHg)"))

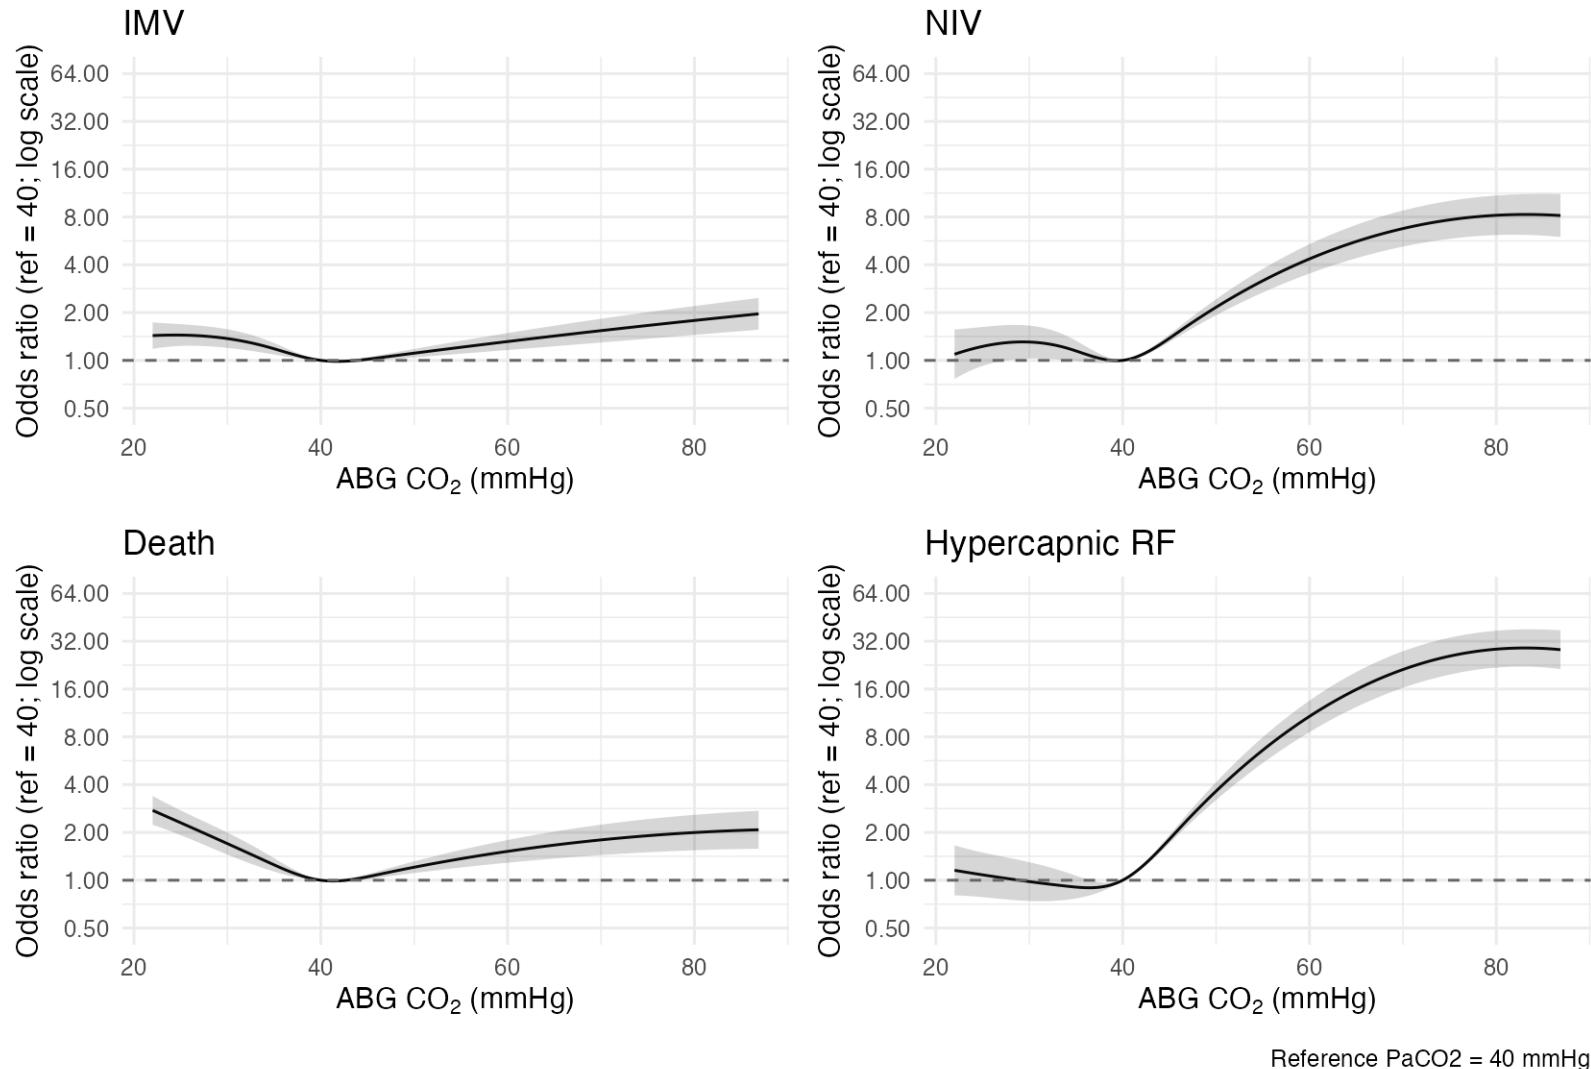
plt <- function(dat, title)
  ggplot(dat, aes(paco2, OR)) +
    geom_line() +
    geom_ribbon(aes(ymin = LCL, ymax = UCL), alpha = 0.2) +
    geom_hline(yintercept = 1, linetype = "dashed", color = "grey40") +
    or_axis_scale(axis_abg_ipw_trim) +
    labs(title = title, x = xlab,
         y = paste0("Odds ratio (ref = ", ABG_CO2_REF, "; log scale)")) +
    theme_minimal()

ipw_abg_panel <- (patchwork::wrap_plots(
  plt(pred_imv_abg,    "IMV"),
  plt(pred_niv_abg,    "NIV"),
  plt(pred_death_abg,  "Death"),
  plt(pred_hcrcf_abg,  "Hypercapnic RF"),
  ncol = 2
)) +
  plot_annotation(
    title = paste0("Propensity-weighted adjusted odds ratios by ABG CO2 (ref = ",
                  ABG_CO2_REF, "; conditional on X; 2-98% range)"),
    caption = paste0("Reference PaCO2 = ", ABG_CO2_REF, " mmHg")
  )

print_plot_once(ipw_abg_panel, "spline-ipw-abg-trimmed", width = 8.5, height = 6)

```

Propensity-weighted adjusted odds ratios by ABG CO₂ (ref = 40; conditional on X; 2–98% range)



VBG uses the same GBM tuning as ABG (shared `gbm_params`).

2.0.4 5.4 VBG IPW weighting and spline models

```
# Inverse-propensity weighting & outcome modelling for **VBG** cohort
#   - mirrored 1-to-1 to the validated ABG workflow

set.seed(42)

# 1. IPW for VBG -----
set.seed(42)
gbm_df_vbg <- subset_data[, c("has_abg", "has_vbg", covars_gbm), drop = FALSE]
gbm_df_vbg <- normalize_types(gbm_df_vbg, levels_ref)
gbm_df_vbg <- droplevels_all(gbm_df_vbg)
gbm_preflight(gbm_df_vbg, covars_gbm, "unimp_vbg")
append_mem_snapshot("gbm_unimp", "unimp_vbg", "pre")
w_vbg <- do.call(
  weightit,
  c(
    list(
      formula_vbg,
      data      = gbm_df_vbg,
      method    = "gbm",
      estimand  = "ATE",
      missing   = "ind",
      include.obj = FALSE
    ),
    gbm_params
  )
)
append_mem_snapshot("gbm_unimp", "unimp_vbg", "post")

# One-sided IPSW (VBG observed only) + truncation of small propensities
ipow_vbg <- compute_ipow_weights(
  w_vbg,
  treat = gbm_df_vbg$has_vbg,
  ps_floor_quantile = ps_trunc_quantile,
```

```

    stabilize = TRUE
)
w_vbg_ipow <- ipow_vbg$weights
ps_floor_vbg <- ipow_vbg$ps_floor
subset_data$trunc_vbg <- ipow_vbg$truncated
subset_data$ps_vbg <- ipow_vbg$ps
subset_data$w_vbg <- w_vbg_ipow
assert_finite_weights(w_vbg_ipow[subset_data$has_vbg == 1], "w_vbg")
rm(w_vbg, gbm_df_vbg)
invisible(gc())

# Balance diagnostics are handled later.

# 2. Analysis set (VBG only) -----
subset_data_vbg <- subset_data %>%
  filter(!is.na(vbg_co2)) %>%
  select(vbg_co2, imv_proc, niv_proc, death_60d,
         hypercap_resp_failure, w_vbg, all_of(adj_core)) %>%
  filter(complete.cases(.))

fitfun <- function(formula, outcome) {
  fit_res <- fit_with_diagnostics(
    function() svyglm(
      formula,
      design = svydesign(ids = ~1, weights = ~w_vbg, data = subset_data_vbg),
      family = quasibinomial(),
      control = glm.control(maxit = 50)
    ),
    context = make_context(
      stage = "outcome",
      component = "spline",
      analysis_variant = "ipw",
      model_type = "spline",
      group = "VBG",
      outcome = outcome,
      imputation = NA_integer_
    )
  )
}
```

```

    batch = NA_integer_
  )
)
append_outcome_diag(fit_res$diag)
fit_res$fit
}

# Model diagrams: IPW VBG spline models
ipw_vbg_rcs_forms <- list(
  "VBG IPW spline (adjusted): IMV ~ CO2 spline + X"      = make_spline_fml("imv_proc", "vbg_co2", adj_core),
  "VBG IPW spline (adjusted): NIV ~ CO2 spline + X"      = make_spline_fml("niv_proc", "vbg_co2", adj_core),
  "VBG IPW spline (adjusted): Death60d ~ CO2 spline + X" = make_spline_fml("death_60d", "vbg_co2", adj_core),
  "VBG IPW spline (adjusted): HCRF ~ CO2 spline + X"      = make_spline_fml("hypercap_resp_failure", "vbg_co2", adj_core)
)
register_model_diagrams(ipw_vbg_rcs_forms)

fit_imv_vbg   <- fitfun(make_spline_fml("imv_proc", "vbg_co2", adj_core), "imv_proc")
fit_niv_vbg   <- fitfun(make_spline_fml("niv_proc", "vbg_co2", adj_core), "niv_proc")
fit_death_vbg <- fitfun(make_spline_fml("death_60d", "vbg_co2", adj_core), "death_60d")
fit_hcrf_vbg  <- fitfun(make_spline_fml("hypercap_resp_failure", "vbg_co2", adj_core),
                           "hypercap_resp_failure")
if (any(vapply(list(fit_imv_vbg, fit_niv_vbg, fit_death_vbg, fit_hcrf_vbg), is.null, logical(1)))) {
  stop("IPW VBG spline fits failed; see model_fit_diagnostics.csv.")
}

# 4. Prediction helper -----
mkpred <- function(fit, data_ref, co2_var, ref_df, co2_ref) {
  co2_seq <- stats::quantile(data_ref[[co2_var]], probs = c(0.02, 0.98), na.rm = TRUE)
  grid_info <- make_co2_grid_ref(
    co2_var,
    seq(co2_seq[1], co2_seq[2], length.out = SPLINE_GRID_N),
    ref_df,
    co2_ref
  )
  predict_or_curve_from_fit(fit, grid_info$grid, grid_info$ref_idx, co2_var)
}

```

```

pred_imv_vbg   <- mkpred(fit_imv_vbg,   subset_data_vbg, "vbg_co2", x_ref_vbg, VBG_CO2_REF)
pred_niv_vbg   <- mkpred(fit_niv_vbg,   subset_data_vbg, "vbg_co2", x_ref_vbg, VBG_CO2_REF)
pred_death_vbg <- mkpred(fit_death_vbg, subset_data_vbg, "vbg_co2", x_ref_vbg, VBG_CO2_REF)
pred_hcrf_vbg  <- mkpred(fit_hcrf_vbg,  subset_data_vbg, "vbg_co2", x_ref_vbg, VBG_CO2_REF)
axis_ipw_common <- compute_or_axis_spec(
  list(pred_imv_abg, pred_niv_abg, pred_death_abg, pred_hcrf_abg,
       pred_imv_vbg, pred_niv_vbg, pred_death_vbg, pred_hcrf_vbg),
  lo_col = "LCL", hi_col = "UCL"
)

# 5. Plotting (gray scheme) -----
xlab <- expression(paste("VBG CO" [2], " (mmHg)"))

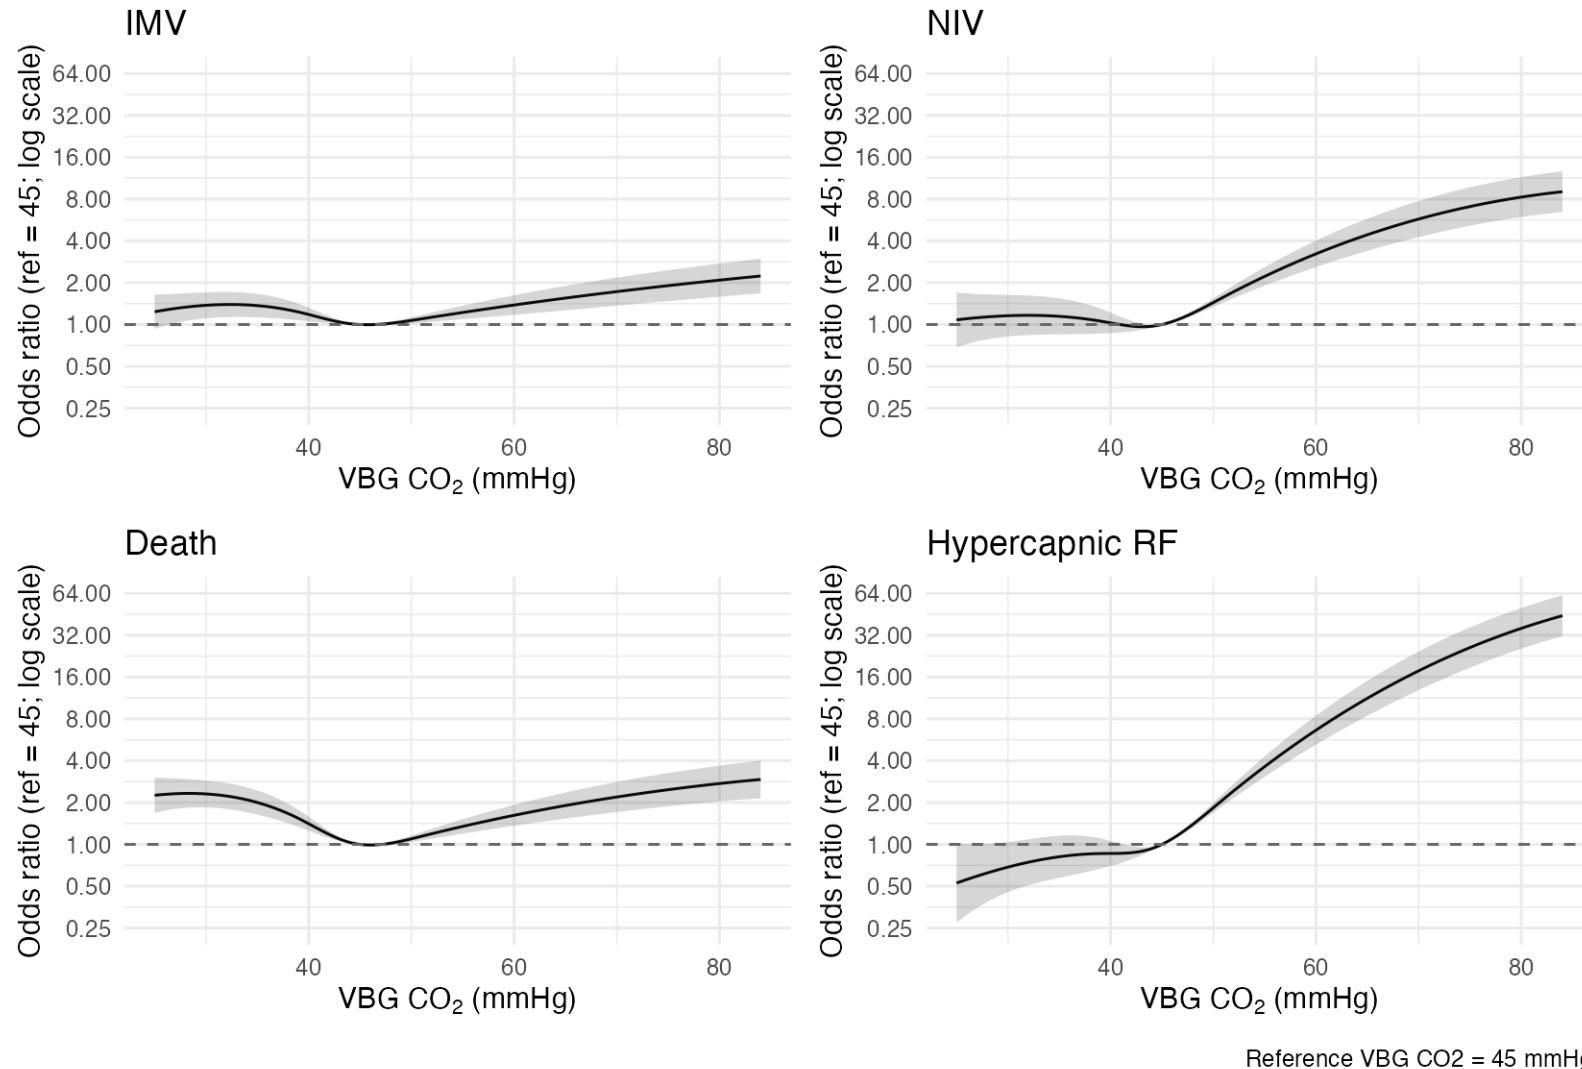
plt <- function(dat, title)
  ggplot(dat, aes(vbg_co2, OR)) +
    geom_line() +
    geom_ribbon(aes(ymin = LCL, ymax = UCL), alpha = 0.2) +
    geom_hline(yintercept = 1, linetype = "dashed", color = "grey40") +
    or_axis_scale(axis_ipw_common) +
    labs(title = title, x = xlab,
         y = paste0("Odds ratio (ref = ", VBG_CO2_REF, "; log scale)")) +
    theme_minimal()

ipw_vbg_panel <- (patchwork::wrap_plots(
  plt(pred_imv_vbg,    "IMV"),
  plt(pred_niv_vbg,    "NIV"),
  plt(pred_death_vbg,  "Death"),
  plt(pred_hcrf_vbg,   "Hypercapnic RF"),
  ncol = 2
)) +
  plot_annotation(
    title = paste0("Propensity-weighted adjusted odds ratios by VBG CO2 (ref = ",
                  VBG_CO2_REF, "; conditional on X)"),
    caption = paste0("Reference VBG CO2 = ", VBG_CO2_REF, " mmHg")
)

```

```
print_plot_once(ipw_vbg_panel, "spline-ipw-vbg", width = 8.5, height = 6)
```

Propensity-weighted adjusted odds ratios by VBG CO₂ (ref = 45; conditional on X)



```

# ABG plots with the same axis (shared with VBG)
xlab_abg <- expression(paste("ABG CO"[2], " (mmHg)"))

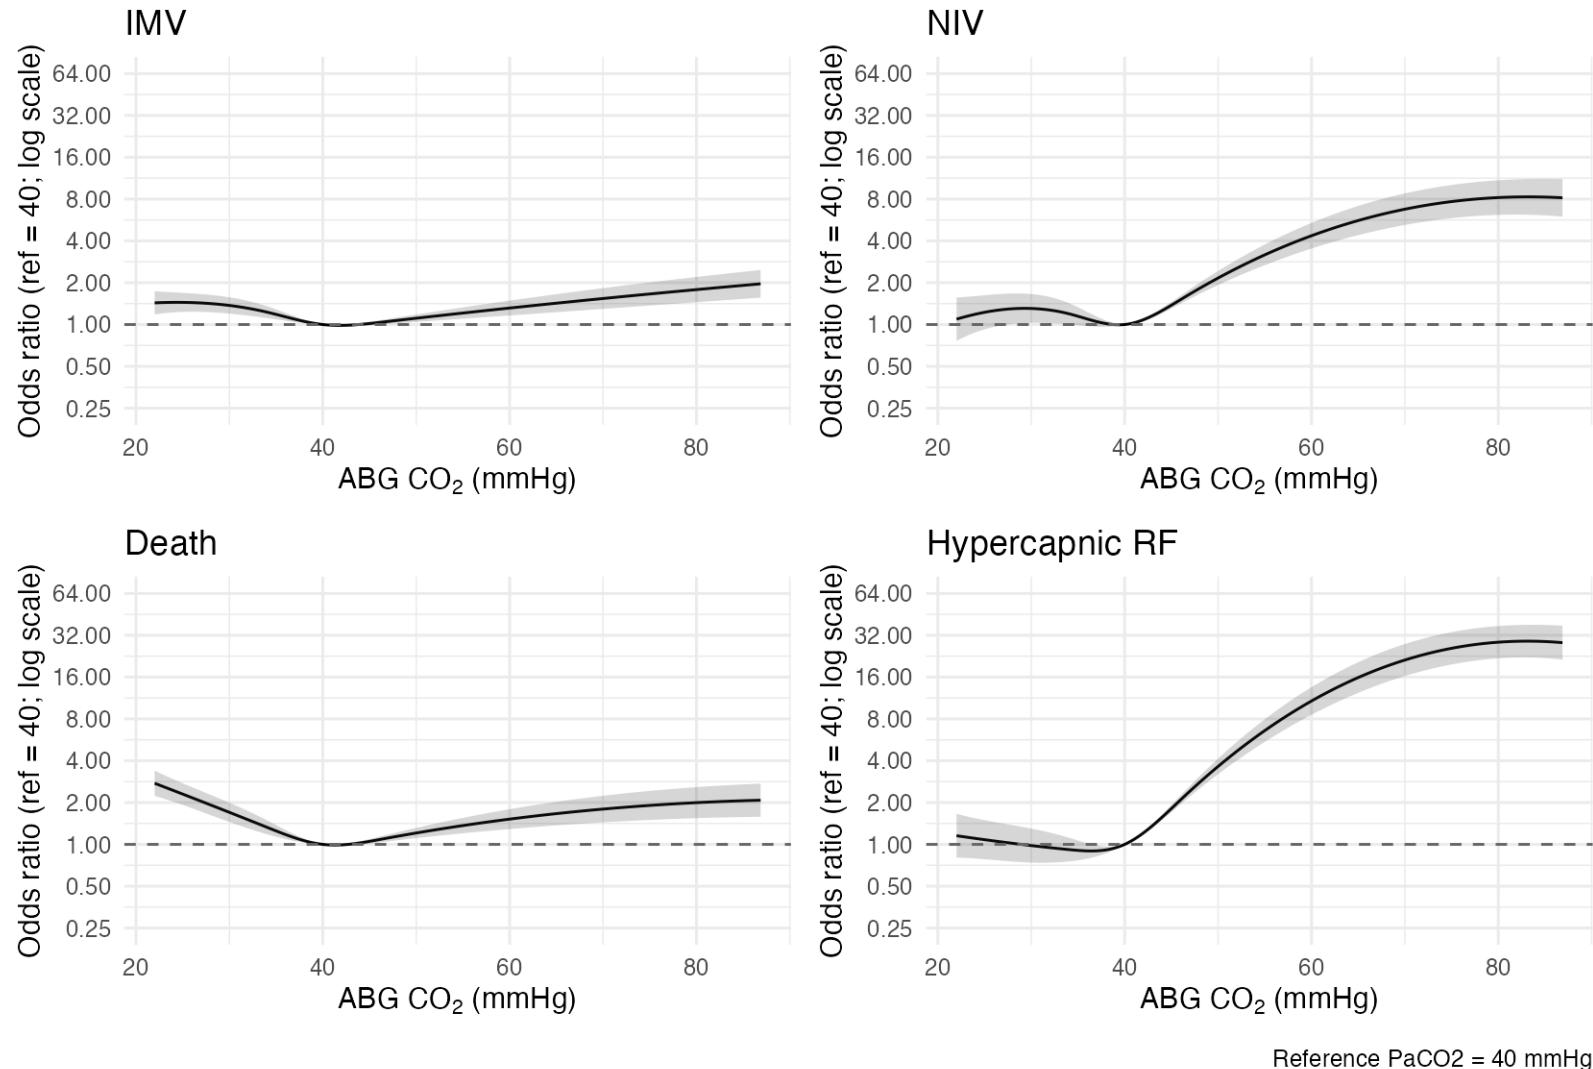
plt_abg <- function(dat, title)
  ggplot(dat, aes(paco2, OR)) +
    geom_line() +
    geom_ribbon(aes(ymin = LCL, ymax = UCL), alpha = 0.2) +
    geom_hline(yintercept = 1, linetype = "dashed", color = "grey40") +
    or_axis_scale(axis_ipw_common) +
    labs(title = title, x = xlab_abg,
         y = paste0("Odds ratio (ref = ", ABG_CO2_REF, "; log scale)")) +
    theme_minimal()

ipw_abg_shared_panel <- (patchwork::wrap_plots(
  plt_abg(pred_imv_abg,    "IMV"),
  plt_abg(pred_niv_abg,    "NIV"),
  plt_abg(pred_death_abg, "Death"),
  plt_abg(pred_hcrf_abg,   "Hypercapnic RF"),
  ncol = 2
)) +
  plot_annotation(
    title = paste0("Propensity-weighted adjusted odds ratios by ABG CO2 (ref = ",
                  ABG_CO2_REF, "; conditional on X)"),
    caption = paste0("Reference PaCO2 = ", ABG_CO2_REF, " mmHg")
  )

print_plot_once(ipw_abg_shared_panel, "spline-ipw-abg-shared", width = 8.5, height = 6)

```

Propensity-weighted adjusted odds ratios by ABG CO₂ (ref = 40; conditional on X)



2.0.5 5.5 Three-level PCO₂ categories (weighted; ABG, VBG)

Three groups with weights and covariate adjustment

```

library(dplyr)
library(survey)
library(broom)
library(ggplot2)
library(scales)

# 1. Ensure PCO2 categories are present
stopifnot(all(c("pc02_cat_abg", "pc02_cat_vbg") %in% names(subset_data)))

# 2. Function: weighted logistic regression & OR extraction
run_weighted_or <- function(data, outcome, cat_var, weight_var, group_name,
                           treat_var, adj_vars) {
  stopifnot(!is.null(treat_var))
  stopifnot(!is.null(adj_vars))
  dat <- data %>%
    filter(
      .data[[treat_var]] == 1,
      !is.na(.data[[cat_var]]),
      !is.na(.data[[outcome]]),
      !is.na(.data[[weight_var]]),
      .data[[weight_var]] > 0
    ) %>%
    mutate(
      !!cat_var := factor(.data[[cat_var]],
                           levels = CO2_CAT_LEVELS)
    ) %>%
    droplevels()

  design <- svydesign(
    ids = ~1,
    weights = as.formula(paste0("~", weight_var)),
    data = dat
  )

  rhs_terms <- c(cat_var, adj_vars)
  fml <- stats::reformulate(rhs_terms, response = outcome)
}

```

```

fit_res <- fit_with_diagnostics(
  function() svyglm(fml, design = design, family = quasibinomial(),
    control = glm.control(maxit = 50)),
  context = make_context(
    stage = "outcome",
    component = "cat3",
    analysis_variant = "ipw",
    model_type = "cat3",
    group = group_name,
    outcome = outcome,
    imputation = NA_integer_,
    batch = NA_integer_
  )
)
append_outcome_diag(fit_res$diag)
if (is.null(fit_res$fit)) {
  stop("run_weighted_or: model fit failed for outcome=", outcome,
    " cat_var=", cat_var, " group=", group_name)
}

tidy(fit_res$fit, exponentiate = TRUE, conf.int = TRUE) %>%
  filter(term != "(Intercept)", startsWith(term, cat_var)) %>%
  mutate(
    group      = group_name,
    outcome    = outcome
  )
}

# 3. Run across outcomes & cohorts
outcomes_ipw <- c("imv_proc", "niv_proc", "death_60d", "hypercap_resp_failure")

ipw_three_level_forms <- list(
  "ABG IPW 3-level: IMV ~ CO2 category + X"      = reformulate(c("pc02_cat_abg", adj_core), response = "imv_proc"),
  "ABG IPW 3-level: NIV ~ CO2 category + X"      = reformulate(c("pc02_cat_abg", adj_core), response = "niv_proc"),
  "ABG IPW 3-level: Death60d ~ CO2 category + X" = reformulate(c("pc02_cat_abg", adj_core), response = "death_60d"),
  "ABG IPW 3-level: HCRF ~ CO2 category + X"     = reformulate(c("pc02_cat_abg", adj_core), response = "hypercap_resp_failure"),
)

```

```

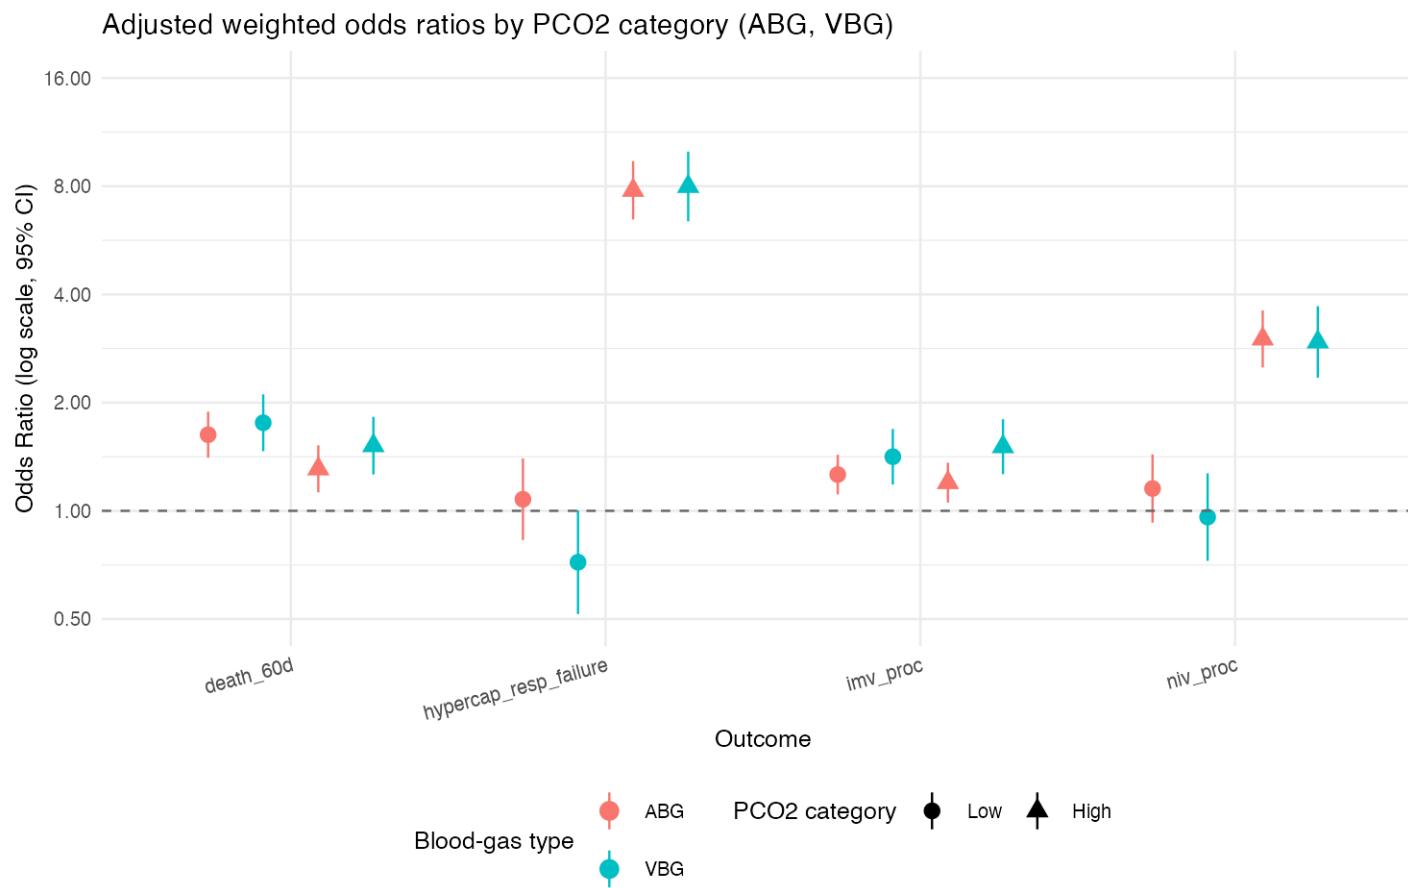
"VBG IPW 3-level: IMV ~ CO2 category + X"      = reformulate(c("pco2_cat_vbg", adj_core), response = "imv_proc"),
"VBG IPW 3-level: NIV ~ CO2 category + X"      = reformulate(c("pco2_cat_vbg", adj_core), response = "niv_proc"),
"VBG IPW 3-level: Death60d ~ CO2 category + X" = reformulate(c("pco2_cat_vbg", adj_core), response = "death_60d"),
"VBG IPW 3-level: HCRF ~ CO2 category + X"      = reformulate(c("pco2_cat_vbg", adj_core), response = "hypercap_resp_failure")
)
register_model_diagrams(ipw_three_level_forms)

ipw_combined_or_df <- bind_rows(
  lapply(outcomes_ipw, function(out)
    run_weighted_or(subset_data, out, "pco2_cat_abg", "w_abg",       "ABG",
                     treat_var = "has_abg", adj_vars = adj_core)),
  lapply(outcomes_ipw, function(out)
    run_weighted_or(subset_data, out, "pco2_cat_vbg", "w_vbg",       "VBG",
                     treat_var = "has_vbg", adj_vars = adj_core))
)
ipw_threoplevel_results <- ipw_combined_or_df %>%
  mutate(method = "IPW adjusted")

ipw_combined_or_df <- map_or_exposure(ipw_combined_or_df, "or-plot-three-level-weighted")
ipw_combined_or_df$group <- factor(ipw_combined_or_df$group, levels = c("ABG", "VBG"))

# 4. Plot weighted odds ratios
ipw_plot_df <- build_or_plot_df(ipw_combined_or_df, "or-plot-three-level-weighted",
                                  expected_exposure_levels = CO2_CAT_CONTRAST_LEVELS)
ipw_axis_spec <- compute_or_axis_spec(ipw_plot_df, lo_col = "conf.low", hi_col = "conf.high",
                                       default_limits = OR_XLIM)
ipw_p_or <- plot_or_safe(
  ipw_plot_df,
  plot_name = "or-plot-three-level-weighted",
  axis_spec = ipw_axis_spec,
  title   = "Adjusted weighted odds ratios by PCO2 category (ABG, VBG)"
)
print_plot_once(ipw_p_or, "or-plot-three-level-weighted", width = 7.5, height = 4.8)

```



2.1 6) Propensity score diagnostics

Plotting propensity scores

```
# --- Propensity score histograms (ABG / VBG) -----
# ABG = arterial blood gas; VBG = venous blood gas

library(dplyr)
library(ggplot2)
```

```

library(scales)

stopifnot("has_abg" %in% names(subset_data))
stopifnot("has_vbg" %in% names(subset_data))
stopifnot(all(c("ps_abg", "ps_vbg") %in% names(subset_data)))

# Build list of per-cohort PS data frames conditionally (so missing cohorts don't error)
ps_dfs_cond <- list(
  ABG = data.frame(
    ps      = subset_data$ps_abg,
    treat   = subset_data$has_abg,
    ScoreType = "ABG"
  ),
  VBG = data.frame(
    ps      = subset_data$ps_vbg,
    treat   = subset_data$has_vbg,
    ScoreType = "VBG"
  )
)

# Bind, clean, and factorize for plotting
df_ps_cond <- bind_rows(ps_dfs_cond) %>%
  filter(!is.na(ps), !is.na(treat)) %>%
  mutate(
    treat   = factor(treat, levels = c(0, 1), labels = c("No Test", "Test")),
    ScoreType = factor(ScoreType, levels = c("ABG", "VBG"))
  )

# Plot
p_ps_cond <- ggplot(df_ps_cond, aes(x = ps, fill = treat)) +
  geom_histogram(aes(y = after_stat(density)), alpha = 0.5,
                 position = "identity", bins = 30) +
  scale_fill_manual(values = c("No Test" = "steelblue", "Test" = "tomato")) +
  facet_wrap(~ScoreType, scales = "free_y") +
  coord_cartesian(xlim = c(0, 1)) +
  labs(

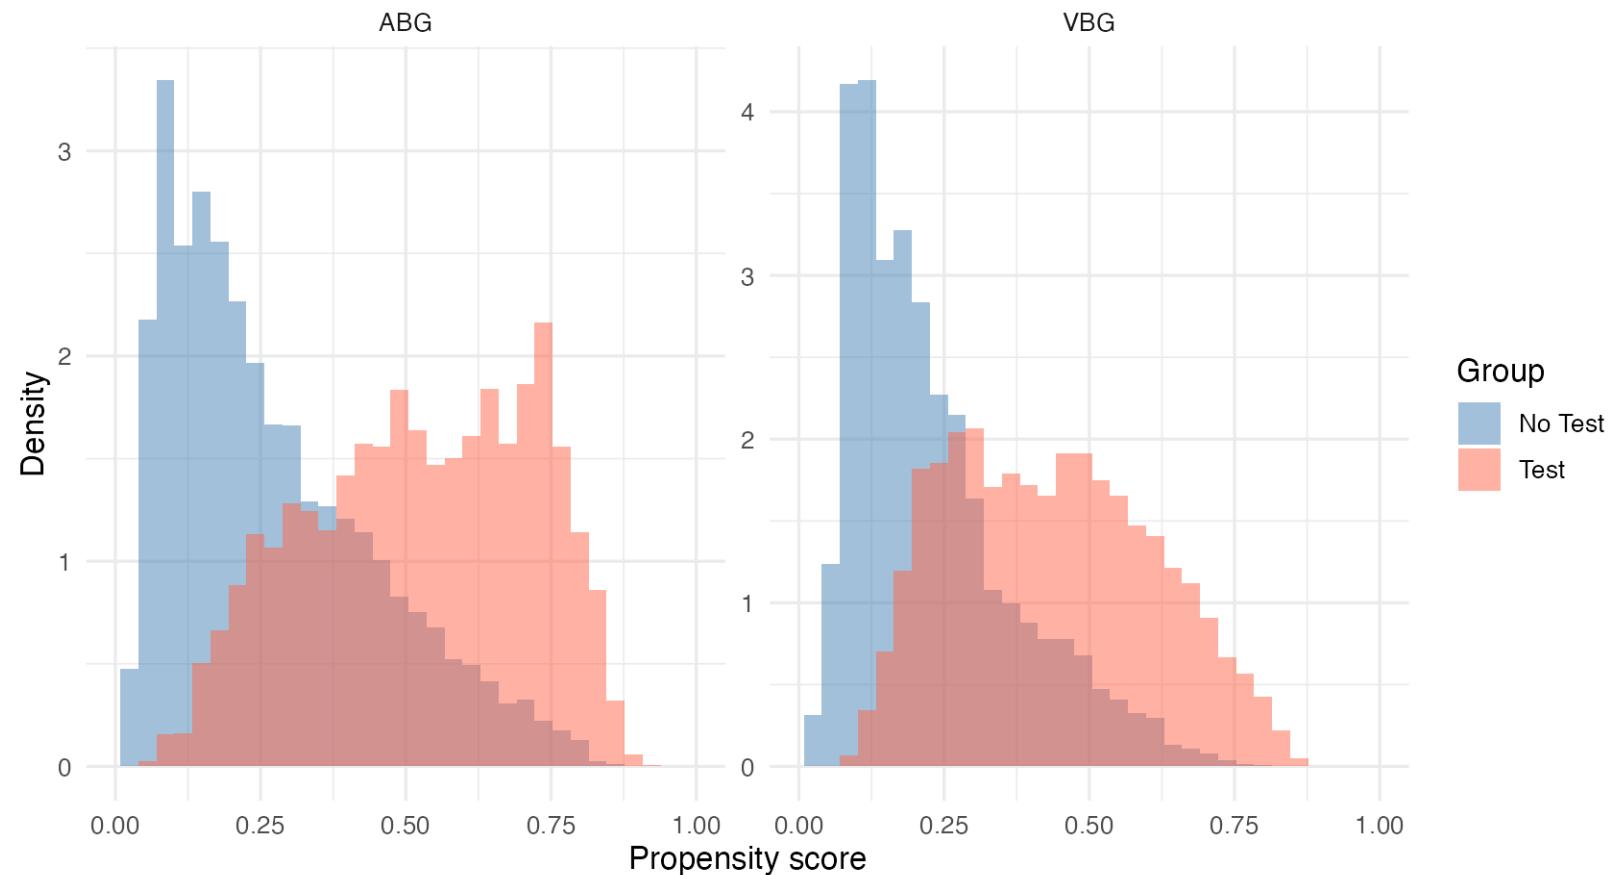
```

```

title = "Propensity Score Distributions",
x     = "Propensity score",
y     = "Density",
fill  = "Group"
) +
theme_minimal(base_size = 12)
print_plot_once(p_ps_cond, "propensity-histograms-conditional", width = 8.5, height = 5)

```

Propensity Score Distributions



3 Multiple Imputation Analysis

added 12/6/2025

3.0.1 7.2 Missingness structure and drivers

Preview (top 20 rows). Full table saved to Results/missingness_by_strata_preview.csv.

Pre-imputation missingness

| Variable | Missing (n) | Missing (%) |
|------------------|-------------|-------------|
| vbg_o2sat | 22,366 | 86.5% |
| bnp | 21,181 | 81.9% |
| vbg_co2 | 18,392 | 71.1% |
| spo2 | 18,319 | 70.9% |
| paco2 | 16,490 | 63.8% |
| serum_lac | 15,571 | 60.2% |
| curr_bmi | 14,723 | 57.0% |
| serum_phos | 13,762 | 53.2% |
| temp_new | 12,460 | 48.2% |
| hr | 9,418 | 36.4% |
| dbp | 7,738 | 29.9% |
| sbp | 7,666 | 29.7% |
| wbc | 4,542 | 17.6% |
| serum_ca | 2,627 | 10.2% |
| serum_cr | 2,396 | 9.3% |
| plt | 2,080 | 8.0% |
| serum_k | 2,021 | 7.8% |
| serum_cl | 1,498 | 5.8% |
| serum_hco3 | 1,460 | 5.6% |
| sodium | 1,326 | 5.1% |
| age_at_encounter | 0 | 0.0% |
| sex | 0 | 0.0% |
| race_ethnicity | 0 | 0.0% |
| copd | 0 | 0.0% |
| asthma | 0 | 0.0% |
| osa | 0 | 0.0% |
| chf | 0 | 0.0% |
| acute_nmd | 0 | 0.0% |
| phtn | 0 | 0.0% |
| ckd | 0 | 0.0% |
| dm | 0 | 0.0% |
| location | 0 | 0.0% |
| encounter_type | 0 | 0.0% |
| has_abg | 0 | 0.0% |
| has_vbg | 0 | 0.0% |
| imv_proc | 0 | 0.0% |

Table 7: Missingness by key strata (pre-imputation; top 10 variables; full table saved to Results/missingness-by-strata.csv).

| level | variable | pct_missing | stratum | pct_missing_overall |
|-------|-----------|-------------|----------|---------------------|
| 0 | vbg_o2sat | 86.4 | has_abg | 86.5 |
| 1 | vbg_o2sat | 86.7 | has_abg | 86.5 |
| 0 | vbg_o2sat | 99.2 | has_vbg | 86.5 |
| 1 | vbg_o2sat | 55.3 | has_vbg | 86.5 |
| 0 | vbg_o2sat | 87.1 | inv_proc | 86.5 |
| 1 | vbg_o2sat | 82.1 | inv_proc | 86.5 |
| 0 | bnp | 84.4 | has_abg | 81.9 |
| 1 | bnp | 77.7 | has_abg | 81.9 |
| 0 | bnp | 82.8 | has_vbg | 81.9 |
| 1 | bnp | 79.8 | has_vbg | 81.9 |
| 0 | bnp | 82.4 | inv_proc | 81.9 |
| 1 | bnp | 78.5 | inv_proc | 81.9 |
| 0 | vbg_co2 | 71.5 | has_abg | 71.1 |
| 1 | vbg_co2 | 70.5 | has_abg | 71.1 |
| 0 | vbg_co2 | 100.0 | has_vbg | 71.1 |
| 1 | vbg_co2 | 0.0 | has_vbg | 71.1 |
| 0 | vbg_co2 | 72.6 | inv_proc | 71.1 |
| 1 | vbg_co2 | 59.8 | inv_proc | 71.1 |
| 0 | spo2 | 69.6 | has_abg | 70.9 |
| 1 | spo2 | 73.1 | has_abg | 70.9 |

3.0.2 7.3 Monte Carlo error check after MI

3.1 8) Pre-imputation data prep (consistent types & predictors)

Why: MI models need coherent types; using exactly the same covariates as the propensity score models avoids model drift.

3.2 9) Imputation model specification (MICE)

3.2.1 9.1 Predictor matrix & methods. Run MICE (moderate settings for scale)

```
# --- variables for GBM propensity (kept identical to main analysis) ---
# ---- MICE setup: include PaC02/VBG CO2 as predictors but do not impute ----
library(mice)
library(dplyr)
```

Table 8: Predictors of missingness (logit OR; top 20 by p-value; full table saved to Results/missingness-drivers.csv).

| variable | term | OR | LCL | UCL | p.value |
|------------|-------------------------|---------|---------|---------|---------|
| temp_new | location1 | 18.41 | 15.92 | 21.29 | 0 |
| temp_new | location3 | 21.44 | 18.95 | 24.26 | 0 |
| serum_phos | encounter_typeInpatient | 0.12 | 0.11 | 0.14 | 0 |
| hr | location1 | 11.68 | 10.14 | 13.46 | 0 |
| sbp | encounter_typeInpatient | 0.04 | 0.04 | 0.05 | 0 |
| dbp | encounter_typeInpatient | 0.05 | 0.04 | 0.06 | 0 |
| serum_lac | has_abg | 0.24 | 0.22 | 0.26 | 0 |
| vbg_o2sat | has_vbg | 0.00 | 0.00 | 0.00 | 0 |
| dbp | location3 | 1607.67 | 847.74 | 3048.81 | 0 |
| serum_lac | location3 | 3.24 | 2.89 | 3.62 | 0 |
| serum_phos | has_abg | 0.34 | 0.31 | 0.38 | 0 |
| sbp | location3 | 3373.24 | 1382.62 | 8229.86 | 0 |
| hr | location2 | 4.34 | 3.68 | 5.13 | 0 |
| wbc | location3 | 3.13 | 2.72 | 3.61 | 0 |
| temp_new | location2 | 3.76 | 3.17 | 4.45 | 0 |
| serum_cl | age_at_encounter | 0.96 | 0.96 | 0.97 | 0 |
| sodium | age_at_encounter | 0.96 | 0.96 | 0.97 | 0 |
| serum_ca | age_at_encounter | 0.97 | 0.97 | 0.97 | 0 |
| serum_hco3 | age_at_encounter | 0.97 | 0.96 | 0.97 | 0 |
| serum_phos | has_vbg | 0.38 | 0.33 | 0.43 | 0 |

```
# --- add analysis targets and CO2 measures explicitly -----
mi_vars <- setdiff(unique(c(
  covars_gbm,
  "has_abg", "has_vbg", # treatments (NOT imputed)
  "imv_proc", "niv_proc", "death_60d", "hypercap_resp_failure", # outcomes (NOT imputed)
  co2_vars
)), drop_vars_ultra_missing)

mi_df <- subset_data[, mi_vars, drop = FALSE]
mi_df <- normalize_types(mi_df, levels_ref)

mi_df_size <- utils::object.size(mi_df)
message("MI data size (bytes): ", format(mi_df_size, units = "auto"))

# Rough memory preflight based on total missing cells (imputed values only).
miss_counts <- vapply(mi_df, function(x) sum(is.na(x)), numeric(1))
miss_total <- sum(miss_counts)
```

```

if (is.finite(miss_total) && miss_total > 0) {
  est_bytes <- miss_total * M_IMP * 8
  message("MI imputation storage estimate: ", format(structure(est_bytes, class = "object_size"), units = "auto"))
  if (est_bytes > MI_MAX_BYTES) {
    m_max <- floor(MI_MAX_BYTES / (miss_total * 8))
    m_target <- max(50L, m_max)
    if (m_target < 50L) {
      stop("Estimated MI storage exceeds memory (m=", M_IMP,
           ", estimated=", format(structure(est_bytes, class = "object_size"), units = "auto"),
           "). Reduce missingness or MI scope, or increase available memory.")
    }
    if (m_target < M_IMP) {
      message("Reducing M_IMP from ", M_IMP, " to ", m_target,
              " to stay within MI_MAX_BYTES.")
      M_IMP <- m_target
    }
  }
}

# Make binary comorbid factors so "logreg" is used (and stays binary)
bin_covars <- c("copd", "asthma", "osa", "chf", "acute_nmd", "phtn", "ckd", "dm")
missing_bin <- setdiff(bin_covars, names(mi_df))
stopifnot(length(missing_bin) == 0)
mi_df[bin_covars] <- lapply(mi_df[bin_covars], function(z) {
  if (is.factor(z)) return(droplevels(z))
  zz <- suppressWarnings(as.integer(z))
  factor(zz, levels = c(0L, 1L), labels = c("0", "1"))
})

# For MICE: convert any remaining characters → factors
mi_df <- dplyr::mutate(mi_df, across(where(is.character), ~ factor(.x)))

# Guardrail: high-cardinality factors can blow up MICE model matrices.
# Exclude them from predictorMatrix and (if missing) make "Missing" explicit.
high_card <- names(which(vapply(mi_df, function(x) is.factor(x) && nlevels(x) > MAX_LEVELS_PRED, logical(1))))
if (length(high_card)) {

```

```

message("MICE: high-cardinality factors detected (nlevels > ", MAX_LEVELS_PRED, "): ",
       paste(high_card, collapse = ", "))
for (v in high_card) {
  if (any(is.na(mi_df[[v]]))) {
    lv <- levels(mi_df[[v]])
    tmp <- as.character(mi_df[[v]])
    tmp[is.na(tmp)] <- "Missing"
    mi_df[[v]] <- factor(tmp, levels = unique(c(lv, "Missing")))
    if (!is.null(levels_ref) && !is.null(levels_ref[[v]])) &&
      !"Missing" %in% levels_ref[[v]]) {
      levels_ref[[v]] <- c(levels_ref[[v]], "Missing")
    }
  }
}

# --- methods & predictor matrix aligned to *mi_df* -----
meth <- mice::make.method(mi_df)

is_fac      <- vapply(mi_df, is.factor, logical(1))
is_num      <- vapply(mi_df, is.numeric, logical(1))
is_bin_fac  <- vapply(mi_df, function(x) is.factor(x) && nlevels(x) == 2, logical(1))
is_multicat <- vapply(mi_df, function(x) is.factor(x) && nlevels(x) > 2, logical(1))

# robust defaults
meth[is_num]      <- "pmm"      # numerics: predictive mean matching
meth[is_multicat] <- "polyreg"   # unordered multicategory
meth[is_bin_fac]  <- "logreg"   # binary factors: logistic regression

# never impute treatments, outcomes, or CO2 exposures
no_imp <- c("has_abg", "has_vbg", "imv_proc", "niv_proc", "death_60d", "hypercap_resp_failure",
           "paco2", "vbg_co2")
if (length(high_card)) no_imp <- unique(c(no_imp, high_card))
meth[intersect(names(meth), no_imp)] <- ""

# predictor matrix; force has_abg/has_vbg as predictors, but do not impute no_imp

```

```

pred <- mice:::quickpred(mi_df, mincor = MINCOR_QUICKPRED, minpuc = MINPUC_QUICKPRED)
for (nm in intersect(c("has_abg", "has_vbg"), colnames(pred))) {
  pred[, nm] <- 1
}
pred[intersect(rownames(pred), no_imp), ] <- 0

if (length(high_card)) {
  pred[, intersect(high_card, colnames(pred))] <- 0
  pred[intersect(high_card, rownames(pred)), ] <- 0
}

# Ensure dropped covariates do not appear as predictors
drop_covars <- intersect(drop_vars_ultra_missing, colnames(pred))
if (length(drop_covars)) {
  pred[, drop_covars] <- 0
}

# Non-imputed variables with missingness should NOT be predictors
no_imp_with_missing <- intersect(no_imp, names(mi_df))
no_imp_with_missing <- no_imp_with_missing[
  vapply(mi_df[no_imp_with_missing], function(x) any(is.na(x)), logical(1))
]
no_imp_with_missing <- setdiff(no_imp_with_missing, c("has_abg", "has_vbg"))

pred_cols_check <- intersect(c("paco2", "vbg_co2"), colnames(pred))
if (length(pred_cols_check)) {
  message(
    "MICE: predictor column sums (pre-exclude) for ",
    paste(pred_cols_check, collapse = ", "),
    ":",
    paste(colSums(pred[, pred_cols_check, drop = FALSE]), collapse = ", ")
  )
}
if (length(no_imp_with_missing)) {
  pred[, intersect(no_imp_with_missing, colnames(pred))] <- 0
  message("MICE: excluded non-imputed missing predictors: ",
}

```

```

        paste(no_imp_with_missing, collapse = " ", ""))
}

if (length(pred_cols_check)) {
  message(
    "MICE: predictor column sums (post-exclude) for ",
    paste(pred_cols_check, collapse = " ", ""),
    ": ",
    paste(colSums(pred[, pred_cols_check, drop = FALSE]), collapse = " ", "")
  )
}

# Ensure key covariates with missingness have predictors (avoid zero-row pred)
core_preds <- intersect(
  c("age_at_encounter", "sex", "race_ethnicity", "location", "encounter_type",
    "has_abg", "has_vbg", "imv_proc", "niv_proc", "death_60d",
    "hypercap_resp_failure", "paco2", "vbg_co2"),
  colnames(pred)
)
core_preds <- setdiff(core_preds, no_imp_with_missing)
vars_need_pred <- intersect(covars_gbm, rownames(pred))
vars_need_pred <- setdiff(vars_need_pred, no_imp)
vars_need_pred <- vars_need_pred[vapply(mi_df[vars_need_pred], function(x) any(is.na(x)), logical(1))]
zero_pred <- vars_need_pred[rowSums(pred[vars_need_pred, , drop = FALSE] != 0, na.rm = TRUE) == 0]
if (length(zero_pred)) {
  message("MICE predictor rows empty for: ", paste(zero_pred, collapse = " ", ),
         ". Adding core predictors.")
  pred[zero_pred, core_preds] <- 1
}
pred[intersect(rownames(pred), no_imp), ] <- 0
pred[c(zero_pred), c(zero_pred)] <- 0

# --- Preflight and cap predictor rows to avoid huge model matrices -------

n_mm_cols <- function(pred_row, df) {
  preds <- names(which(pred_row != 0))
  cols <- 1L
}

```

```

for (p in preds) {
  x <- df[[p]]
  if (is.factor(x)) {
    lv <- nlevels(x)
    cols <- cols + max(1L, lv - 1L)
  } else {
    cols <- cols + 1L
  }
}
cols
}

pred_width_preflight <- function(pred, df, meth) {
  vars <- names(meth)[meth != ""]
  vars <- vars[vapply(df[vars], function(x) any(is.na(x)), logical(1))]
  rows <- lapply(vars, function(v) {
    pred_row <- pred[v, ]
    n_pred <- sum(pred_row != 0, na.rm = TRUE)
    mm_cols <- n_mm_cols(pred_row, df)
    x <- df[[v]]
    nlevels_v <- if (is.factor(x)) nlevels(x) else NA_integer_
    miss_n <- sum(is.na(x))
    data.frame(
      variable = v,
      method = meth[[v]],
      n_pred = n_pred,
      mm_cols = mm_cols,
      nlevels_v = nlevels_v,
      miss_n = miss_n,
      stringsAsFactors = FALSE
    )
  })
  if (!length(rows)) {
    return(data.frame(variable = character(), method = character(), n_pred = integer(),
                     mm_cols = integer(), nlevels_v = integer(), miss_n = integer()))
  }
}

```

```

    dplyr::bind_rows(rows)
}

preflight_pre <- pred_width_preflight(pred, mi_df, meth) |>
  dplyr::mutate(stage = "pre")
if (nrow(preflight_pre)) {
  top_mm <- preflight_pre |> dplyr::arrange(dplyr::desc(mm_cols)) |> dplyr::slice_head(n = 10)
  top_np <- preflight_pre |> dplyr::arrange(dplyr::desc(n_pred)) |> dplyr::slice_head(n = 10)
  print(top_mm)
  print(top_np)
}

core_preds <- intersect(
  c("age_at_encounter", "sex", "race_ethnicity", "location", "encounter_type",
    "has_abg", "has_vbg", "imv_proc", "niv_proc", "death_60d",
    "hypercap_resp_failure", "paco2", "vbg_co2"),
  colnames(pred))
)
core_preds <- setdiff(core_preds, no_imp_with_missing)

set.seed(MI_SEED + 100)
idx <- sample.int(nrow(mi_df), min(COR_SAMPLE_N, nrow(mi_df)))

vars_cap <- names(meth)[meth != ""]
vars_cap <- vars_cap[vapply(mi_df[vars_cap], function(x) any(is.na(x)), logical(1))]

pred_nlevels <- function(x) {
  if (is.factor(x)) nlevels(x) else 1L
}

for (v in vars_cap) {
  cand <- names(which(pred[v, ] != 0))
  if (!length(cand)) next
  keep <- intersect(core_preds, cand)
  rem  <- setdiff(cand, keep)
}

```

```

if (length(rem)) {
  y_raw <- mi_df[[v]][idx]
  y <- if (is.factor(y_raw)) as.integer(y_raw) else suppressWarnings(as.numeric(y_raw))
  scores <- vapply(rem, function(r) {
    x_raw <- mi_df[[r]][idx]
    x <- if (is.factor(x_raw)) as.integer(x_raw) else suppressWarnings(as.numeric(x_raw))
    suppressWarnings(abs(stats::cor(y, x, use = "pairwise.complete.obs")))
  }, numeric(1))
  ord <- order(is.na(scores), -scores)
  rem_keep <- rem[ord]
  rem_keep <- rem_keep[seq_len(min(length(rem_keep), max(0L, MAX_PRED_PER_VAR - length(keep))))]
} else {
  rem_keep <- character()
}

keep_all <- unique(c(keep, rem_keep))
pred[v, ] <- 0
pred[v, keep_all] <- 1
pred[v, v] <- 0

mm_cols <- n_mm_cols(pred[v, ], mi_df)
if (mm_cols > MAX_MM_COLS) {
  drop_pool <- setdiff(keep_all, keep)
  if (!length(drop_pool)) drop_pool <- keep
  drop_order <- drop_pool[order(vapply(drop_pool, function(p) pred_nlevels(mi_df[[p]]), integer(1)),
                                 decreasing = TRUE))]
  for (p in drop_order) {
    if (p %in% keep && length(keep) == 1) break
    pred[v, p] <- 0
    mm_cols <- n_mm_cols(pred[v, ], mi_df)
    if (mm_cols <= MAX_MM_COLS) break
  }
}
}

preflight_post <- pred_width_preflight(pred, mi_df, meth) |>

```

```

dplyr::mutate(stage = "post")
preflight_all <- dplyr::bind_rows(preflight_pre, preflight_post)
write_csv_safely(preflight_all, results_path("mice_pred_width_preflight.csv"), row_names = FALSE)
if (nrow(preflight_post)) {
  top_mm_post <- preflight_post |> dplyr::arrange(dplyr::desc(mm_cols)) |> dplyr::slice_head(n = 10)
  print(top_mm_post)
}

# MI integrity: treatments/outcomes excluded; binaries use logreg
stopifnot(all(meth[no_imp] == ""))
stopifnot(all(rowSums(pred[intersect(rownames(pred), no_imp), , drop = FALSE]) == 0))
bin_fac <- names(which(vapply(mi_df, function(x) is.factor(x) && nlevels(x) == 2, logical(1))))
bin_fac <- setdiff(bin_fac, no_imp)
stopifnot(all(meth[bin_fac] == "logreg"))

# integrity checks
stopifnot(
  ncol(pred) == ncol(mi_df),
  nrow(pred) == ncol(mi_df),
  length(meth) == ncol(mi_df),
  identical(names(meth), colnames(mi_df)))
)

# --- run MICE -----
mi_df_run <- mi_df
M_IMP_RUN <- M_IMP
MAXIT_RUN <- MAXIT_MI
FORCE_MI_BATCHED <- nrow(mi_df_run) > MI_BATCH_THRESHOLD

if (!requireNamespace("digest", quietly = TRUE)) {
  stop("Package 'digest' is required for MI checkpoint signatures.")
}

make_mi_signature <- function(df) {
  classes <- vapply(df, function(x) class(x)[1], character(1))
  na_counts <- vapply(df, function(x) sum(is.na(x)), integer(1))
}

```

```

nlevels <- vapply(df, function(x) if (is.factor(x)) nlevels(x) else NA_integer_, integer(1))
lvl_hash <- vapply(df, function(x) {
  if (!is.factor(x)) return(NA_character_)
  digest::digest(levels(x), algo = "xxhash64")
}, character(1))
col_sig <- data.frame(
  name = names(df),
  class = classes,
  na = na_counts,
  nlevels = nlevels,
  lvl_hash = lvl_hash,
  stringsAsFactors = FALSE
)
hash <- digest::digest(
  list(dim = dim(df), col_sig = col_sig),
  algo = "xxhash64"
)
list(
  hash = hash,
  nrow = nrow(df),
  ncol = ncol(df),
  run_mode = RUN_MODE,
  pilot_frac = PILOT_FRAC,
  mi_pilot_mode = RUN_MODE,
  sample_seed = SAMPLE_SEED,
  mi_seed = MI_SEED,
  col_sig = col_sig
)
}

mi_smoke_log_path <- results_path("mice_smoketest.log")
write_smoke_log <- function(lines) {
  write_diag_lines(lines, mi_smoke_log_path)
}

mem_max <- tryCatch(mem.maxVSize(), error = function(e) NA_real_)
```

```

mem_env <- Sys.getenv("R_MAX_VSIZE", unset = NA_character_)
gc_pre <- utils::capture.output(gc())
write_smoke_log(c(
  paste0("MI smoke test log: ", Sys.time()),
  paste0("mem.maxVSize: ", ifelse(is.na(mem_max), "NA", format(mem_max, scientific = FALSE))),
  paste0("R_MAX_VSIZE env: ", ifelse(nchar(mem_env), mem_env, "NA")),
  paste0("mi_df size: ", format(utils::object.size(mi_df), units = "auto")),
  paste0("mi_df_run size: ", format(utils::object.size(mi_df_run), units = "auto")),
  "gc() pre:",
  gc_pre
))
if (RUN_MODE == "full" && is.finite(mem_max) && mem_max < 2.2e10) {
  message("Full run on ~16GB memory: consider running MI/weights on a >32GB machine for speed.")
}

subset_data_saved <- FALSE
subset_data_path <- results_path("subset_data_pre_mi.rds")
if (MI_MEMORY_HYGIENE) {
  saveRDS(subset_data, subset_data_path)
  rm(subset_data)
  subset_data_saved <- TRUE
  invisible(gc())
}

run_mice_call <- function(m_val, maxit, seed_val, print_flag = FALSE) {
  mice::mice(
    data      = mi_df_run,
    m         = m_val,
    maxit     = maxit,
    predictorMatrix = pred,
    method    = meth,
    printFlag = print_flag,
    seed      = seed_val
  )
}

```

```

if (MI_SMOKE_TEST) {
  smoke_con <- file(mi_smoke_log_path, open = "at")
  sink(smoke_con, type = "output")

  cap_smoke <- capture_warnings(
    tryCatch(
      run_mice_call(m_val = 1L, maxit = 1L, seed_val = MI_SEED, print_flag = TRUE),
      error = function(e) e
    ),
    context = make_context(
      stage = "MI",
      component = "mice_smoketest",
      analysis_variant = "mi",
      model_type = "mice",
      group = NA_character_,
      outcome = NA_character_,
      imputation = NA_integer_,
      batch = NA_integer_
    )
  )
  append_warnings(cap_smoke$warnings)
  sink(type = "output")
  close(smoke_con)

  if (inherits(cap_smoke$value, "error")) {
    smoke_msg <- conditionMessage(cap_smoke$value)
    message("MICE smoke test failed: ", smoke_msg)
    write_smoke_log(c(
      "Smoke test failed.",
      paste0("Error: ", smoke_msg)
    ))
    if (grepl("vector memory limit", smoke_msg, fixed = TRUE)) {
      MI_DEBUG_PRINTFAG <- TRUE
    }
    if (subset_data_saved) {
      subset_data <- readRDS(subset_data_path)
    }
  }
}

```

```

}

if (MI_DEBUG_PRINTFLAG) {
  debug_path <- results_path("mice_debug_print.txt")
  dbg_con <- file(debug_path, open = "wt")
  sink(dbg_con, type = "output")
  message("Debug MI: running m=1, maxit=1 with printFlag=TRUE.")
  tryCatch(
    run_mice_call(m_val = 1L, maxit = 1L, seed_val = MI_SEED, print_flag = TRUE),
    error = function(e) e
  )
  sink(type = "output")
  close(dbg_con)
}
stop("MICE smoke test failed; see ", mi_smoke_log_path, " and Results/mice_debug_print.txt.")
} else {
  message("MICE smoke test succeeded (m=1, maxit=1).")
}
}

get_vcells_stats <- function() {
  g <- gc()
  cn <- colnames(g)
  mb_cols <- cn[grep("\\\\(Mb\\\\)", cn)]
  used_mb_col <- if (length(mb_cols)) mb_cols[1] else NA_character_
  limit_col <- cn[grep("limit", cn, ignore.case = TRUE)]
  limit_col <- if (length(limit_col)) limit_col[1] else NA_character_
  trig_col <- cn[grep("trigger", cn, ignore.case = TRUE)]
  trig_col <- if (length(trig_col)) trig_col[1] else NA_character_
  max_col <- cn[grep("max", cn, ignore.case = TRUE)]
  max_col <- if (length(max_col)) max_col[1] else NA_character_

  used_mb <- if (!is.na(used_mb_col)) as.numeric(g["Vcells", used_mb_col]) else NA_real_
  limit_mb <- if (!is.na(limit_col)) as.numeric(g["Vcells", limit_col]) else NA_real_
  trig_mb <- if (!is.na(trig_col)) as.numeric(g["Vcells", trig_col]) else NA_real_
  max_mb <- if (!is.na(max_col)) as.numeric(g["Vcells", max_col]) else NA_real_
}

```

```

if (!is.finite(limit_mb) || limit_mb <= 0) {
  mem_lim <- tryCatch(mem.maxVSize(), error = function(e) NA_real_)
  if (is.finite(mem_lim) && mem_lim > 0) {
    limit_mb <- if (mem_lim < 1e8) mem_lim else mem_lim / 1024^2
  }
}

data.frame(
  gc_vcells_used_mb = used_mb,
  gc_vcells_limit_mb = limit_mb,
  gc_vcells_frac = ifelse(is.finite(limit_mb) && limit_mb > 0, used_mb / limit_mb, NA_real_),
  gc_vcells_trigger_mb = trig_mb,
  gc_vcells_max_used_mb = max_mb,
  stringsAsFactors = FALSE
)
}

set.seed(MI_SEED)
run_mice_single <- function(m_val) {
  runtime_logger(
    "mice_imputation",
    run_mice_call(m_val = m_val, maxit = MAXIT_RUN, seed_val = MI_SEED, print_flag = FALSE),
    notes = paste0("m=", m_val, "; maxit=", MAXIT_RUN)
  )
}

mc_progress <- list()
sentinel_specs <- list(
  list(name = "abg_imv", outcome = "imv_proc", treat = "has_abg", co2_var = "paco2",
       low = ABG_CO2_LOW, high = ABG_CO2_HIGH),
  list(name = "vbg_imv", outcome = "imv_proc", treat = "has_vbg", co2_var = "vbg_co2",
       low = VBG_CO2_LOW, high = VBG_CO2_HIGH)
)

mcerr_ratio_for_spec <- function(imp_obj, spec) {
  fits <- lapply(seq_len(imp_obj$m), function(i) {

```

```

d <- mice::complete(imp_obj, action = i)
d <- d[, c(spec$outcome, spec$treat, spec$co2_var), drop = FALSE]
d <- d[d[[spec$treat]] == 1 & is.finite(d[[spec$co2_var]]), , drop = FALSE]
if (nrow(d) == 0L) return(NULL)
d$co2_cat <- make_co2_cat3(d[[spec$co2_var]], spec$low, spec$high)
d$co2_cat <- stats::relevel(base::droplevels(d$co2_cat), ref = "Normal")
if (dplyr::n_distinct(d[[spec$outcome]]) < 2L) return(NULL)
fit <- tryCatch(
  stats::glm(stats::reformulate("co2_cat", response = spec$outcome),
             data = d, family = binomial(), x = FALSE, y = FALSE, model = FALSE),
  error = function(e) NULL
)
if (is.null(fit)) return(NULL)
list(coef = stats::coef(fit), vcov = stats::vcov(fit))
})
fits <- fits[!vapply(fits, is.null, logical(1))]
if (length(fits) < 2L) return(NA_real_)
results <- lapply(fits, `[[`, "coef")
variances <- lapply(fits, `[[`, "vcov")
pooled <- mitools::MIcombine(results = results, variances = variances)
est <- as.numeric(stats::coef(pooled))
names(est) <- names(stats::coef(pooled))
se <- sqrt(diag(pooled$variance))
names(se) <- names(stats::coef(pooled))
coef_mat <- do.call(cbind, lapply(results, function(v) v[names(est)]))
B <- apply(coef_mat, 1, stats::var, na.rm = TRUE)
mcerr <- sqrt(B / length(results))
ratio <- mcerr / se
idx <- grep("co2_cat", names(ratio))
if (!any(idx)) return(max(ratio, na.rm = TRUE))
max(ratio[idx], na.rm = TRUE)
}

run_mice_batched <- function(m_total, m_batch_start, maxit, base_seed) {
  imp_acc <- NULL
  m_done <- 0L

```

```

batch_attempt_idx <- 0L
m_batch <- m_batch_start
batch_log <- list()
logged_events_acc <- list()

while (m_done < m_total) {
  if (MI_GC_EVERY_BATCH) invisible(gc())
  mem_stats <- get_vcells_stats()
  if (MI_PREEMPTIVE_BATCH_REDUCE &&
      is.finite(mem_stats$gc_vcells_frac) &&
      mem_stats$gc_vcells_frac > MI_VCELLS_FRAC_THRESHOLD &&
      m_batch > MI_BATCH_MIN) {
    m_batch <- max(MI_BATCH_MIN, floor(m_batch / 2))
    message("Preemptively reducing MI batch size to ", m_batch,
           " (Vcells pressure: ", round(mem_stats$gc_vcells_frac, 2), ".)")
  }

  m_b <- min(m_batch, m_total - m_done)
  batch_attempt_idx <- batch_attempt_idx + 1L
  seed_b <- base_seed + batch_attempt_idx * MI_BATCH_SEED_STRIDE
  message("MICE batch ", batch_attempt_idx, " (m=", m_b, ", seed=", seed_b, ")")

  t0 <- Sys.time()
  mem_pre <- get_vcells_stats()
  cap <- capture_warnings(
    tryCatch(
      runtime_logger(
        paste0("mice_batch_", batch_attempt_idx),
        mice::mice(
          data = mi_df_run,
          m = m_b,
          maxit = maxit,
          predictorMatrix = pred,
          method = meth,
          printFlag = FALSE,
          seed = seed_b

```

```

),
notes = paste0("batch=", batch_attempt_idx, "; m=", m_b, "; maxit=", maxit)
),
error = function(e) e
),
context = make_context(
  stage = "MI",
  component = "mice_batch",
  analysis_variant = "mi",
  model_type = "mice",
  group = NA_character_,
  outcome = NA_character_,
  imputation = NA_integer_,
  batch = batch_attempt_idx
)
)
append_warnings(cap$warnings)

imp_b <- cap$value
if (inherits(imp_b, "error")) {
  err_msg <- conditionMessage(imp_b)
  message("MICE batch ", batch_attempt_idx, " failed: ", err_msg)
  mem_stats <- get_vcells_stats()
  batch_log[[batch_attempt_idx]] <- data.frame(
    batch = batch_attempt_idx,
    m_batch = m_b,
    seed = seed_b,
    ok = FALSE,
    error_message = err_msg,
    seconds = as.numeric(difftime(Sys.time(), t0, units = "secs")),
    gc_vcells_used_mb_pre = mem_pre$gc_vcells_used_mb,
    gc_vcells_limit_mb_pre = mem_pre$gc_vcells_limit_mb,
    gc_vcells_frac_pre = mem_pre$gc_vcells_frac,
    gc_vcells_used_mb_post = mem_stats$gc_vcells_used_mb,
    gc_vcells_limit_mb_post = mem_stats$gc_vcells_limit_mb,
    gc_vcells_frac_post = mem_stats$gc_vcells_frac,

```

```

    stringsAsFactors = FALSE
)
if (grepl("vector memory limit", err_msg, fixed = TRUE) && m_batch > MI_BATCH_MIN) {
  m_batch <- max(MI_BATCH_MIN, floor(m_batch / 2))
  message("Reducing MI batch size to ", m_batch, " and retrying.")
  invisible(gc())
  next
}
write_csv_safely(dplyr::bind_rows(batch_log), results_path("mice_batches_log.csv"), row_names = FALSE)
stop("MICE batch ", batch_attempt_idx, " failed; see log: ", results_path("mice_batches_log.csv"))
}

if (is.null(imp_acc)) {
  imp_acc <- imp_b
} else {
  imp_acc <- mice::ibind(imp_acc, imp_b)
}

le_b <- imp_b$loggedEvents
le_b <- if (is.null(le_b)) data.frame() else as.data.frame(le_b)
if (NROW(le_b) > 0) {
  le_b <- le_b |>
    dplyr::mutate(
      batch = batch_attempt_idx,
      seed = seed_b,
      m_global_start = m_done + 1L
    )
  logged_events_acc[[length(logged_events_acc) + 1L]] <- le_b
}

m_done <- imp_acc$m
mem_stats <- get_vcells_stats()
batch_log[[batch_attempt_idx]] <- data.frame(
  batch = batch_attempt_idx,
  m_batch = m_b,
  seed = seed_b,

```

```

ok = TRUE,
error_message = NA_character_,
seconds = as.numeric(difftime(Sys.time(), t0, units = "secs")),
gc_vcells_used_mb_pre = mem_pre$gc_vcells_used_mb,
gc_vcells_limit_mb_pre = mem_pre$gc_vcells_limit_mb,
gc_vcells_frac_pre = mem_pre$gc_vcells_frac,
gc_vcells_used_mb_post = mem_stats$gc_vcells_used_mb,
gc_vcells_limit_mb_post = mem_stats$gc_vcells_limit_mb,
gc_vcells_frac_post = mem_stats$gc_vcells_frac,
stringsAsFactors = FALSE
)
rm(imp_b)
if (MI_GC_EVERY_BATCH) invisible(gc())

if (ALLOW_M_IMP_EARLY_STOP &&
  m_done >= M_IMP_MIN &&
  (m_done %% M_IMP_STEP == 0 || m_done == m_total)) {
  ratios <- vapply(sentinel_specs, function(s) mcerr_ratio_for_spec(imp_acc, s), numeric(1))
  mc_progress[[length(mc_progress) + 1L]] <- data.frame(
    m = m_done,
    abg_ratio = ratios[["abg_imv"]],
    vbg_ratio = ratios[["vbg_imv"]],
    max_ratio = max(ratios, na.rm = TRUE),
    stringsAsFactors = FALSE
  )
  write_csv_safely(dplyr::bind_rows(mc_progress),
                  results_path("mi_mcerr_progress.csv"),
                  row_names = FALSE)
  if (all(is.finite(ratios)) && max(ratios, na.rm = TRUE) <= MCERR_RATIO_TARGET) {
    message("MC error criterion met at m=", m_done,
           " (max MCerr/SE=", round(max(ratios, na.rm = TRUE), 3), "). Stopping early.")
    break
  }
}
}

```

```

write_csv_safely(dplyr::bind_rows(batch_log), results_path("mice_batches_log.csv"), row_names = FALSE)
log_events_raw_batched <- dplyr::bind_rows(logged_events_acc)
attr(imp_acc, "logged_events_batched") <- log_events_raw_batched
imp_acc
}

imp <- NULL
use_batched <- isTRUE(FORCE_MI_BATCHED)
if (!use_batched) {
  cap_mice <- capture_warnings(
    tryCatch(run_mice_single(M_IMP_RUN), error = function(e) e),
    context = make_context(
      stage = "MI",
      component = "mice",
      analysis_variant = "mi",
      model_type = "mice",
      group = NA_character_,
      outcome = NA_character_,
      imputation = NA_integer_,
      batch = NA_integer_
    )
  )
  append_warnings(cap_mice$warnings)
  imp <- cap_mice$value
  if (inherits(imp, "error") && grepl("vector memory limit", conditionMessage(imp), fixed = TRUE)) {
    message("MICE memory limit hit; switching to batched mode.")
    use_batched <- TRUE
  }
}
if (use_batched) {
  message("MICE: running in batches (start=", MI_BATCH_START, ").")
  imp <- run_mice_batched(M_IMP_RUN, MI_BATCH_START, MAXIT_RUN, MI_SEED)
} else {
  write_csv_safely(data.frame(), results_path("mice_batches_log.csv"), row_names = FALSE)
}

```

```

if (inherits(imp, "error")) stop(imp)
stopifnot(inherits(imp, "mids"))
if (ALLOW_M_IMP_EARLY_STOP && imp$m < M_IMP_RUN) {
  message("Early stop: stopping at m=", imp$m, " (target ", M_IMP_RUN, ").")
  M_IMP_RUN <- imp$m
} else {
  stopifnot(imp$m == M_IMP_RUN)
}
if (M_IMP != M_IMP_RUN) M_IMP <- M_IMP_RUN
if (MAXIT_MI != MAXIT_RUN) MAXIT_MI <- MAXIT_RUN
write_csv_safely(dplyr::bind_rows(mc_progress), results_path("mi_mcerr_progress.csv"), row_names = FALSE)
saveRDS(imp, file = mi_mids_file)

# Save MICE spec for reproducibility
saveRDS(
  list(method = imp$method, predictorMatrix = imp$predictorMatrix),
  results_path("mice_spec.rds")
)
if (use_batched) {
  message("Multiple imputation was run in batches and combined via mice::ibind().")
}
if (subset_data_saved) {
  subset_data <- readRDS(subset_data_path)
}

# Logged events: raw + summary (by dep/out/meth)
log_events_raw <- as.data.frame(imp$loggedEvents)
log_events_batched <- as.data.frame(attr(imp, "logged_events_batched"))
log_events_raw <- dplyr::bind_rows(log_events_raw, log_events_batched)

if (nrow(log_events_raw)) {
  write_csv_safely(log_events_raw, results_path("mice_logged_events_raw.csv"), row_names = FALSE,
                  required_cols = c("dep", "out", "meth"))
  log_events_summary <- log_events_raw |>
    dplyr::count(dep, out, meth, name = "n") |>
    dplyr::mutate(variable = dep) |>

```

```

dplyr::arrange(dplyr::desc(n)) |>
  dplyr::mutate(pct = n / sum(n))
  write_csv_safely(log_events_summary, results_path("mice_logged_events_summary.csv"), row_names = FALSE,
                  required_cols = c("variable", "n", "pct"))
} else {
  log_events_raw_empty <- data.frame(
    dep = character(), out = character(), meth = character(),
    stringsAsFactors = FALSE
  )
  log_events_summary_empty <- data.frame(
    variable = character(), n = integer(), pct = numeric(),
    stringsAsFactors = FALSE
  )
  write_csv_safely(log_events_raw_empty, results_path("mice_logged_events_raw.csv"), row_names = FALSE,
                  required_cols = c("dep", "out", "meth"))
  write_csv_safely(log_events_summary_empty, results_path("mice_logged_events_summary.csv"), row_names = FALSE,
                  required_cols = c("variable", "n", "pct"))
  log_events_summary <- log_events_summary_empty
}
if (nrow(mi_info_log)) {
  warns_events <- mi_info_log |>
    dplyr::filter(stage == "MI", component %in% c("mice", "mice_batch"))
  if (nrow(warns_events) && nrow(log_events_raw) == 0L) {
    warning("Mismatch: main MI run reported logged events but loggedEvents table is empty; ",
            "check batch capture and loggedEvents exports.", call. = FALSE)
  }
}

# Chain diagnostics (lightweight; no complete("long"))
chain_diag <- data.frame()
chain_diag_stats <- list(
  n_imputed_vars = 0L,
  n_with_chainMean = 0L,
  n_with_drift_tail = 0L,
  drift_tail_na_frac = NA_real_,
  tail_window_na_mean = NA_real_

```

```

)
stopifnot(!is.null(imp$method))
impute_vars <- names(imp$method)[imp$method != ""]
impute_vars <- intersect(impute_vars, names(imp$data))
if (length(impute_vars)) {
  impute_vars <- impute_vars[vapply(imp$data[impute_vars], function(x) any(is.na(x)), logical(1))]
}
stopifnot(!is.null(imp$chainMean))
{
  cm <- imp$chainMean
  dims <- dim(cm)
  dn <- dimnames(cm)
  iter_candidates <- unique(c(imp$iteration, MAXIT_MI, MAXIT_MI + 1L))
  iter_candidates <- iter_candidates[is.finite(iter_candidates) & iter_candidates > 0]
  imp_m <- imp$m

  extract_chain_mean <- function(cm_obj, dims_obj, dn_obj, imp_m_val, iter_cand) {
    mean_chain <- NULL
    var_names <- NULL
    iter_dim <- NA_integer_
    var_dim <- NA_integer_
    m_dim <- NA_integer_
    if (!is.null(dims_obj) && length(dims_obj) == 2) {
      iter_dim <- which(dims_obj %in% iter_cand)[1]
      if (length(iter_dim) == 0) iter_dim <- 1L
      var_dim <- setdiff(seq_along(dims_obj), iter_dim)[1]
      mean_chain <- cm_obj
      if (iter_dim == 2L) mean_chain <- t(mean_chain)
    } else if (!is.null(dims_obj) && length(dims_obj) == 3) {
      m_dim <- which(dims_obj == imp_m_val)
      if (length(m_dim)) {
        m_dim <- m_dim[1]
      } else {
        m_dim <- 3L
      }
      iter_dim <- setdiff(which(dims_obj %in% iter_cand), m_dim)[1]
    }
  }
}

```

```

if (length(iter_dim) == 0) iter_dim <- setdiff(1:3, c(m_dim, NA_integer_))[1]
var_dim <- setdiff(1:3, c(iter_dim, m_dim))[1]
if (all(is.finite(c(iter_dim, var_dim, m_dim)))) {
  cm_std <- aperm(cm_obj, c(iter_dim, var_dim, m_dim))
  mean_chain <- apply(cm_std, c(1, 2), mean, na.rm = TRUE)
  if (!is.null(dimnames(cm_std))) {
    if (!is.null(dimnames(cm_std)[[2]])) {
      var_names <- dimnames(cm_std)[[2]]
    }
    if (!is.null(dimnames(cm_std)[[1]])) {
      rownames(mean_chain) <- dimnames(cm_std)[[1]]
    }
  }
}
if (is.null(mean_chain)) {
  return(list(mean_chain = NULL, var_names = NULL, iter_dim = iter_dim, var_dim = var_dim, m_dim = m_dim))
}
if (is.null(var_names) && !is.null(dn_obj) && length(dn_obj) >= var_dim) {
  var_names <- dn_obj[[var_dim]]
}
if (is.null(var_names)) {
  var_names <- colnames(mean_chain)
}
list(mean_chain = mean_chain, var_names = var_names, iter_dim = iter_dim, var_dim = var_dim, m_dim = m_dim)
}

res_chain <- extract_chain_mean(cm, dims, dn, imp_m, iter_candidates)
mean_chain <- res_chain$mean_chain
var_names <- res_chain$var_names
if (is.null(mean_chain)) {
  stop("Chain diagnostics: unable to construct mean_chain from imp$chainMean.")
}
{
  iter_idx <- seq_len(nrow(mean_chain))

```

```

numeric_names <- !is.null(var_names) && all(grepl("^\\d+$", var_names))
if (is.null(var_names) || numeric_names) {
  if (length(impute_vars) && length(impute_vars) == ncol(mean_chain)) {
    var_names <- impute_vars
    numeric_names <- FALSE
  }
}
if (is.null(var_names)) {
  stop("Chain diagnostics: variable names missing and could not be matched to imp$method.")
}
colnames(mean_chain) <- var_names

if (numeric_names) {
  warn_df <- data.frame(
    time = as.character(Sys.time()),
    stage = "MI",
    component = "chain_diagnostics",
    analysis_variant = NA_character_,
    model_type = NA_character_,
    group = NA_character_,
    outcome = NA_character_,
    imputation = NA_integer_,
    batch = NA_integer_,
    message = "Chain diagnostics variable name mapping failed; using numeric/fallback names.",
    stringsAsFactors = FALSE
  )
  append_warnings(warn_df)
}

vars_imputed <- impute_vars
chain_diag_stats$n_imputed_vars <- length(vars_imputed)
keep_vars <- intersect(vars_imputed, var_names)
missing_in_chain <- setdiff(vars_imputed, var_names)
if (length(missing_in_chain)) {
  missing_df <- data.frame(
    variable = missing_in_chain,

```

```

        stringsAsFactors = FALSE
    )
    write_csv_safely(missing_df, results_path("mice_chain_diagnostics_missing_vars.csv"), row_names = FALSE)
}
if (length(vars_imputed) && length(keep_vars) == 0L) {
  warn_df <- data.frame(
    time = as.character(Sys.time()),
    stage = "MI",
    component = "chain_diagnostics",
    analysis_variant = NA_character_,
    model_type = NA_character_,
    group = NA_character_,
    outcome = NA_character_,
    imputation = NA_integer_,
    batch = NA_integer_,
    message = "Chain diagnostics: no imputed variables matched chainMean names; skipping drift metrics.",
    stringsAsFactors = FALSE
  )
  append_warnings(warn_df)
  mean_chain <- mean_chain[, 0, drop = FALSE]
  var_names <- character()
} else if (length(keep_vars)) {
  mean_chain <- mean_chain[, keep_vars, drop = FALSE]
  var_names <- keep_vars
}

safe_sd <- function(x) if (sum(is.finite(x)) < 2) NA_real_ else stats::sd(x, na.rm = TRUE)
safe_maxdiff <- function(x) {
  x <- x[is.finite(x)]
  if (length(x) < 2) return(NA_real_)
  max(abs(diff(x)))
}
safe_slope <- function(x, iter) {
  ok <- is.finite(x)
  if (sum(ok) < 2) return(NA_real_)
  coef(stats::lm(x[ok] ~ iter[ok]))[2]
}

```

```

}

tail_finite <- function(x, k) tail(x[is.finite(x)], k)
safe_maxdiff_tail <- function(x, k) {
  xf <- tail_finite(x, k)
  if (length(xf) < 2) return(NA_real_)
  max(abs(diff(xf)))
}

n_vars <- ncol(mean_chain)
n_finite <- integer(n_vars)
drift_all <- numeric(n_vars)
drift_tail <- numeric(n_vars)
tail_n_finite <- integer(n_vars)
tail_window_na_frac <- numeric(n_vars)
slope <- numeric(n_vars)
sd_chain <- numeric(n_vars)
overall_reason <- character(n_vars)
tail_reason <- character(n_vars)
flag_reason <- character(n_vars)
flag <- logical(n_vars)
diagnostic_available <- logical(n_vars)
sd_obs <- rep(NA_real_, n_vars)

chain_src_df <- mi_df_run
stopifnot(is.data.frame(chain_src_df))
sd_obs <- vapply(var_names, function(v) {
  stopifnot(v %in% names(chain_src_df))
  x <- chain_src_df[[v]]
  if (inherits(x, "haven_labelled")) x <- suppressWarnings(as.numeric(x))
  if (!is.numeric(x)) return(NA_real_)
  safe_sd(x)
}, numeric(1))

for (j in seq_len(n_vars)) {
  x <- mean_chain[, j]
  n_finite[j] <- sum(is.finite(x))
}

```

```

drift_all[j] <- safe_maxdiff(x)
slope[j] <- safe_slope(x, iter_idx)
sd_chain[j] <- safe_sd(x)
tail_vals <- tail_finite(x, TAIL_NFINITE)
tail_n_finite[j] <- length(tail_vals)
drift_tail[j] <- safe_maxdiff_tail(x, TAIL_NFINITE)
tail_window <- tail(x, min(TAIL_WINDOW_ITERS, length(x)))
tail_window_na_frac[j] <- mean(!is.finite(tail_window))
overall_reason[j] <- if (n_finite[j] >= 2) "ok" else "insufficient_finite_overall"
tail_reason[j] <- if (tail_n_finite[j] >= 2) "ok" else "insufficient_finite_tail"
}

drift_all_scaled <- drift_all / sd_obs
drift_tail_scaled <- drift_tail / sd_obs
slope_scaled <- slope / sd_obs

for (j in seq_len(n_vars)) {
  if (!is.finite(sd_obs[j]) || sd_obs[j] <= 0) {
    flag[j] <- NA
    flag_reason[j] <- "missing_scale"
    diagnostic_available[j] <- FALSE
    next
  }
  if (tail_n_finite[j] < 2) {
    flag[j] <- NA
    flag_reason[j] <- "tail_insufficient"
    diagnostic_available[j] <- FALSE
    next
  }
  if (!is.finite(drift_tail_scaled[j]) || !is.finite(slope_scaled[j])) {
    flag[j] <- NA
    flag_reason[j] <- "insufficient_data"
    diagnostic_available[j] <- FALSE
    next
  }
  flag_tail <- drift_tail_scaled[j] > 0.01
}

```

```

flag_slope <- abs(slope_scaled[j]) > 0.001
flag[j] <- flag_tail | flag_slope
flag_reason[j] <- if (flag_tail & flag_slope) {
  "both"
} else if (flag_tail) {
  "tail_drift"
} else if (flag_slope) {
  "slope"
} else {
  "none"
}
diagnostic_available[j] <- TRUE
}

chain_diag <- data.frame(
  variable = var_names,
  method = imp$method[var_names],
  n_finite = n_finite,
  drift_all = drift_all,
  drift_tail = drift_tail,
  drift_all_scaled = drift_all_scaled,
  drift_tail_scaled = drift_tail_scaled,
  tail_n_finite = tail_n_finite,
  tail_window_na_frac = tail_window_na_frac,
  slope = slope,
  slope_scaled = slope_scaled,
  sd_chain = sd_chain,
  sd_obs = sd_obs,
  overall_reason = overall_reason,
  tail_reason = tail_reason,
  flag_reason = flag_reason,
  diagnostic_available = diagnostic_available,
  flag = flag,
  stringsAsFactors = FALSE
)

```

```

chain_diag_stats$n_with_chainMean <- nrow(chain_diag)
chain_diag_stats$n_with_drift_tail <- sum(is.finite(chain_diag$drift_tail))
if (chain_diag_stats$n_with_chainMean > 0) {
  chain_diag_stats$drift_tail_na_frac <- 1 - (chain_diag_stats$n_with_drift_tail /
                                                 chain_diag_stats$n_with_chainMean)
  chain_diag_stats$tail_window_na_mean <- mean(chain_diag$tail_window_na_frac, na.rm = TRUE)
}
write_csv_safely(chain_diag, results_path("mice_chain_diagnostics.csv"), row.names = FALSE)
if (any(isTRUE(chain_diag$flag), na.rm = TRUE)) {
  frac_flag <- mean(chain_diag$flag, na.rm = TRUE)
  if (is.finite(frac_flag) && frac_flag < 0.05) {
    message("Chain diagnostics: low drift flags (", round(frac_flag * 100, 1),
           "%). MAXIT_MI may be reduced safely (consider 10 or 5).")
  }
}
}

# quick sanity: these must exist and be numeric in completed data
imp_n <- imp$m
get_imp_stats <- list(count = 0L, seconds = 0)
get_imp <- function(i, imp_obj = imp) {
  t0 <- Sys.time()
  d <- normalize_types(mice:::complete(imp_obj, action = i), levels_ref)
  get_imp_stats$count <-> get_imp_stats$count + 1L
  get_imp_stats$seconds <-> get_imp_stats$seconds +
    as.numeric(difftime(Sys.time(), t0, units = "secs"))
  d
}
d1 <- get_imp(1)
stopifnot(all(c("paco2", "vbg_co2") %in% names(d1)))
stopifnot(is.numeric(d1$paco2), is.numeric(d1$vbg_co2))

# post-MICE sanity: no remaining NA in covars_gbm
covars_check <- intersect(covars_gbm, names(d1))
na_counts <- vapply(d1[, covars_check, drop = FALSE], function(x) sum(is.na(x)), numeric(1))

```

Monte Carlo error vs SE (diagnostic only)

| Term | Estimate | SE | MC error | MC error / SE | 2.5% | 97.5% |
|-------------------------|----------|-------|----------|---------------|-------------|---------------|
| (Intercept) | -3.821 | 0.189 | 0.014 | 0.074 | -4.19136254 | -3.4497745159 |
| has_abg | 2.166 | 0.053 | 0.001 | 0.016 | 2.06223732 | 2.2704301035 |
| age_at_encounter | -0.003 | 0.001 | 0.000 | 0.014 | -0.00596037 | -0.0009504728 |
| curr_bmi | -0.009 | 0.005 | 0.000 | 0.086 | -0.01787779 | 0.0001613787 |
| sexMale | 0.236 | 0.043 | 0.001 | 0.014 | 0.15097989 | 0.3212478672 |
| encounter_typeInpatient | 1.044 | 0.068 | 0.000 | 0.005 | 0.91072250 | 1.1766322121 |

```

na_counts <- na_counts[na_counts > 0]
if (length(na_counts)) {
  message("Post-MICE NA counts (covars_gbm): ",
         paste(names(na_counts), na_counts, collapse = ", "))
  ev_sum <- summarize_logged_events(imp)
  if (nrow(ev_sum)) {
    ev_sub <- ev_sum[ev_sum$variable %in% names(na_counts), , drop = FALSE]
    if (nrow(ev_sub)) {
      print(utils::head(ev_sub, 10))
    } else {
      message("No loggedEvents entries for covars_gbm with NA.")
    }
  } else {
    message("No loggedEvents recorded.")
  }
  stop("Post-MICE check failed: remaining NA in covars_gbm. See loggedEvents summary above.")
}

```

MC error diagnostic model: Diagnostic model: imv_proc ~ has_abg + age_at_encounter + curr_bmi + sex + encounter_type (unweighted).

3.3 10) Refit propensity models within each imputation

MI propensity scores use logistic regression with restricted cubic splines (`rms::rcs`, 4 knots by default) for continuous covariates; the same covariate set used in non-MI models is reused here (`covars_ps`). IPSW truncation rules are unchanged.

Note: the MI computations below run in a single pass per imputation (weights, balance, cat3, spline). Subsequent MI sections reuse those outputs and will stop if they are missing.

3.3.1 FAIL-FAST CHECKS

| | used (Mb) | gc trigger (Mb) | limit (Mb) | max used (Mb) |
|--------|-----------|-----------------|------------|-------------------|
| Ncells | 6628748 | 354.1 | 12180610 | 650.6 |
| Vcells | 610437722 | 4657.3 | 1053929965 | 8040.9 |
| | | | NA | 12180610 650.6 |
| | | | 16384 | 1053929965 8040.9 |

3.3.2 10.1 ABG propensity (`has_abg`)

```
stopifnot(exists("mi_single_pass_done"), isTRUE(mi_single_pass_done))
stopifnot(exists("W_abg_list"))
message("ABG MI weights were computed in the single-pass MI loop above.")
```

Table 9: ABG weight diagnostics (MI; median across imputations)

| n | min | p99.99. | max | ess |
|------|-------|---------|-------|----------|
| 9362 | 0.384 | 3.819 | 3.821 | 6847.997 |

3.3.3 10.2 Balance diagnostics across imputations

```
stopifnot(exists("target_balance_table"))

# Use a fixed x-axis max so ABG/VBG balance plots are directly comparable.
```

```

stopifnot(exists("bal_imp_abg"))

bal_abg_pooled <- bal_imp_abg$bal_long |>
  mutate(label = ifelse(is.na(level), variable, paste0(variable, ":", level))) |>
  group_by(label) |>
  summarise(
    pre_med = median(abs(smd_pre), na.rm = TRUE),
    post_med = median(abs(smd_post), na.rm = TRUE),
    pre_mean = mean(abs(smd_pre), na.rm = TRUE),
    post_mean = mean(abs(smd_post), na.rm = TRUE),
    post_max = max(abs(smd_post), na.rm = TRUE),
    .groups = "drop"
  )

bal_abg_plot <- bal_abg_pooled |>
  mutate(label = factor(label, levels = label[order(post_med, decreasing = TRUE)])) |>
  pivot_longer(c(pre_med, post_med), names_to = "type", values_to = "smd") |>
  mutate(type = recode(type, pre_med = "Pre", post_med = "Post"))

p_abg <- ggplot(bal_abg_plot, aes(x = smd, y = label, shape = type)) +
  geom_vline(xintercept = 0.1, linetype = 2, linewidth = 0.3) +
  geom_point(size = 1.2) +
  labs(title = "MI target balance (ABG): pooled |SMD|", x = "|Target SMD|", y = NULL, shape = "Stage") +
  scale_x_continuous(limits = c(0, BAL_XLIM_MAX),
                     expand = expansion(mult = c(0, 0.02))) +
  theme_minimal(base_size = 10)
save_diag_plot(p_abg, results_path("figs", "diag-mi-balance-pooled-abg.png"), width = 7, height = 6)

knitr::kable(bal_imp_abg$worst_rows_overall, caption = "ABG: worst target SMD rows across imputations (top 10)")

```

Table 10: ABG: worst target SMD rows across imputations (top 10)

| variable | level | type | smd_pre | smd_post | group | imp | abs_post |
|----------------|-----------|--------|------------|------------|-------|-----|-----------|
| encounter_type | Inpatient | factor | 0.3704053 | 0.1001128 | ABG | 45 | 0.1001128 |
| encounter_type | Emergency | factor | -0.3704053 | -0.1001128 | ABG | 45 | 0.1001128 |

| variable | level | type | smd_pre | smd_post | group | imp | abs_post |
|----------------|-----------|--------|------------|------------|-------|-----|-----------|
| encounter_type | Inpatient | factor | 0.3704053 | 0.0979839 | ABG | 40 | 0.0979839 |
| encounter_type | Emergency | factor | -0.3704053 | -0.0979839 | ABG | 40 | 0.0979839 |
| encounter_type | Inpatient | factor | 0.3704053 | 0.0959636 | ABG | 38 | 0.0959636 |
| encounter_type | Emergency | factor | -0.3704053 | -0.0959636 | ABG | 38 | 0.0959636 |
| encounter_type | Emergency | factor | -0.3704053 | -0.0952833 | ABG | 51 | 0.0952833 |
| encounter_type | Inpatient | factor | 0.3704053 | 0.0952833 | ABG | 51 | 0.0952833 |
| encounter_type | Emergency | factor | -0.3704053 | -0.0952657 | ABG | 8 | 0.0952657 |
| encounter_type | Inpatient | factor | 0.3704053 | 0.0952657 | ABG | 8 | 0.0952657 |

```
abg_imp_summary <- bal_imp_abg$bal_imp_summary |>
  summarise(
    med = median(max_abs_post, na.rm = TRUE),
    p90 = quantile(max_abs_post, 0.9, na.rm = TRUE),
    max = max(max_abs_post, na.rm = TRUE)
  )
knitr::kable(abg_imp_summary, caption = "ABG: max |Target SMD| summary across imputations")
```

Table 11: ABG: max |Target SMD| summary across imputations

| med | p90 | max |
|-----------|-----------|-----------|
| 0.0911384 | 0.0940709 | 0.1001128 |

3.3.4 10.3 VBG propensity (has_vbg)

```
stopifnot(exists("mi_single_pass_done"), isTRUE(mi_single_pass_done))
stopifnot(exists("W_vbg_list"))
message("VBG MI weights were computed in the single-pass MI loop above.")
```

Table 12: VBG weight diagnostics (MI; median across imputations)

| | n | min | p99.99. | max | ess |
|--|------|-------|---------|-------|---------|
| | 7460 | 0.308 | 3.892 | 3.895 | 5076.89 |

3.3.5 10.4 VBG balance

```

stopifnot(exists("bal_imp_vbg"))

bal_vbg_pooled <- bal_imp_vbg$bal_long |>
  mutate(label = ifelse(is.na(level), variable, paste0(variable, ":", level))) |>
  group_by(label) |>
  summarise(
    pre_med = median(abs(smd_pre), na.rm = TRUE),
    post_med = median(abs(smd_post), na.rm = TRUE),
    pre_mean = mean(abs(smd_pre), na.rm = TRUE),
    post_mean = mean(abs(smd_post), na.rm = TRUE),
    post_max = max(abs(smd_post), na.rm = TRUE),
    .groups = "drop"
  )

bal_vbg_plot <- bal_vbg_pooled |>
  mutate(label = factor(label, levels = label[order(post_med, decreasing = TRUE)])) |>
  pivot_longer(c(pre_med, post_med), names_to = "type", values_to = "smd") |>
  mutate(type = recode(type, pre_med = "Pre", post_med = "Post"))

p_vbg <- ggplot(bal_vbg_plot, aes(x = smd, y = label, shape = type)) +
  geom_vline(xintercept = 0.1, linetype = 2, linewidth = 0.3) +
  geom_point(size = 1.2) +
  labs(title = "MI target balance (VBG): pooled |SMD|", x = "|Target SMD|", y = NULL, shape = "Stage") +
  scale_x_continuous(limits = c(0, BAL_XLIM_MAX),
                     expand = expansion(mult = c(0, 0.02))) +
  theme_minimal(base_size = 10)
save_diag_plot(p_vbg, results_path("figs", "diag-mi-balance-pooled-vbg.png"), width = 7, height = 6)

```

```
knitr::kable(bal_imp_vbg$worst_rows_overall, caption = "VBG: worst target SMD rows across imputations (top 10)")
```

Table 13: VBG: worst target SMD rows across imputations (top 10)

| variable | level | type | smd_pre | smd_post | group | imp | abs_post |
|----------------|-----------|--------|------------|------------|-------|-----|-----------|
| encounter_type | Emergency | factor | -0.0375275 | -0.0505432 | VBG | 26 | 0.0505432 |
| encounter_type | Inpatient | factor | 0.0375275 | 0.0505432 | VBG | 26 | 0.0505432 |
| encounter_type | Emergency | factor | -0.0375275 | -0.0493456 | VBG | 68 | 0.0493456 |
| encounter_type | Inpatient | factor | 0.0375275 | 0.0493456 | VBG | 68 | 0.0493456 |
| encounter_type | Emergency | factor | -0.0375275 | -0.0489478 | VBG | 36 | 0.0489478 |
| encounter_type | Inpatient | factor | 0.0375275 | 0.0489478 | VBG | 36 | 0.0489478 |
| encounter_type | Inpatient | factor | 0.0375275 | 0.0484923 | VBG | 9 | 0.0484923 |
| encounter_type | Emergency | factor | -0.0375275 | -0.0484923 | VBG | 9 | 0.0484923 |
| encounter_type | Inpatient | factor | 0.0375275 | 0.0479686 | VBG | 53 | 0.0479686 |
| encounter_type | Emergency | factor | -0.0375275 | -0.0479686 | VBG | 53 | 0.0479686 |

```
vbg_imp_summary <- bal_imp_vbg$bal_imp_summary |>
  summarise(
    med = median(max_abs_post, na.rm = TRUE),
    p90 = quantile(max_abs_post, 0.9, na.rm = TRUE),
    max = max(max_abs_post, na.rm = TRUE)
  )
knitr::kable(vbg_imp_summary, caption = "VBG: max |Target SMD| summary across imputations")
```

Table 14: VBG: max |Target SMD| summary across imputations

| med | p90 | max |
|----------|-----------|-----------|
| 0.044039 | 0.0474142 | 0.0505432 |

3.4 11) Weighted outcome models within each imputation + pooling

Within each imputation, fit covariate-adjusted CO2 spline outcome models **only in the measured cohort** (has_abg==1 for PaCO2; has_vbg==1 for VBG CO2), using IPSW weights to address nonrandom testing. Curves are pooled pointwise across imputations (Rubin's

rules on the log-OR scale) and displayed as odds ratios relative to CO2_ref at a reference covariate profile.

3.4.1 11.1 Helper: fit + extract log-OR and SE from svyglm

```
stopifnot(exists("fit_spline_imp"), exists("pool_spline_curve"),
          exists("pool_spline_coefs"), exists("pool_terms"))

stopifnot(exists("mi_single_pass_done"), isTRUE(mi_single_pass_done))
stopifnot(exists("abg_spline"), exists("abg_curves"), exists("abg_coefs"))
message("ABG MI spline results were computed in the single-pass MI loop above.")
```

3.4.2 11.3 VBG: MI pooled spline models (treated cohort only)

```
stopifnot(exists("mi_single_pass_done"), isTRUE(mi_single_pass_done))
stopifnot(exists("vbg_spline"), exists("vbg_curves"), exists("vbg_coefs"))
message("VBG MI spline results were computed in the single-pass MI loop above.")
```

3.5 12) Explainability on one representative imputation

MI propensity scores are estimated via logistic regression with splines; SHAP summaries are not computed for the MI PS model. A placeholder diagnostic file is written for audit completeness.

3.6 13) MI three-level PCO2 helpers and checks

```
stopifnot(exists("pool_terms"))
message("MI 3-level helpers defined in the single-pass MI section.")
```

3.7 14) MI + IPW three-level PCO2 (ABG & VBG)

3.7.1 14.1 ABG: MI + IPW, three-level PCO2 outcomes

```
stopifnot(exists("mi_single_pass_done"), isTRUE(mi_single_pass_done))
stopifnot(exists("abg_cat_results"))

mi_abg_cat_forms <- list(
  "MI IPW ABG 3-level: IMV ~ CO2 category + X"      = reformulate(c("co2_cat", adj_core), response = "imv_proc"),
  "MI IPW ABG 3-level: NIV ~ CO2 category + X"      = reformulate(c("co2_cat", adj_core), response = "niv_proc"),
  "MI IPW ABG 3-level: Death60d ~ CO2 category + X" = reformulate(c("co2_cat", adj_core), response = "death_60d"),
  "MI IPW ABG 3-level: HCRF ~ CO2 category + X"     = reformulate(c("co2_cat", adj_core), response = "hypercap_resp_failure")
)
register_model_diagrams(mi_abg_cat_forms)
message("ABG MI 3-level results were computed in the single-pass MI loop above.")
```

3.7.2 14.2 VBG: MI + IPW, three-level PCO2 outcomes

```
stopifnot(exists("mi_single_pass_done"), isTRUE(mi_single_pass_done))
stopifnot(exists("vbg_cat_results"))

mi_vbg_cat_forms <- list(
  "MI IPW VBG 3-level: IMV ~ CO2 category + X"      = reformulate(c("co2_cat", adj_core), response = "imv_proc"),
  "MI IPW VBG 3-level: NIV ~ CO2 category + X"      = reformulate(c("co2_cat", adj_core), response = "niv_proc"),
  "MI IPW VBG 3-level: Death60d ~ CO2 category + X" = reformulate(c("co2_cat", adj_core), response = "death_60d"),
  "MI IPW VBG 3-level: HCRF ~ CO2 category + X"     = reformulate(c("co2_cat", adj_core), response = "hypercap_resp_failure")
)
register_model_diagrams(mi_vbg_cat_forms)
message("VBG MI 3-level results were computed in the single-pass MI loop above.")
```

Table 3. MI-pooled IPW associations between CO category and outcomes (adjusted)

| Cohort | Outcome | Low vs normal OR (95% CI) | High vs normal OR (95% CI) |
|--------|----------------|---------------------------|----------------------------|
| ABG | IMV | 1.27 (1.11, 1.45) | 1.26 (1.10, 1.44) |
| ABG | NIV | 1.14 (0.91, 1.44) | 3.06 (2.52, 3.70) |
| ABG | Death (60d) | 1.63 (1.40, 1.90) | 1.30 (1.12, 1.52) |
| ABG | Hypercapnic RF | 1.03 (0.79, 1.35) | 7.66 (6.29, 9.32) |
| VBG | IMV | 1.50 (1.23, 1.82) | 1.54 (1.27, 1.87) |
| VBG | NIV | 0.91 (0.68, 1.24) | 2.83 (2.21, 3.63) |
| VBG | Death (60d) | 1.75 (1.42, 2.15) | 1.42 (1.16, 1.74) |
| VBG | Hypercapnic RF | 0.68 (0.46, 1.00) | 6.53 (5.08, 8.38) |

Weighted survey GLMs adjusted for baseline covariates; weights = MI-specific GLM (RCS) IPW; m = 80 imputations (seed 20251206); reference = Normal.

3.7.3 14.3 Table 3: MI-pooled IPW associations (3-level CO)

3.7.4 14.4 Summary: adjusted CO2-category associations across analysis tracks

Table 15: Adjusted odds ratios (low/high vs normal) across analysis tracks; n/events reflect model sample size (median across imputations for MI).

| method | group | outcome_label | n_model | events | Low vs normal | High vs normal |
|---------------------|-------|----------------|---------|--------|-------------------|-------------------|
| Unweighted adjusted | ABG | IMV | 9362 | 2391 | 1.28 (1.14, 1.43) | 1.27 (1.14, 1.43) |
| Unweighted adjusted | ABG | NIV | 9362 | 896 | 1.08 (0.89, 1.31) | 2.90 (2.46, 3.41) |
| Unweighted adjusted | ABG | Death (60d) | 9362 | 1645 | 1.72 (1.51, 1.97) | 1.38 (1.20, 1.57) |
| Unweighted adjusted | ABG | Hypercapnic RF | 9362 | 992 | 0.96 (0.76, 1.21) | 6.78 (5.75, 8.03) |
| Unweighted adjusted | VBG | IMV | 7460 | 1162 | 1.30 (1.11, 1.53) | 1.54 (1.32, 1.80) |
| Unweighted adjusted | VBG | NIV | 7460 | 554 | 0.98 (0.76, 1.25) | 2.74 (2.22, 3.37) |
| Unweighted adjusted | VBG | Death (60d) | 7460 | 1044 | 1.71 (1.45, 2.01) | 1.35 (1.14, 1.59) |
| Unweighted adjusted | VBG | Hypercapnic RF | 7460 | 692 | 0.66 (0.49, 0.89) | 7.21 (5.92, 8.82) |
| IPW adjusted | ABG | IMV | 9362 | 2391 | 1.26 (1.11, 1.43) | 1.20 (1.05, 1.36) |
| IPW adjusted | ABG | NIV | 9362 | 896 | 1.15 (0.93, 1.44) | 3.01 (2.50, 3.61) |
| IPW adjusted | ABG | Death (60d) | 9362 | 1645 | 1.63 (1.41, 1.89) | 1.31 (1.13, 1.52) |

| method | group | outcome_label | n_model | events | Low vs normal | High vs normal |
|-------------------|-------|----------------|---------|--------|-------------------|-------------------|
| IPW adjusted | ABG | Hypercapnic RF | 9362 | 992 | 1.08 (0.83, 1.40) | 7.79 (6.46, 9.39) |
| IPW adjusted | VBG | IMV | 7460 | 1162 | 1.41 (1.18, 1.69) | 1.51 (1.26, 1.80) |
| IPW adjusted | VBG | NIV | 7460 | 554 | 0.96 (0.73, 1.27) | 2.95 (2.35, 3.71) |
| IPW adjusted | VBG | Death (60d) | 7460 | 1044 | 1.76 (1.46, 2.11) | 1.52 (1.26, 1.83) |
| IPW adjusted | VBG | Hypercapnic RF | 7460 | 692 | 0.72 (0.52, 1.00) | 7.99 (6.39, 9.98) |
| IPW + MI adjusted | ABG | IMV | 9362 | 2391 | 1.27 (1.11, 1.45) | 1.26 (1.10, 1.44) |
| IPW + MI adjusted | ABG | NIV | 9362 | 896 | 1.14 (0.91, 1.44) | 3.06 (2.52, 3.70) |
| IPW + MI adjusted | ABG | Death (60d) | 9362 | 1645 | 1.63 (1.40, 1.90) | 1.30 (1.12, 1.52) |
| IPW + MI adjusted | ABG | Hypercapnic RF | 9362 | 992 | 1.03 (0.79, 1.35) | 7.66 (6.29, 9.32) |
| IPW + MI adjusted | VBG | IMV | 7460 | 1162 | 1.50 (1.23, 1.82) | 1.54 (1.27, 1.87) |
| IPW + MI adjusted | VBG | NIV | 7460 | 554 | 0.91 (0.68, 1.24) | 2.83 (2.21, 3.63) |
| IPW + MI adjusted | VBG | Death (60d) | 7460 | 1044 | 1.75 (1.42, 2.15) | 1.42 (1.16, 1.74) |
| IPW + MI adjusted | VBG | Hypercapnic RF | 7460 | 692 | 0.68 (0.46, 1.00) | 6.53 (5.08, 8.38) |

3.8 Manuscript outputs summary

```
# Cohort flow / sample sizes
flow_tbl <- tibble::tibble(
  metric = c(
    "Full cohort (raw)",
    "Analytic subset",
    "ABG tested",
    "ABG with PaCO2",
    "VBG tested",
    "VBG with VBG CO2"
  ),
  n = c(
    nrow(stata_data),
    nrow(subset_data),
    sum(subset_data$has_abg == 1, na.rm = TRUE),
    sum(subset_data$has_abg == 1 & !is.na(subset_data$paco2), na.rm = TRUE),
    sum(subset_data$has_vbg == 1, na.rm = TRUE),
    sum(subset_data$has_vbg == 1 & !is.na(subset_data$vbg_co2), na.rm = TRUE)
  )
)
```

```

    )
)
render_table_pdf(flow_tbl, "Cohort flow summary", "cohort_flow_summary",
                 preview_rows = 10, digits = 0)

```

Table 16: Cohort flow summary

| metric | n |
|-------------------|--------|
| Full cohort (raw) | 833476 |
| Analytic subset | 25852 |
| ABG tested | 9362 |
| ABG with PaCO2 | 9362 |
| VBG tested | 7460 |
| VBG with VBG CO2 | 7460 |

```

# Event counts by cohort
outcome_vars <- c("imv_proc", "niv_proc", "death_60d", "hypercap_resp_failure")
outcome_labels <- c(
  imv_proc = "IMV",
  niv_proc = "NIV",
  death_60d = "Death (60d)",
  hypercap_resp_failure = "Hypercapnic RF"
)
event_tbl <- dplyr::bind_rows(
  lapply(outcome_vars, function(o) {
    dplyr::tibble(
      outcome = outcome_labels[[o]],
      group = "ABG",
      n = sum(subset_data$has_abg == 1 & !is.na(subset_data[[o]]), na.rm = TRUE),
      events = sum(subset_data$has_abg == 1 & subset_data[[o]] == 1, na.rm = TRUE)
    )
  }),
  lapply(outcome_vars, function(o) {
    dplyr::tibble(
      outcome = outcome_labels[[o]],
      group = "VBG",
      n = sum(subset_data$has_vbg == 1 & !is.na(subset_data[[o]]), na.rm = TRUE),
      events = sum(subset_data$has_vbg == 1 & subset_data[[o]] == 1, na.rm = TRUE)
    )
  })
)

```

```

    events = sum(subset_data$has_vbg == 1 & subset_data[[o]] == 1, na.rm = TRUE)
  )
})
)
event_tbl <- event_tbl |>
  dplyr::mutate(pct = ifelse(n > 0, 100 * events / n, NA_real_))
render_table_pdf(event_tbl,
  "Outcome counts by cohort (ABG/VBG tested)",
  "outcome_counts_by_cohort",
  preview_rows = 20,
  digits = 1)

```

Table 17: Outcome counts by cohort (ABG/VBG tested)

| outcome | group | n | events | pct |
|----------------|-------|------|--------|------|
| IMV | ABG | 9362 | 2391 | 25.5 |
| NIV | ABG | 9362 | 896 | 9.6 |
| Death (60d) | ABG | 9362 | 1645 | 17.6 |
| Hypercapnic RF | ABG | 9362 | 992 | 10.6 |
| IMV | VBG | 7460 | 1162 | 15.6 |
| NIV | VBG | 7460 | 554 | 7.4 |
| Death (60d) | VBG | 7460 | 1044 | 14.0 |
| Hypercapnic RF | VBG | 7460 | 692 | 9.3 |

```

# Weighting diagnostics (non-MI weights)
stopifnot(all(c("w_abg", "w_vbg", "ps_abg", "ps_vbg") %in% names(subset_data)))
wt_abg <- subset_data$w_abg[subset_data$has_abg == 1]
wt_vbg <- subset_data$w_vbg[subset_data$has_vbg == 1]
ps_abg <- subset_data$ps_abg[subset_data$has_abg == 1]
ps_vbg <- subset_data$ps_vbg[subset_data$has_vbg == 1]
trunc_abg <- subset_data$trunc_abg[subset_data$has_abg == 1]
trunc_vbg <- subset_data$trunc_vbg[subset_data$has_vbg == 1]

wt_sum <- dplyr::bind_rows(
  weight_summary(wt_abg, ps = ps_abg, ps_floor = ps_floor_abg,
                 truncated = trunc_abg) |>
  dplyr::mutate(group = "ABG"),

```

```

    weight_summary(wt_vbg, ps = ps_vbg, ps_floor = ps_floor_vbg,
                   truncated = trunc_vbg) |>
    dplyr::mutate(group = "VBG")
)
wt_sum_display <- wt_sum |>
  dplyr::select(group, n, ess, min, p99, max, trunc_rate, ps_p01) |>
  dplyr::rename(
    trunc = trunc_rate,
    p01_ps = ps_p01
)
render_table_pdf(wt_sum_display,
                 "Weighting diagnostics summary (non-MI)",
                 "weighting_diagnostics_non_mi",
                 preview_rows = 10,
                 preview_cols = c("group", "n", "ess", "min", "p99", "max", "trunc", "p01_ps"),
                 wide = TRUE,
                 digits = 3)

```

Preview (top 10 rows). Full table saved to Results/weighting_diagnostics_non_mi.csv.

Table 18: Weighting diagnostics summary (non-MI)

| group | n | ess | min | p99 | max | trunc | p01_ps |
|--------|-----|------|----------|-------|-------|-------|--------|
| 1%...1 | ABG | 9362 | 7199.526 | 0.468 | 3.330 | 3.334 | 0.01 |
| 1%...2 | VBG | 7460 | 5982.854 | 0.410 | 2.784 | 2.784 | 0.01 |

```

# Missingness + MI spec summary
miss_vars <- c(covars_gbm, "paco2", "vbg_co2",
                 "imv_proc", "niv_proc", "death_60d", "hypercap_resp_failure")
miss_vars <- intersect(miss_vars, names(subset_data_raw))
miss_tbl <- subset_data_raw |>
  dplyr::summarise(dplyr::across(dplyr::all_of(miss_vars), ~ mean(is.na(.)) * 100)) |>
  tidyr::pivot_longer(dplyr::everything(), names_to = "variable", values_to = "pct_missing") |>
  dplyr::arrange(dplyr::desc(pct_missing)) |>
  dplyr::slice_head(n = 10)
render_table_pdf(miss_tbl,

```

```

"Top 10 variables by missingness (pre-imputation)",
"missingness_top10",
preview_rows = 10,
digits = 1)

```

Table 19: Top 10 variables by missingness (pre-imputation)

| variable | pct_missing |
|------------|-------------|
| vbg_co2 | 71.1 |
| paco2 | 63.8 |
| serum_lac | 60.2 |
| curr_bmi | 57.0 |
| serum_phos | 53.2 |
| temp_new | 48.2 |
| hr | 36.4 |
| dbp | 29.9 |
| sbp | 29.7 |
| wbc | 17.6 |

```

mi_spec_tbl <- tibble::tibble(
  m = imp$m,
  maxit = MAXIT_MI,
  methods = paste(unique(imp$method[imp$method != ""]),
                 collapse = ", "),
  ps_model = MI_PS_METHOD,
  ps_spline_k = MI_PS_SPLINE_K,
  ps_glm_maxit = MI_GLM_MAXIT
)
render_table_pdf(mi_spec_tbl,
                 "MI specification (methods used)",
                 "mi_specification",
                 preview_rows = 10)

```

Table 20: MI specification (methods used)

| m | maxit | methods | ps_model | ps_spline_k | ps_glm_maxit |
|----|-------|---------|----------|-------------|--------------|
| 80 | 20 | pmm | glm_rcs4 | 4 | 25 |

3.8.1 14.3 Visualization: pooled three-level ORs

```
stopifnot(exists("abg_cat_results"), exists("vbg_cat_results"))

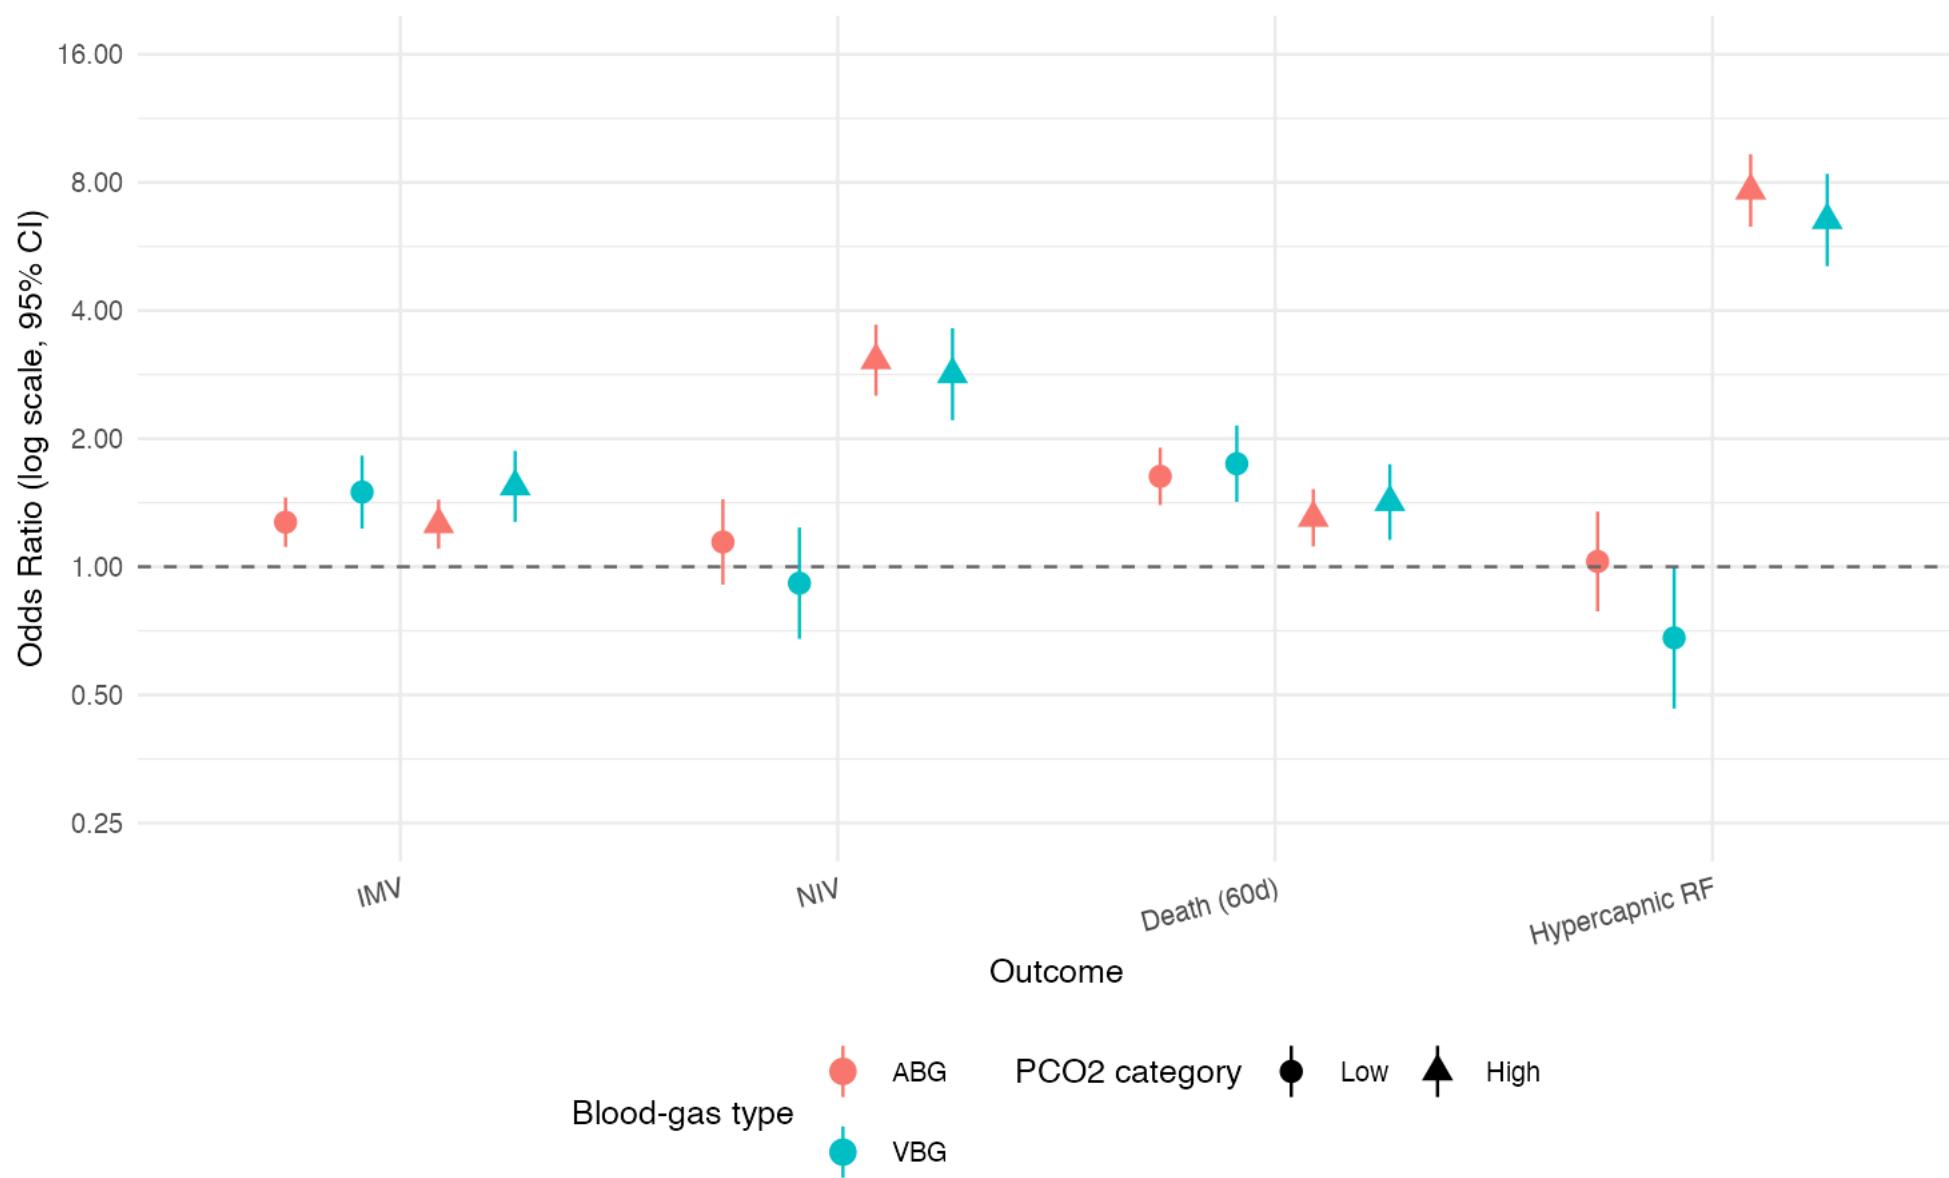
mi_combined_or_df <- dplyr::bind_rows(
  dplyr::mutate(abg_cat_results, group = "ABG"),
  dplyr::mutate(vbg_cat_results, group = "VBG")
) |>
  dplyr::mutate(
    outcome = factor(outcome,
                     levels = c("IMV", "NIV", "Death60d", "HCRF"),
                     labels = c("IMV", "NIV", "Death (60d)", "Hypercapnic RF")),
    group = factor(group, levels = c("ABG", "VBG"))
  )

mi_combined_or_df <- map_or_exposure(mi_combined_or_df, "or-plot-mi-weighted")
mi_plot_df <- mi_combined_or_df |>
  dplyr::mutate(estimate = OR, conf.low = LCL, conf.high = UCL) |>
  dplyr::select(-OR, -LCL, -UCL)

mi_plot_df <- build_or_plot_df(
  mi_plot_df,
  "or-plot-mi-weighted",
  expected_exposure_levels = CO2_CAT_CONTRAST_LEVELS
)
mi_axis_spec <- compute_or_axis_spec(list(mi_plot_df), lo_col = "conf.low", hi_col = "conf.high")

mi_p_or <- plot_or_safe(
  mi_plot_df,
  plot_name = "or-plot-mi-weighted",
  axis_spec = mi_axis_spec,
  title = "MI-pooled, IPW-adjusted odds ratios by PCO2 category (ABG vs VBG)"
)
print_plot_once(mi_p_or, "or-plot-mi-weighted", width = 7.5, height = 4.8)
```

MI-pooled, IPW-adjusted odds ratios by PCO₂ category (ABG vs VBG)



```
float_barrier()
```

3.8.2 15.3 Visualization

```
library(dplyr)
library(ggplot2)
library(patchwork)
library(purrr)

mi_ipw_rcs_forms <- list(
  "MI IPW spline (adjusted) ABG: IMV ~ CO2 spline + X"      = make_spline_fml("imv_proc", "paco2", adj_core),
  "MI IPW spline (adjusted) ABG: NIV ~ CO2 spline + X"      = make_spline_fml("niv_proc", "paco2", adj_core),
  "MI IPW spline (adjusted) ABG: Death60d ~ CO2 spline + X" = make_spline_fml("death_60d", "paco2", adj_core),
  "MI IPW spline (adjusted) ABG: HCRF ~ CO2 spline + X"     = make_spline_fml("hypercap_resp_failure", "paco2", adj_core),
  "MI IPW spline (adjusted) VBG: IMV ~ CO2 spline + X"       = make_spline_fml("imv_proc", "vbg_co2", adj_core),
  "MI IPW spline (adjusted) VBG: NIV ~ CO2 spline + X"       = make_spline_fml("niv_proc", "vbg_co2", adj_core),
  "MI IPW spline (adjusted) VBG: Death60d ~ CO2 spline + X" = make_spline_fml("death_60d", "vbg_co2", adj_core),
  "MI IPW spline (adjusted) VBG: HCRF ~ CO2 spline + X"     = make_spline_fml("hypercap_resp_failure", "vbg_co2", adj_core)
)
register_model_diagrams(mi_ipw_rcs_forms)

stopifnot(exists("abg_curves"), exists("vbg_curves"))

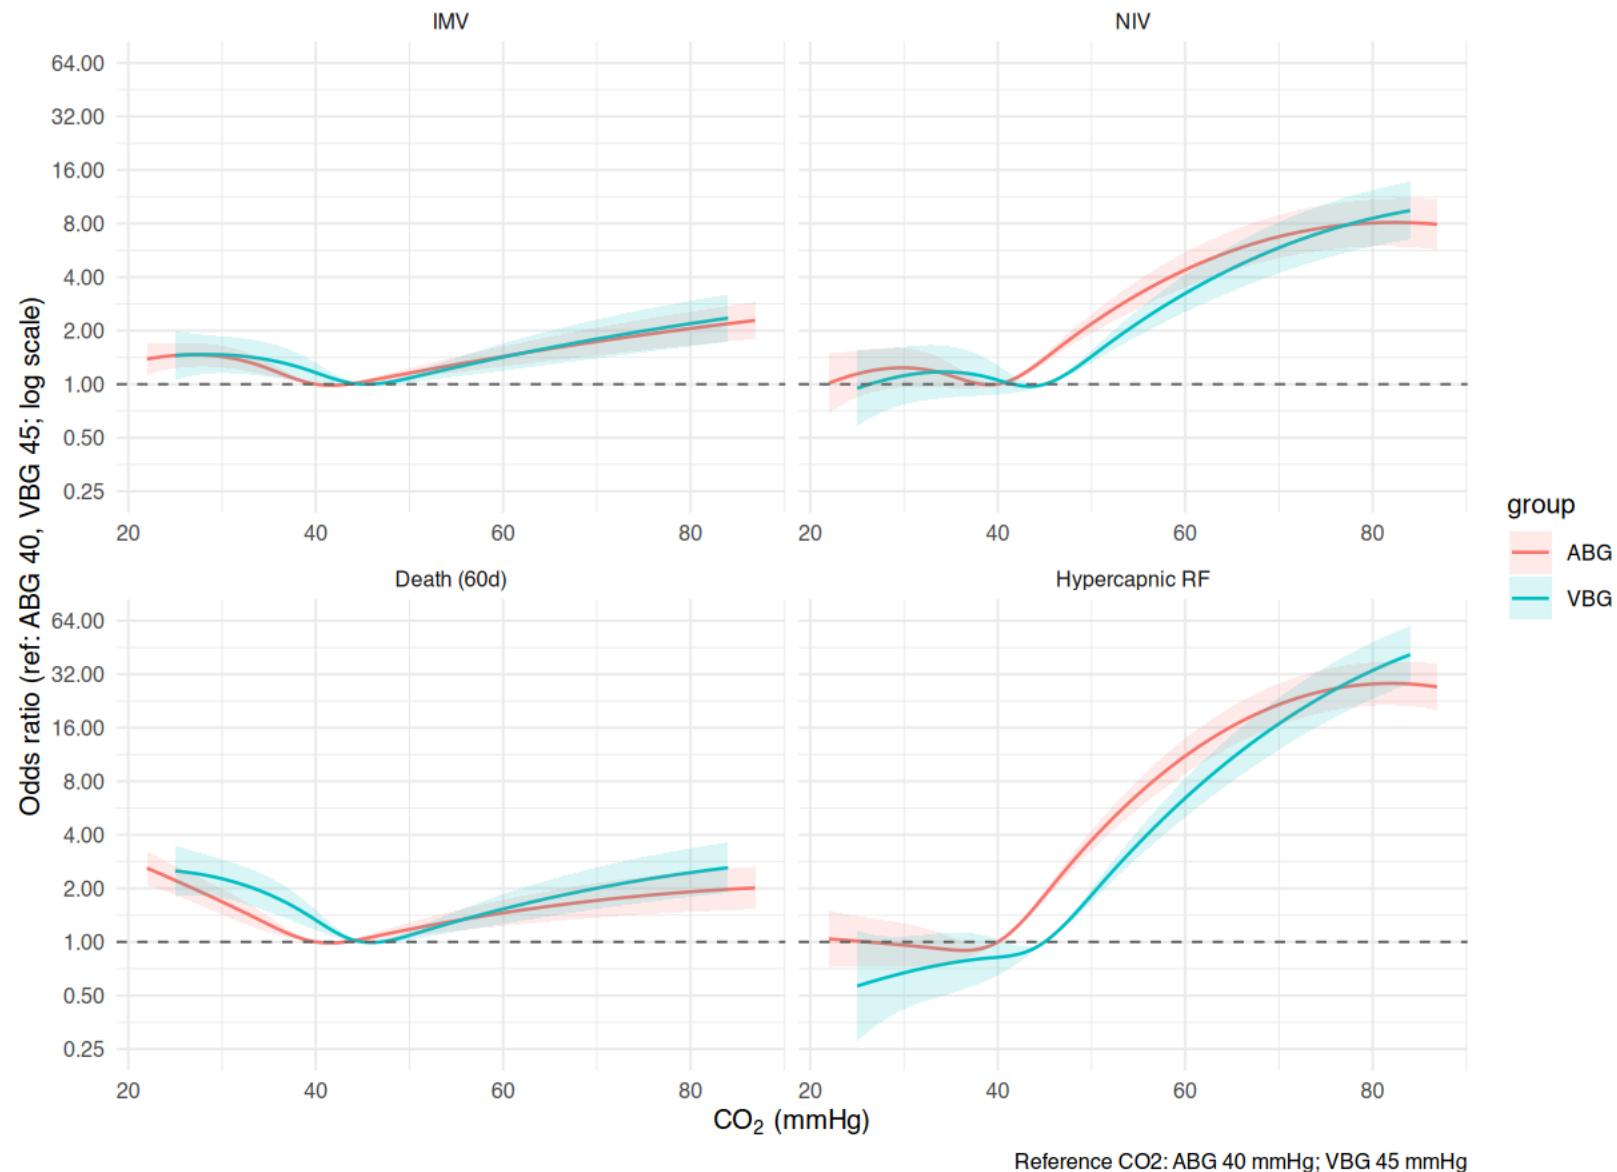
curve_abg <- abg_curves |>
  mutate(co2 = paco2) |>
  select(-paco2)
curve_vbg <- vbg_curves |>
  mutate(co2 = vbg_co2) |>
  select(-vbg_co2)
curve_all <- bind_rows(curve_abg, curve_vbg) |>
  mutate(outcome = factor(outcome,
                          levels = c("imv_proc", "niv_proc", "death_60d", "hypercap_resp_failure"),
                          labels = c("IMV", "NIV", "Death (60d)", "Hypercapnic RF")))
axis_mi_ipw <- compute_or_axis_spec(list(curve_abg, curve_vbg), lo_col = "LCL", hi_col = "UCL")
```

```

ggplot(curve_all, aes(x = co2, y = OR, color = group, fill = group)) +
  geom_line(linewidth = 0.6) +
  geom_ribbon(aes(ymin = LCL, ymax = UCL), alpha = 0.15, color = NA) +
  geom_hline(yintercept = 1, linetype = "dashed", color = "grey40") +
  or_axis_scale(axis_mi_ipw) +
  facet_wrap(~ outcome, scales = "free_x") +
  labs(
    title = expression(
      paste("MI-pooled, IPSW-adjusted spline odds ratios: ABG vs VBG CO"[2])
    ),
    x = expression(CO[2] ~ "(mmHg)" ),
    y = paste0("Odds ratio (ref: ABG ", ABG_CO2_REF, ", VBG ", VBG_CO2_REF, "; log scale)"),
    caption = paste0("Reference CO2: ABG ", ABG_CO2_REF, " mmHg; VBG ", VBG_CO2_REF, " mmHg")
  ) +
  theme_minimal(base_size = 10)

```

MI-pooled, IPSW-adjusted spline odds ratios: ABG vs VBG CO₂



```
float_barrier()
```

3.9 Diagnostics

3.9.1 MI convergence and mixing

3.9.2 MI stability across m

3.9.3 MI maxit sensitivity (sampled)

3.9.4 Balance diagnostics

Table 21: Target balance (top 10 by max |SMD|)

| group | variable | max_abs_pre | max_abs_post |
|-------|------------------|-------------|--------------|
| ABG | curr_bmi | 0.2691867 | 0.1938276 |
| ABG | encounter_type | 0.3704053 | 0.1924810 |
| ABG | location | 0.1663410 | 0.1140054 |
| ABG | sbp | 0.1792273 | 0.0872274 |
| ABG | serum_ca | 0.2235344 | 0.0851847 |
| ABG | age_at_encounter | 0.1055368 | 0.0815615 |
| ABG | copd | 0.0561648 | 0.0782714 |
| ABG | chf | 0.0722894 | 0.0742121 |
| ABG | dbp | 0.1615796 | 0.0728835 |
| ABG | sex | 0.0904038 | 0.0668771 |
| VBG | curr_bmi | 0.2734533 | 0.1797502 |
| VBG | location | 0.3988553 | 0.1506953 |
| VBG | serum_cl | 0.1455225 | 0.0789445 |
| VBG | race_ethnicity | 0.2049873 | 0.0784599 |
| VBG | hr | 0.1307602 | 0.0744904 |
| VBG | dm | 0.0822989 | 0.0716654 |
| VBG | dbp | 0.0927166 | 0.0671185 |
| VBG | sbp | 0.1213141 | 0.0556533 |
| VBG | ckd | 0.0689314 | 0.0537622 |

| group | variable | max_abs_pre | max_abs_post |
|-------|-----------|-------------|--------------|
| VBG | serum_lac | 0.0969344 | 0.0483619 |

Table 22: Worst target SMD rows across imputations (top 10)

| variable | level | type | smd_pre | smd_post | group | imp | abs_post |
|----------------|-----------|--------|------------|------------|-------|-----|-----------|
| encounter_type | Inpatient | factor | 0.3704053 | 0.1001128 | ABG | 45 | 0.1001128 |
| encounter_type | Emergency | factor | -0.3704053 | -0.1001128 | ABG | 45 | 0.1001128 |
| encounter_type | Inpatient | factor | 0.3704053 | 0.0979839 | ABG | 40 | 0.0979839 |
| encounter_type | Emergency | factor | -0.3704053 | -0.0979839 | ABG | 40 | 0.0979839 |
| encounter_type | Inpatient | factor | 0.3704053 | 0.0959636 | ABG | 38 | 0.0959636 |
| encounter_type | Emergency | factor | -0.3704053 | -0.0959636 | ABG | 38 | 0.0959636 |
| encounter_type | Emergency | factor | -0.3704053 | -0.0952833 | ABG | 51 | 0.0952833 |
| encounter_type | Inpatient | factor | 0.3704053 | 0.0952833 | ABG | 51 | 0.0952833 |
| encounter_type | Emergency | factor | -0.3704053 | -0.0952657 | ABG | 8 | 0.0952657 |
| encounter_type | Inpatient | factor | 0.3704053 | 0.0952657 | ABG | 8 | 0.0952657 |
| encounter_type | Emergency | factor | -0.0375275 | -0.0505432 | VBG | 26 | 0.0505432 |
| encounter_type | Inpatient | factor | 0.0375275 | 0.0505432 | VBG | 26 | 0.0505432 |
| encounter_type | Emergency | factor | -0.0375275 | -0.0493456 | VBG | 68 | 0.0493456 |
| encounter_type | Inpatient | factor | 0.0375275 | 0.0493456 | VBG | 68 | 0.0493456 |
| encounter_type | Emergency | factor | -0.0375275 | -0.0489478 | VBG | 36 | 0.0489478 |
| encounter_type | Inpatient | factor | 0.0375275 | 0.0489478 | VBG | 36 | 0.0489478 |
| encounter_type | Inpatient | factor | 0.0375275 | 0.0484923 | VBG | 9 | 0.0484923 |
| encounter_type | Emergency | factor | -0.0375275 | -0.0484923 | VBG | 9 | 0.0484923 |
| encounter_type | Inpatient | factor | 0.0375275 | 0.0479686 | VBG | 53 | 0.0479686 |
| encounter_type | Emergency | factor | -0.0375275 | -0.0479686 | VBG | 53 | 0.0479686 |

Table 23: Distribution of max |Target SMD| across imputations

| group | med | iqr | max |
|-------|-----------|-----------|-----------|
| ABG | 0.0911384 | 0.0034712 | 0.1001128 |
| VBG | 0.0440390 | 0.0037890 | 0.0505432 |

Table 24: Most frequent worst-balance terms (top 10 per group)

| group | term | n |
|-------|--------------------------|----|
| ABG | age_at_encounter | 80 |
| ABG | curr_bmi | 80 |
| ABG | encounter_type:Emergency | 80 |
| ABG | encounter_type:Inpatient | 80 |
| ABG | location:0 | 80 |
| ABG | location:1 | 80 |
| ABG | chf:1 | 76 |
| ABG | chf:0 | 74 |
| ABG | ckd:0 | 68 |
| ABG | ckd:1 | 61 |
| VBG | age_at_encounter | 80 |
| VBG | encounter_type:Emergency | 80 |
| VBG | encounter_type:Inpatient | 80 |
| VBG | location:3 | 80 |
| VBG | serum_ca | 80 |
| VBG | serum_phos | 51 |
| VBG | race_ethnicity:6 | 46 |
| VBG | temp_new | 45 |
| VBG | serum_cl | 43 |
| VBG | location:0 | 40 |

3.9.5 Outcome diagnostics

3.9.6 Diagnostics summary and audit

Preview (top 10 rows). Full table saved to Results/diagnostics_summary_display.csv.

3.9.7 Performance / runtime log

Table 25: Diagnostics summary (IPSW + MI)

| block | m | ess | trunc | med_max_smd | max_max_smd |
|--------------|----|----------|-------|-------------|-------------|
| ABG weights | 80 | 7199.526 | 0.01 | 0.091 | 0.100 |
| VBG weights | 80 | 5982.854 | 0.01 | 0.044 | 0.051 |
| ABG outcomes | 80 | 7199.526 | 0.01 | 0.091 | 0.100 |
| VBG outcomes | 80 | 5982.854 | 0.01 | 0.044 | 0.051 |

Table 26: Worst fitted-probability extremes (top 10)

| com- | analy- | im- | er- | po- | sis_vari- | puta- | con- | non- | warn- | top_warn | war | mes- | extre- | | | |
|-------|--------|--------|------------|------------------|-----------|-------|--------|--------|-------|----------|-----------|----------|-----------|-----|------|-------|
| stage | ment | ant | model_type | out- | come | n_use | events | verged | iter | sep_flag | conv_flag | min_flag | tx_flag | ing | sage | treme |
| out- | spline | ipw | spline | ABG hyper- | NA | 9362 | 992 | TRUE | 7 | TRUE | FALSE | 0 | 0.7398233 | 0 | NA | 0 |
| come | | | | cap_resp_failure | | | | | | | | | | | | |
| out- | spline | ipw | spline | ABG hyper- | NA | 9362 | 992 | TRUE | 7 | TRUE | FALSE | 0 | 0.7398233 | 0 | NA | 0 |
| come | | | | cap_resp_failure | | | | | | | | | | | | |
| out- | spline | mi_ipw | spline | ABG hyper- | 1 | 9362 | 992 | TRUE | 7 | TRUE | FALSE | 0 | 0.7837141 | 0 | NA | 0 |
| come | | | | cap_resp_failure | | | | | | | | | | | | |
| out- | spline | mi_ipw | spline | ABG hyper- | 2 | 9362 | 992 | TRUE | 7 | TRUE | FALSE | 0 | 0.7842440 | 0 | NA | 0 |
| come | | | | cap_resp_failure | | | | | | | | | | | | |
| out- | spline | mi_ipw | spline | ABG hyper- | 3 | 9362 | 992 | TRUE | 7 | TRUE | FALSE | 0 | 0.7898782 | 0 | NA | 0 |
| come | | | | cap_resp_failure | | | | | | | | | | | | |
| out- | spline | mi_ipw | spline | ABG hyper- | 4 | 9362 | 992 | TRUE | 7 | TRUE | FALSE | 0 | 0.7820752 | 0 | NA | 0 |
| come | | | | cap_resp_failure | | | | | | | | | | | | |
| out- | spline | mi_ipw | spline | ABG hyper- | 5 | 9362 | 992 | TRUE | 7 | TRUE | FALSE | 0 | 0.7788758 | 0 | NA | 0 |
| come | | | | cap_resp_failure | | | | | | | | | | | | |
| out- | spline | mi_ipw | spline | ABG hyper- | 6 | 9362 | 992 | TRUE | 7 | TRUE | FALSE | 0 | 0.7855521 | 0 | NA | 0 |
| come | | | | cap_resp_failure | | | | | | | | | | | | |

| com- | analy- | | | im- | | | | | | | | er- | |
|-------|-----------|--------|------------|----------------------------|-------|--------|--------|------|----------|------------|------------|-----------|-----------|
| po- | sis_vari- | | | puta- | | con- | | non- | | warn- | top_wa r n | mes- | |
| stage | ment | ant | model_type | put_outcome | n_use | events | verged | iter | sep_flag | onv_flag | min_phatx | flag_n | ing |
| out- | spline | mi_ipw | spline | ABG hyper-cap_resp_failure | 7 | 9362 | 992 | TRUE | 7 | TRUE FALSE | 0 | 0.7769311 | 0 NA NA 0 |
| out- | spline | mi_ipw | spline | ABG hyper-cap_resp_failure | 8 | 9362 | 992 | TRUE | 7 | TRUE FALSE | 0 | 0.7807939 | 0 NA NA 0 |

3.9.8 Performance / runtime log

Table 27: Top runtime steps (seconds)

| step_name | seconds | start_time | end_time | notes | run_id | run_mode | n_sub-set |
|---------------|-----------|----------------------------|----------------------------|-------------------------|----------------------|----------|-----------|
| mi_singl_pass | 195.22984 | 2026-02-04 07:39:15.373768 | 2026-02-04 07:42:30.603606 | m=80 | 20260204_072236pilot | | 25852 |
| mice_batch_36 | 18.20452 | 2026-02-04 07:36:52.950617 | 2026-02-04 07:37:11.155134 | batch=36; m=2; maxit=20 | 20260204_072236pilot | | 25852 |
| mice_batch_37 | 18.12006 | 2026-02-04 07:37:12.528671 | 2026-02-04 07:37:30.648726 | batch=37; m=2; maxit=20 | 20260204_072236pilot | | 25852 |
| mice_batch_35 | 18.06604 | 2026-02-04 07:36:33.511232 | 2026-02-04 07:36:51.577272 | batch=35; m=2; maxit=20 | 20260204_072236pilot | | 25852 |
| mice_batch_34 | 18.02548 | 2026-02-04 07:36:14.115993 | 2026-02-04 07:36:32.141469 | batch=34; m=2; maxit=20 | 20260204_072236pilot | | 25852 |
| mice_batch_38 | 18.01687 | 2026-02-04 07:37:31.977403 | 2026-02-04 07:37:49.994278 | batch=38; m=2; maxit=20 | 20260204_072236pilot | | 25852 |
| mice_batch_1 | 17.99264 | 2026-02-04 07:25:50.91151 | 2026-02-04 07:26:08.90415 | batch=1; m=2; maxit=20 | 20260204_072236pilot | | 25852 |
| mice_batch_40 | 17.78155 | 2026-02-04 07:38:10.337816 | 2026-02-04 07:38:28.119368 | batch=40; m=2; maxit=20 | 20260204_072236pilot | | 25852 |
| mice_batch_6 | 17.76175 | 2026-02-04 07:27:25.250658 | 2026-02-04 07:27:43.012407 | batch=6; m=2; maxit=20 | 20260204_072236pilot | | 25852 |

| step_name | seconds | start_time | end_time | notes | run_id | run_mode | n_sub-set |
|---------------|----------|-------------------------------|-------------------------------|----------------------------|----------------------|----------|-----------|
| mice_batch_14 | 17.74737 | 2026-02-04 07:29:55.877432 | 2026-02-04 07:30:13.624806 | batch=14; m=2; maxit=20 | 20260204_072236pilot | pilot | 25852 |
| mice_batch_18 | 17.71192 | 2026-02-04 07:31:11.903159 | 2026-02-04 07:31:29.615077 | batch=18; m=2; maxit=20 | 20260204_072236pilot | pilot | 25852 |
| mice_batch_39 | 17.67599 | 2026-02-04 07:37:51.287994 | 2026-02-04 07:38:08.963982 | batch=39; m=2; maxit=20 | 20260204_072236pilot | pilot | 25852 |
| mice_batch_10 | 17.65232 | 2026-02-04 07:28:40.65581 | 2026-02-04 07:28:58.308133 | batch=10; m=2; maxit=20 | 20260204_072236pilot | pilot | 25852 |
| mice_batch_31 | 17.65083 | 2026-02-04 07:35:17.437083 | 2026-02-04 07:35:35.087911 | batch=31; m=2; maxit=20 | 20260204_072236pilot | pilot | 25852 |
| mice_batch_15 | 17.64427 | 2026-02-04 07:30:15.026542 | 2026-02-04 07:30:32.67081 | batch=15; m=2; maxit=20 | 20260204_072236pilot | pilot | 25852 |

Table 28: Runtime summary (total + top 5 steps)

| step_name | seconds |
|----------------|-----------|
| TOTAL | 900.53355 |
| mi_single_pass | 195.22984 |
| mice_batch_36 | 18.20452 |
| mice_batch_37 | 18.12006 |
| mice_batch_35 | 18.06604 |
| mice_batch_34 | 18.02548 |

3.10 16) Save, export, and session info

```
audit_issues <- results_path("diagnostics_audit_issues.csv")
if (file.exists(audit_issues)) {
  issues_df <- read.csv(audit_issues)
  issues_df <- issues_df |>
    dplyr::arrange(factor(severity, levels = c("blocker", "high", "medium", "low")))) |>
```

```

dplyr::slice_head(n = 10)
knitr::kable(issues_df, caption = "Diagnostics audit: top issues (see Results/diagnostics_audit.md for full details)")
} else {
  cat("Diagnostics audit summary not available.\n")
}

```

Table 29: Diagnostics audit: top issues (see Results/diagnostics_audit.md for full details)

| com-sever-ity | com-po-ny | evidence_file | evidence_snippet | why_it_matters | recommended_fix |
|---------------|-----------|---|-----------------------------------|---|---|
| high | Bal-ance | Results/balance_target_imputation_summary.csv | ABG max SMD =0.100 | ABG target balance exceeds 0.10 threshold across imputations. | Revisit GBM tuning, covariate set, or truncation to improve ABG balance. |
| high | Out-come | Re-sults/model_fit_diagnostics.csv | sep_flag TRUE for 905 / 1324 fits | High rate of separation/near-separation can bias ORs and CIs. | Inspect flagged outcomes; consider penalized fits or check data sparsity. |

```

stopifnot(exists("abg_curves"), exists("vbg_curves"), exists("abg_coefs"), exists("vbg_coefs"))
saveRDS(
  list(
    abg_curves = abg_curves,
    vbg_curves = vbg_curves,
    abg_coefs = abg_coefs,
    vbg_coefs = vbg_coefs
  ),
  mi_pooled_file
)

```

Software: R 4.5.2 ; key packages: mice, WeightIt, cobalt, survey, rms.